

Course reader Basic Machine Learning for Bioinformatics 2021



Figure 1. Pretentious Machine Learning images on the covers of course readers, who can do without them? Taken from: <https://www.synaptica.com/machine-and-deep-learning/>

Table of Contents

Course reader Basic Machine Learning for Bioinformatics 2021.....	1
Introduction.....	3
Course structure.....	4
Before the course starts.....	5
Downloading practical materials and slides.....	5
Day 1 materials.....	6
Day 2 materials.....	8
Day 3 materials.....	9
Day 4 materials.....	11
Day 5 materials.....	12
Day 6 materials.....	14
Day 7 materials.....	16
Day 8 materials.....	16
Project description and instruction.....	17
Extra Resources.....	18
Words of gratitude.....	19
Stock photos of a woman threatening a goldfish with a gun.....	19

Introduction

This course focusses on the very basics of Machine Learning (ML). This means that we'll implement linear regression (yes, that's ML), logistic regression (which, confusingly, is a classification method), a simple neural network, K-means clustering and hierarchical clustering, and Principal Component Analysis (PCA) ourselves using only basic functions and packages like Numpy and Matplotlib for plotting. Along the way we look at overfitting and underfitting and related bias and variance, look at how we measure model performance, and other things to keep in mind.

Machine Learning algorithms are implemented using Linear Algebra computations. Linear Algebra is a whole domain of mathematics with many interesting applications. If you are a studying bioinformatics, you have already come across its use in 2D and/or multi-D systems of ODEs. A proper course about Linear Algebra can take weeks, and we don't have that time. Therefore, for our purposes, Linear Algebra is just a nice way of manipulating numbers (such as multiplying the weights for each feature in a multivariate linear regression with the feature values for all training examples) in a form that's efficient for computation. The main ideas from Linear Algebra that are relevant for us will be introduced in a small crash course during the lectures, supplemented with suggestions for videos and other resources for further study.

We focus on the basics because it is important that you get a feel for what's really going on. There is an aura of magic around ML, and the excitement around AI has reached a fever pitch. True, advances like AlphaFold 2 and some of Google's and Facebook's convolutional neural nets are true marvels of engineering and a sight to behold. Yet at its core, we're just shuffling numbers around. That's why we work on implementing things ourselves: to get a feel for what is really happening.

It is likely that most of what you'll be doing during your studies and later on is applied ML: taking a problem from your domain of study, using an ML algorithm on it, and thereby solving said problem. Finding clusters (subtypes) in single cell cancer sequencing data, training a classifier to predict these different types of cancers based on biopsy sequencing data, that sort of thing. If you're going to train a neural network, you won't implement it yourself, you'll use Keras and TensorFlow. If you're making a Support Vector Machine, Random Forest, Linear Regressor, using PCA or a clustering of any kind, you're going to use Scikit-learn (or R equivalents). There are free, open-source packages for almost every popular ML algorithm, and you'd be foolish not to use them. For this reason, in the second week, we give a bird's eye view of Scikit-learn and Keras, so you'll have a basic familiarity and know enough to dive deeper yourself after the course.

The rest of the second week consists of a two-day ML project on Kaggle, where you're going to compete in groups to build the best classifier for a biological dataset using Scikit-learn, topped off with a short presentation about what you did. Finally, on the 23rd in the afternoon, there's a pen and paper exam that will probe your knowledge about what you've done.

Course structure

General

The first three days of week 1 cover supervised learning methods. The last two days of week 1 cover unsupervised learning methods. Week 2 switches from implementing the methods ourselves to using modern established ML libraries (Scikit-learn, Keras and TensorFlow) that you'll use for your applied ML projects. There's a two day project on Kaggle where you work in teams to get the best classification performance on a certain biological dataset and present your work afterwards. The pen and paper exam is on Thursday the 23rd in the afternoon.

Daily structure

We start at 9:00. Lectures and short practicals applying the lecture concepts are interleaved in the morning session. This means ~1 hour of lecture, followed by ~45 minutes of direct application in programming exercises, repeated about 2-3 times. We have lunch from ~12:15-13:00. The afternoon practical starts at 13:00 and runs until (at most) 17:00. This schedule is approximate: this is the first time the course is given so mismatches in actual timing versus planned timing are to be expected! Note also that sometimes there'll be 2 lectures rather than 3, or only 2 (somewhat longer) practicals.

Topics Week 1

- Day 1: Linear regression, gradient descent, introduction to linear algebra
- Day 2: Logistic regression, regularisation, ROC curve, introduction to neural networks (NNs)
- Day 3: NN Backpropagation algorithm, convolutional neural networks explained
- Day 4: K-means clustering, hierarchical clustering, deep dive into phylogenetics
- Day 5: Problems with high-dimensional data, Principal Component Analysis (PCA)

Topics Week 2

- Day 1: Working with scikit-learn, introduction to Keras and TensorFlow, classifier calibration, invited speaker talk, project introduction and start
- Day 2: Working on the project
- Day 3: Working on the project, project presentation in the afternoon
- Day 4: Q&A session in the morning, exam in the afternoon

Before the course starts

Software requirements

Make sure you have a working installation of Python 3 (preferably the newest version, which is python 3.8 at the time of writing), and have installed [numpy](#), [matplotlib](#), [scikit-learn](#), and [TensorFlow](#) (which automatically also installs [Keras](#)) in your working environment. Rather than manually setting this up, please download [Anaconda](#). It comes with many of the most-used packages pre-installed, a graphical user interface to manage packages, and many other tools. You might want to make an environment for the course specifically, but that is optional. Environments allow you to keep projects separate: rather than installing packages globally, you install them only for a specific environment, in this case the environment you make for this course. See [here](#) if you want to do that.

Downloading practical materials and slides

All practical material, slides, and this reader are on GitHub [here](#). Clone the git repository locally such that you can work with the files. If you do not know how to do that, do the following:

1. If on Windows, install [Git for Windows](#); if on Mac, install [Git for Mac](#); if on Linux (Ubuntu distro), run `sudo apt update`, followed by `sudo apt-get install git`. Follow all instructions for your respective case.
2. Go to the folder where you want the files to be downloaded. In Windows, you can open Git Bash there by right clicking (Figure 2). This opens a console. In all cases, you can type `git clone`, followed by <https://github.com/DieStok/Basic-Machine-Learning-for-Bioinformatics.git> in the console in the directory where you want the files. This will automatically download everything there.

3. Alternatively, if this doesn't work, you can download a .zip of all files manually, by clicking on the area indicated by the red rectangle in Figure 3.

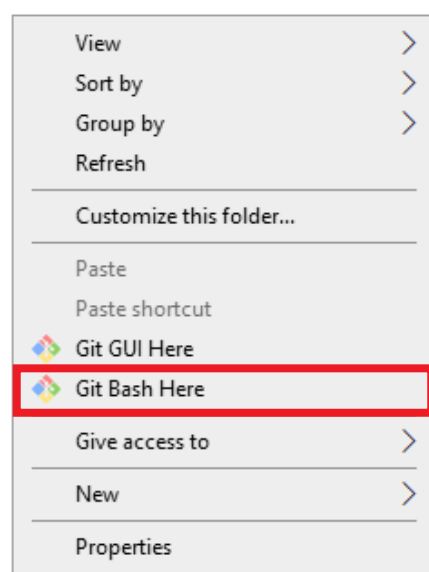


Figure 2. Opening Git Bash in a folder by right clicking in Windows.

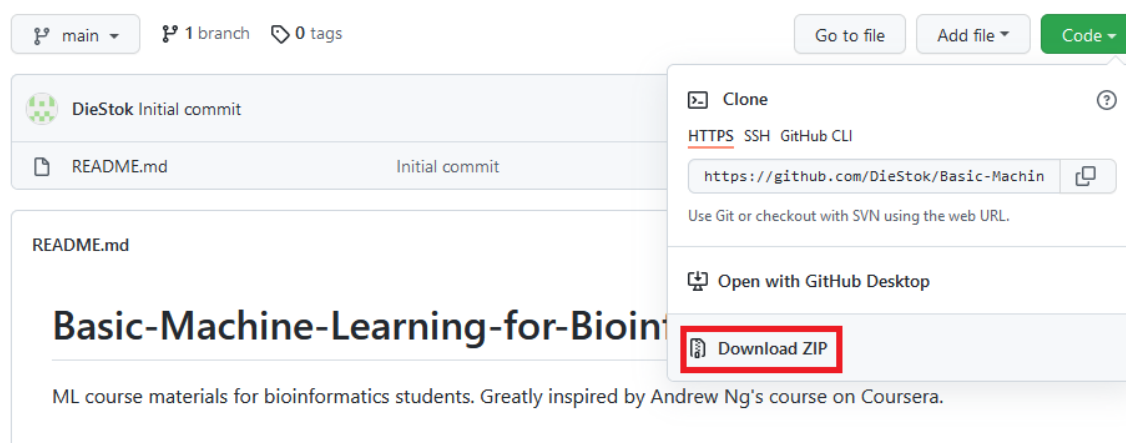


Figure 3. Manually downloading the files from the GitHub page.

Day 1 materials

Description

Today we focus on the most basic supervised ML algorithm: linear regression. We talk about the equations for linear regression, its cost function, using gradient descent to minimise the cost function, going from univariate linear regression (one feature, x , one predicted quantity based on it, y) to multivariate linear regression (predicting, say, disease severity based on levels of 10 biomarkers in the blood). We also discuss the bias-variance trade-off and using cross-validation for measuring ML performance.

Recommended videos for review:

[Video 1](#); [Video 2](#); [Video 3](#); [Video 4](#) (~55 minutes)

Extra resources <https://www.youtube.com/watch?v=ISBGFY-gBug>

Videos

- Linear Regression short StatQuest: <https://www.youtube.com/watch?v=PaFPbb66DxQ>
- More in-depth StatQuest Linear Regression: https://www.youtube.com/watch?v=nk2CQITm_eo
- Cross-validation: <https://www.youtube.com/watch?v=fSytzGwwBVw>
- Hyperparameters: <https://www.youtube.com/watch?v=VTE2KlfoO3Q>
- Bias and variance (Video 2 above): <https://www.youtube.com/watch?v=EuBBz3bI-aA>
- Linear algebra: https://www.youtube.com/watch?v=fNk_zzaMoSs&list=PLZHQObOWTQDPD3MizzM2xVFitgF8hE_ab (3Blue1Brown playlist)

Reading

- Cross-validation: <https://machinelearningmastery.com/k-fold-cross-validation/>
- Hyperparameters: <https://machinelearningmastery.com/difference-between-a-parameter-and-a-hyperparameter/>
- In-depth treatment of bias-variance trade-off: <http://scott.fortmann-roe.com/docs/BiasVariance.html>
- Linear algebra: <https://cbmm.mit.edu/sites/default/files/documents/algebra.pdf> (slide deck with a lot of depth) ; https://minireference.com/static/tutorials/linear_algebra_in_4_pages.pdf
- Math symbols primer: <https://amitnss.com/2019/08/math-for-programmers/>
- L2-norm: <https://kawahara.ca/what-does-the-l2-or-euclidean-norm-mean/>

Other

- Bias-variance trade-off and overfitting discussed with interactive examples: https://machinelearningcompass.com/model_optimization/bias_and_variance/
- Jupyter Guide to Linear Algebra: https://bvanderlei.github.io/jupyter-guide-to-linear-algebra/Matrix_Algebra.html
- Linear algebra course Khan Academy: <https://www.khanacademy.org/math/linear-algebra>

Day 2 materials

Description

Today we focus on a simple ML algorithm for classification: predicting discrete classes from data (as opposed to continuous variables, which we covered yesterday). This is logistic regression, which you've probably heard of before. It will turn out that the implementation is extremely similar to that of linear regression, just with a different cost function. After this, we will look into regularisation to prevent overfitting, and get started on the basics of neural networks.

Recommended videos for review:

[Video 1](#); [Video 2](#); [Video 3](#); [Video 4](#) (~60 minutes)

Extra resources

Videos

- Regularisation: <https://www.youtube.com/watch?v=KvtGD37Rm5I> ; <https://www.youtube.com/watch?v=Q81RR3yKn30>
- Neural networks: <https://www.youtube.com/watch?v=HGwBXDKFk9I&list=PLblh5JKOoLUIxGDQs4LFFD--41Vzf-ME1> (StatQuest playlist)
- (AUC) ROC: <https://www.youtube.com/watch?v=4jRBRDbJemM>
- Nested cross-validation: <https://www.youtube.com/watch?v=LpOxxBeggM0>; <https://www.youtube.com/watch?v=DuDtXtKNpZs>

Reading

- Logistic regression: <https://towardsdatascience.com/logistic-regression-detailed-overview-46c4da4303bc>
- Regularisation: <https://explained.ai/regularization/index.html> (highly recommended for a *deep* understanding)
- Neural networks: <http://neuralnetworksanddeeplearning.com/>
- (AUC) ROC and (AUC) PRC: <https://glassboxmedicine.com/2019/02/23/measuring-performance-auc-auROC/> ; <https://glassboxmedicine.com/2019/03/02/measuring-performance-auprc/> ; <https://machinelearningmastery.com/roc-curves-and-precision-recall-curves-for-classification-in-python/>

Other

Day 3 materials

Description

Today we focus on neural networks in a bit more depth. Specifically, we take a look at backpropagation and how it allows us to optimise the weights and biases of a neural network. You'll be implementing backpropagation yourself and training a simple network. However, simple (or dense) neural networks are not what have made the huge strides in image recognition (and some other areas) in the last ~10 years possible. For that, we need to look to convolutional neural networks. We'll discuss what convolution is, and the core ideas that make these networks tick. You'll see how you can train a simple convolutional neural net quickly and easily in Keras.

Recommended videos for review:

[Video 1](#); [Video 2](#); [Video 3](#) (~50 minutes)

Extra resources

Videos

- Whole Neural Network playlist StatQuest: <https://www.youtube.com/watch?v=CqOfi41LfDw&list=PLblh5JKOoLUixGDQs4LFFD--41Vzf-ME1>
- Backpropagation: <https://www.youtube.com/watch?v=tIeHLnjs5U8> ; https://www.youtube.com/watch?v=8d6jf7s6_Qs&t=0s ; <https://www.youtube.com/watch?v=GlcnxUIrtek&list=PLiaHhY2iBX9hdHaRr6b7XevZtgZRa1PoU&index=4>
- Numerical gradient checking: <https://www.youtube.com/watch?v=pHMzNW8Agq4&list=PLiaHhY2iBX9hdHaRr6b7XevZtgZRa1PoU&index=5>
- Batch normalisation: <https://youtu.be/DtEq44FTPM4?t=127>
- Dropout: <https://www.youtube.com/watch?v=ARq74QuavAo> ; <https://www.youtube.com/watch?v=vAVOY8frLIQ> (quite in-depth)

Reading

- Backpropagation: <http://neuralnetworksanddeeplearning.com/chap2.html> ; <https://mattmazur.com/2015/03/17/a-step-by-step-backpropagation-example/>
- Neural networks in general: <http://neuralnetworksanddeeplearning.com>
- SENet architecture (complicated!): <https://github.com/hujie-frank/SENet>
- Understanding convolutions intuitively: <https://towardsdatascience.com/intuitively-understanding-convolutions-for-deep-learning-1f6f42face1>
- Batch normalisation: <https://towardsdatascience.com/batch-normalisation-explained-5f4bd9de5feb>
- Dropout: <https://towardsdatascience.com/simplified-math-behind-dropout-in-deep-learning-6d50f3f47275>

- Learning logical functions with neurons illustrated (until halfway through subheading 4): [Solving XOR with a single Perceptron | by Lucas Araújo | Medium](#)

Other

- Interactive TensorFlow Neural Network: <http://playground.tensorflow.org/>

Day 4 materials

Description

Today we change gears completely and switch from supervised to unsupervised learning. In unsupervised learning, there are no labels, no known training data on which to train our algorithms. Instead, we want to find some structure in the data, without knowing what we are looking for: clustering. This means we must contend with the fact that there is no 'correct' clustering, and that depending on your exact criterion for clustering, there may be many possible clusters to make. We'll look at prototype clustering (K-means as an example) and hierarchical clustering. We'll then take a slightly in-depth look at how we use clustering in phylogeny.

Recommended videos for review:

[Video 1](#); [Video 2](#); [Video 3](#); [Video 4](#) (~45 minutes)

Extra resources

Videos

- Clustering algorithms overview: https://www.youtube.com/watch?v=Se28XHI2_xE
- K-means maths: <https://www.youtube.com/watch?v=Z2VXrwMYHFk> ;
<https://www.youtube.com/watch?v=0MQEt10e4NM>
- Lecture series hierarchical clustering: https://www.youtube.com/watch?v=GVz6Y8r5AkY&list=PLBv09BD7ez_7qIbBhyQDr-LAKWUeycZtx

Reading

- Clustering algorithms overview: <https://towardsdatascience.com/the-5-clustering-algorithms-data-scientists-need-to-know-a36d136ef68>
- K-means clustering: <https://towardsdatascience.com/k-means-clustering-algorithm-applications-evaluation-methods-and-drawbacks-aa03e644b48a>
- Hierarchical clustering:
https://www.tutorialspoint.com/machine_learning_with_python/machine_learning_with_python_clustering_algorithms_hierarchical.htm

Other

Day 5 materials

Description

Today we continue on with unsupervised learning, now focussing on dimensionality reduction and its flagship/most well-known method: Principal Component Analysis. We first take a look at the curse of dimensionality: how more dimensions cause problems by exponentially increasing the need for training data, inducing overfitting, and making it impossible to use distance metrics. Luckily, we have intelligent ways to reduce the dimensionality of data. Using StatQuest's excellent basis, we'll dive a bit deeper into performing PCA ourselves, and the mathematics behind it. Finally, with knowledge of PCA and multivariate linear regression in hand, we know enough to perform our very own GWAS, which we will do in the afternoon practical.

Recommended videos for review:

[Video 1](#) ; [Video 2](#) (28 minutes)

Extra resources

Videos

- Counter-intuitive phenomena in high dimensions: <https://www.youtube.com/watch?v=dr2sIoD7eeU>
- Worked example of PCA calculations: <https://www.youtube.com/watch?v=S51bTyIwxFs>
- PCA lecture series: https://www.youtube.com/watch?v=IbE0tbjy6JQ&list=PLBv09BD7ez_5_yapAg86Od6JeypkS4YM
- PCA in GWAS: <https://youtu.be/qMkuYrlh7jw?t=370>
- t-SNE and UMAP: <https://www.youtube.com/watch?v=NEaUSP4YerM> ; <https://www.youtube.com/watch?v=6BP181wGGP8>
- Eigenvalue and eigenvector calculations: <https://www.youtube.com/watch?v=TQvxWaQnrqI>
- Why eigenvalues correspond to variance on each component:

Reading

- Curse of dimensionality: <https://www.visiondummy.com/2014/04/curse-dimensionality-affect-classification/>
- PCA in GWAS: <https://pubmed.ncbi.nlm.nih.gov/16862161/> ; <https://stats.stackexchange.com/a/8780>
- PCA stackoverflow gems: <https://stats.stackexchange.com/questions/217995/what-is-an-intuitive-explanation-for-how-pca-turns-from-a-geometric-problem-wit?noredirect=1&lq=1> (mathy) ; <https://stats.stackexchange.com/questions/117695/why-is-the-eigenvector-in-pca-taken-to-be-unit-norm> (unit length vector) ; <https://stats.stackexchange.com/questions/2691/making-sense-of-principal-component->

[analysis-eigenvectors-eigenvalues/140579#140579](#) (simple explanation from the level you'd tell your grandmother to more of an expert)

- In-depth PCA tutorial with examples:
http://www.cs.otago.ac.nz/cosc453/student_tutorials/principal_components.pdf
(unfortunately some images are missing but other than that it is a great read!)

Other

- Interactive exploration of the UMAP technique: <https://pair-code.github.io/understanding-umap/>

Day 6 materials

Description

Today we switch gears again. You've now been duly introduced to the basics of ML: supervised and unsupervised learning, (nested) cross-validation, bias and variance, and implementation of various classifiers yourself. There are great open-source libraries that vastly speed up Machine Learning workflows and do all the heavy lifting for you. We didn't start with those because it is important to know what's really going on, but when applying ML in future projects you'd be foolish to do everything by yourself. Therefore, today we switch resolutely to Scikit-learn and Keras, and also discuss some other best practices (such as classifier calibration).

Recommended videos for review:

[Video 1](#) ; [Video 2 \(you can probably watch on 1.5 speed\)](#); [Video 3](#) (~5 hours 10 minutes: You don't need to watch all of this! The latter two videos are full courses on scikit-learn and Keras which you can browse as you see fit.)

Extra resources

Videos

- Scikit-learn pipelines tutorial (with code): <https://www.youtube.com/watch?v=XvnkUg1yVmk> ; <https://www.youtube.com/watch?v=9vz0n1cyMbc>
- Freecodecamp 2+-hour Keras introduction: <https://www.youtube.com/watch?v=qFJeN9V1ZsI>
- What a tensor is: <https://www.youtube.com/watch?v=f5liqUk0ZTw>
- Optimizers (gradient descent, stochastic gradient descent, mini-batch gradient descent, AdaGrad, Adam): <https://www.youtube.com/watch?v=mdKjMPmcWjY>
- Pandas: <https://www.youtube.com/watch?v=vmEHCJofslg>

Reading

- Hyperparameter optimisation strategies (GridSearchCV and/or RandomSearchCV): <https://machinelearningmastery.com/hyperparameter-optimization-with-random-search-and-grid-search/> ; https://scikit-learn.org/stable/modules/grid_search.html
- One-hot encoding <https://machinelearningmastery.com/why-one-hot-encode-data-in-machine-learning/>
- Nested cross-validation: <https://weina.me/nested-cross-validation/> ; https://scikit-learn.org/stable/auto_examples/model_selection/plot_nested_cross_validation_iris.html#sphx-glr-auto-examples-model-selection-plot-nested-cross-validation-iris-py ; <https://stats.stackexchange.com/questions/65128/nested-cross-validation-for-model-selection>

- Global average pooling: <https://adventuresinmachinelearning.com/global-average-pooling-convolutional-neural-networks/> ; <https://paperswithcode.com/method/global-average-pooling>
- Keras training loop from scratch: https://keras.io/guides/writing_a_training_loop_from_scratch/
- RMSprop (explanation and Python implementation): <https://machinelearningmastery.com/gradient-descent-with-rmsprop-from-scratch/>
- Adam (standard optimisation method deep learning): <https://machinelearningmastery.com/adam-optimization-algorithm-for-deep-learning/>
- Mini-batch gradient descent: <https://towardsdatascience.com/batch-mini-batch-stochastic-gradient-descent-7a62ecba642a>
- Tuning models with KerasTuner: https://keras.io/guides/keras_tuner/getting_started/
- Gradient descent total overview: <https://runder.io/optimizing-gradient-descent/>
- Detailed hyperparameter optimisation guide: <https://nanonets.com/blog/hyperparameter-optimization/>

Other

- Adam gradient optimizer slide deck: https://moodle2.cs.huji.ac.il/nu15/pluginfile.php/316969/mod_resource/content/1/adam_pres.pdf
- Different optimizers visualised (very cool, though does not translate 1-on-1 to the real high-dimensional case): <https://github.com/Jaewan-Yun/optimizer-visualization>

Day 7 materials

Description

We start with a short lecture lightly explaining scikit-learn and Keras. After a short practical about them, you can work on the project. First step through the guided exercises, where you are asked to train a Random Forest classifier and an SVM (with polynomial and radial basis function) on the data. For a refresher or first introduction to these two:

Random Forests: https://www.youtube.com/watch?v=J4Wdy0Wc_xQ;
<https://www.youtube.com/watch?v=sQ870aTKqiM>

SVMs: <https://www.youtube.com/watch?v=efR1C6CvhmE>; <https://www.youtube.com/watch?v=Toet3EiSFcM>; https://www.youtube.com/watch?v=Qc5IyLW_hns

After that, the sky is the limit. Try to get the best classification performance on the data we give you. You'll have to deal with missing values, imbalanced data (see [imbalanced-learn](#)), nested cross-validation, use pipelines ([1](#), [2](#), [3](#)), etc. to properly train your classifier. Also think about trying different types, and even combining separate classifiers into one voting classifier ([1](#)). Good luck!

Day 8 materials

Description

Work some more on your classifier. We will have group presentations from 15:00-17:00. They should showcase your approach, snags you ran into and how you did or did not solve them, and what you learned. Grading of the project will be mostly on the basis of how well you implemented the basics: did you correctly do the guided exercises, did you correctly cross-validate, did you correctly normalise and correct missing values, etc. A small component will be your creativity and use of advanced methods (if any), and another small component how clear your presentation was. The Kaggle competition closes at 12:30. Be sure to look at least as professional and pleased with yourself when presenting as this chimpanzee:



Figure 4: Taken from: <http://media.gettyimages.com/photos/male-chimpanzee-in-business-clothes-picture-id184941527>

Project description and instruction

Introduction

After all you've learned, the proof should be in the pudding. Which is to say: no better way to test what you've learned than by applying it. For that, we host an InClass competition on [Kaggle](#). Kaggle is a huge online machine learning platform that has thousands of (cleaned) datasets from many domains and regularly hosts machine learning competitions. Competitors share their code and approach, can discuss why they did what they did, and learn from each other. There are also tutorials and many plug-and-play examples of ML pipelines. A good place to know for the aspiring ML practitioner!

You can sign up to the competition [here](#). Just as a reminder: manipulating tabular data in Python means [pandas](#). [Here](#) is a pandas cheat sheet. You can also search through a slew of data science cheat sheets on Kaggle [here](#). It is highly recommended to look through these. Be sure to share good ones with your fellow students!

Instructions

You will collaborate in teams of two or three people. Ideally, these are the same groups that you have been in for asking questions since the beginning, but making new groups is allowed. Since it's rather a big step from the practice problems to this huge Kaggle dataset, we will start with some guided exercises. You are required to train a Random Forest classifier and an SVM on this data, with the correct nested cross-validation schema for finding the optimal hyperparameters. You'll also be required to answer and hand in some questions about the data structure and meaning of the different data files.

After that, you are free to do whatever. Try different classifiers, train a neural network, combine different classifiers through voting, use unsupervised learning to find some structure and investigate that further: your imagination (and time, and programming prowess, to be fair) is the limit. The goal is to get the best predictive performance on **Something**

We will walk you through a sample submission process on Kaggle, so you know what to do and how to do it.

Beginning exercises

The questions you need to answer are here (**VOEG LINK IN ZODRA OP GITHUB STAAT**).
Voor nu heet de file **BESTAAT NOG NIET**

Extra Resources

Basic

- Simple overview of steps of Machine Learning: <https://www.youtube.com/watch?v=nKW8Ndu7Mjw> + <https://towardsdatascience.com/the-7-steps-of-machine-learning-2877d7e5548e>
- Numpy video tutorial: <https://www.youtube.com/watch?v=GB9ByFAIAH4>
- Pandas video tutorial: <https://www.youtube.com/watch?v=vmEHCJofslg>
-

Words of gratitude

I am indebted to Dr. Jeroen de Ridder for expert support in preparing the lectures and practical material. The teachings in the course follow [Andrew Ng's Coursera course](#) in much of the material, which I highly recommend for further learning and consolidation of the material presented here. The PCA section is based mostly on [the lecture series by Victor Lavrenko](#). I highly recommend both this lecture series and the other content on this channel for further viewing. [StatQuest](#) is invaluable for intuitive and simple explanations of the ideas behind various ML methods, stripped of their mathy mystery. I'd like to thank the countless volunteers who make and maintain open-source libraries like Matplotlib, Keras, and Scikit-learn and prevent the world from falling apart ([relevant XKCD](#)). Finally, I'd like to thank Prof. Dr. Berend Snel and Dr. Bas van Breukelen for giving me the opportunity to make this course (and/or cobble it together based on many different sources!).

Stock photos of a woman threatening a goldfish with a gun

What course reader can finish without them? There's a whole slew of these. Amazing. That's life for ya. No ML algorithm can hope to explain why this is the case. Note: I am against harming goldfish.



www.shutterstock.com - 55445056



www.shutterstock.com - 55445089



Download from
Dreamstime.com
The mentioned content is for personal use only.

1488417
igor Magpies | Dreamstime.com



shutterstock

IMAGE ID: 1488417
www.shutterstock.com