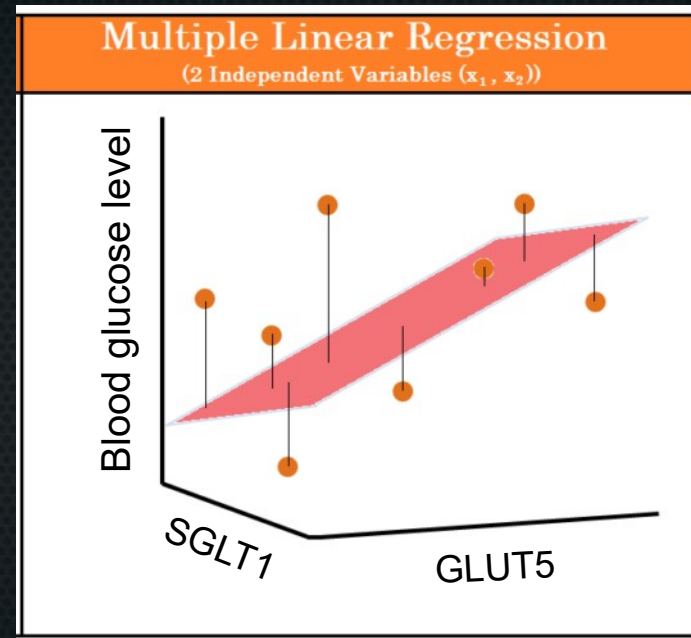


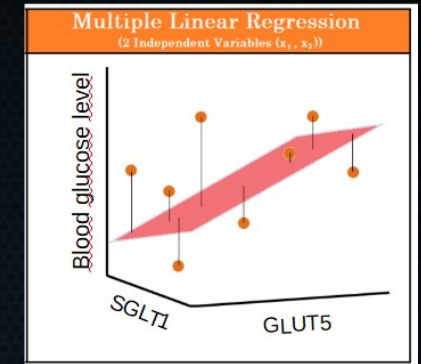
Multiple variables



Multiple variables

- Simple:

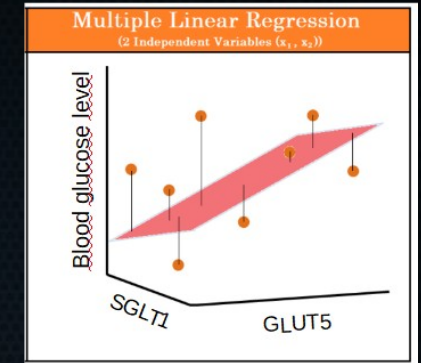
$$h_{\theta}(x) = \theta_0 + \theta_1 x_1 + \theta_2 x_2$$



Multiple variables

- Simple:

$$h_{\theta}(x) = \theta_0 + \theta_1 x_1 + \theta_2 x_2 + \dots + \theta_n x_n$$

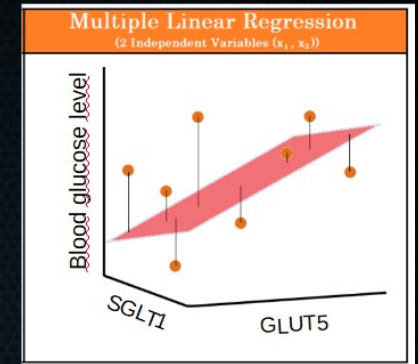


Multiple variables

- Simple:

$$h_{\theta}(x) = \theta_0 + \theta_1 x_1 + \theta_2 x_2 + \dots + \theta_n x_n$$

$$J(\theta_0, \theta_1, \dots, \theta_n) = \frac{1}{2m} \sum_{i=1}^m (h_{\theta}(x) - y^{(i)})^2$$



Multiple variables

- Simple:

$$h_{\theta}(x) = \theta_0 + \theta_1 x_1 + \theta_2 x_2 + \dots + \theta_n x_n$$

$$\theta_0 = \theta_0 - \alpha \frac{\partial}{\partial \theta_0} J(\theta_0, \theta_1, \theta_2, \dots, \theta_n) = \theta_0 - \frac{\alpha}{m} \sum_{i=1}^m (h_{\theta}(x^{(i)}) - y^{(i)})$$

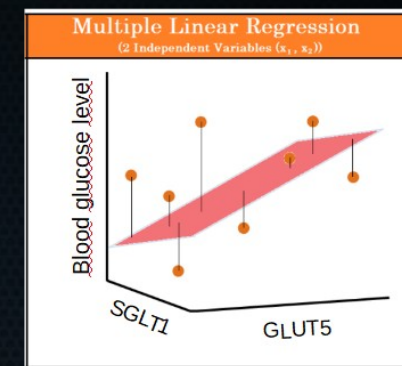
$$\theta_1 = \theta_1 - \alpha \frac{\partial}{\partial \theta_1} J(\theta_0, \theta_1, \theta_2, \dots, \theta_n) = \theta_1 - \frac{\alpha}{m} \sum_{i=1}^m ((h_{\theta}(x^{(i)}) - y^{(i)}) \cdot x_1^{(i)})$$

.

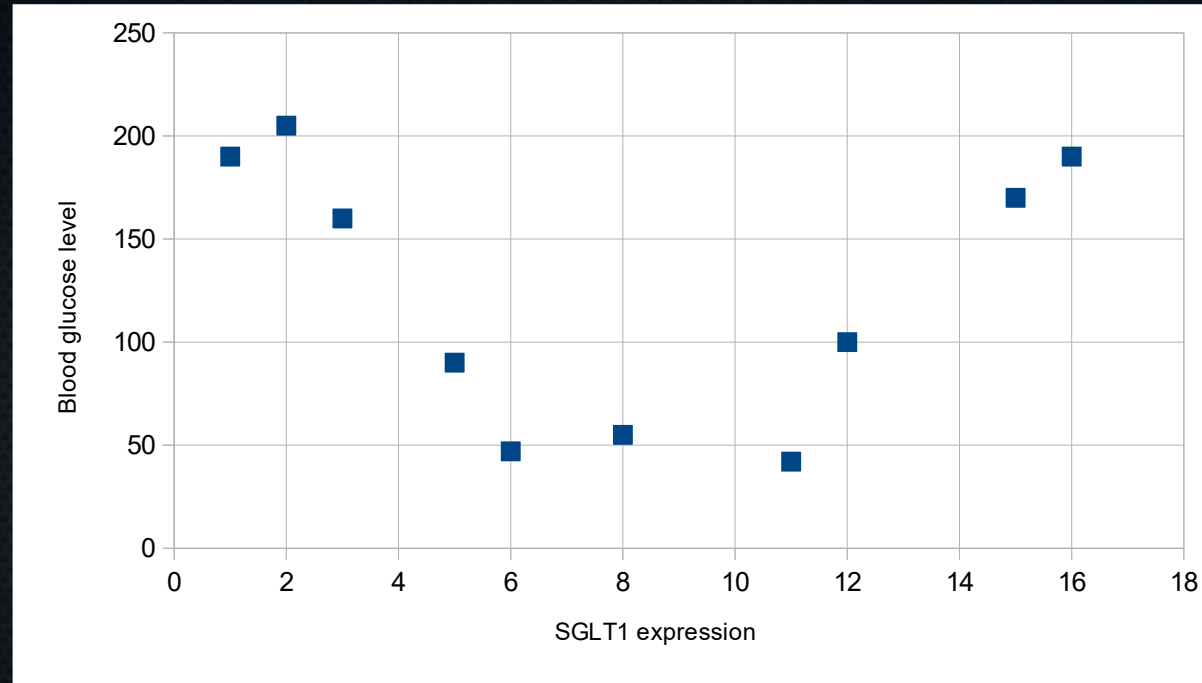
.

.

$$\theta_n = \theta_n - \alpha \frac{\partial}{\partial \theta_n} J(\theta_0, \theta_1, \theta_2, \dots, \theta_n) = \theta_n - \frac{\alpha}{m} \sum_{i=1}^m ((h_{\theta}(x^{(i)}) - y^{(i)}) \cdot x_n^{(i)})$$



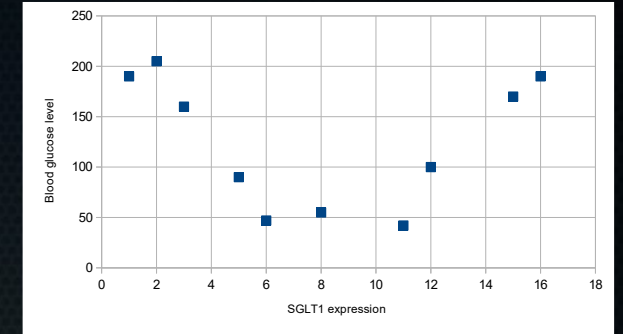
Polynomials/power functions



Polynomials/power functions

- Simple:

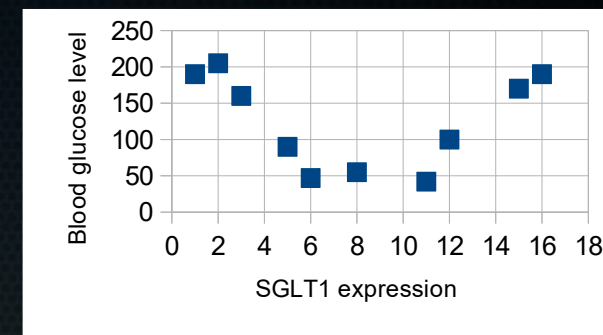
$$h_{\theta}(x) = \theta_0 + \theta_1 x_1 + \theta_2 x_2$$



Polynomials/power functions

- Simple:

$$h_{\theta}(x) = \theta_0 + \theta_1 x_1 + \theta_2 x_2$$

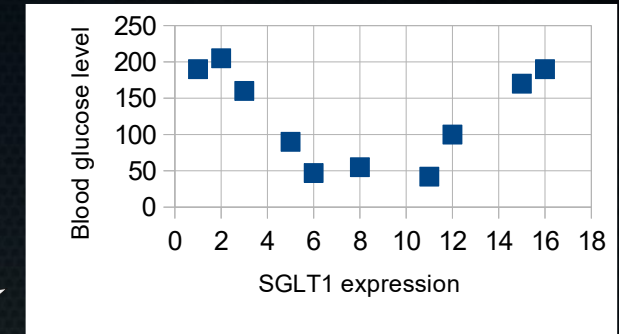


Sample #	SGLT1_linear (x1)	Blood glucose level (mg/dL)
1	3	155
2	8	55
3	12	101
4	2	200

Polynomials/power functions

- Simple:

$$h_{\theta}(x) = \theta_0 + \theta_1 x_1 + \theta_2 x_2$$

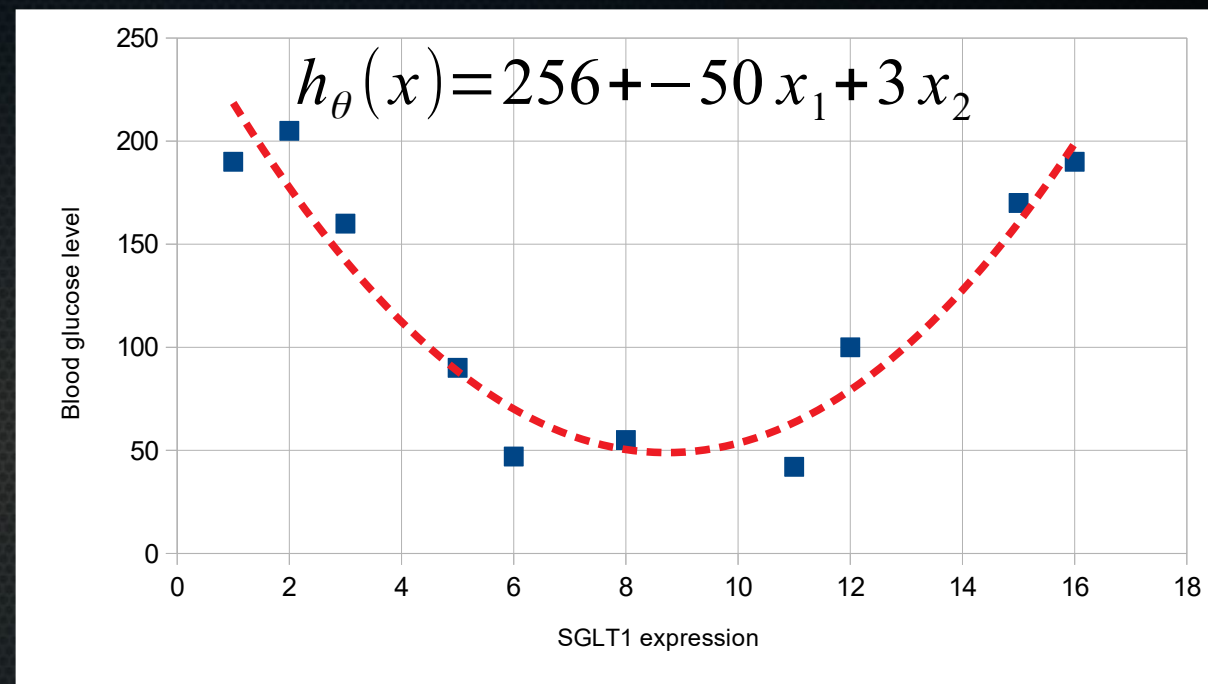


Sample #	SGLT1_linear (x1)	SGLT1_square (x2)	Blood glucose level (mg/dL)
1	3	9	155
2	8	64	55
3	12	144	101
4	2	4	200

Polynomials/power functions

- Simple:

$$h_{\theta}(x) = \theta_0 + \theta_1 x_1 + \theta_2 x_2$$



Sample #	SGLT1_linear (x1)	SGLT1_square (x2)	Blood glucose level (mg/dL)
1	3	9	155
2	8	64	55
3	12	144	101
4	2	4	200

Note on feature scaling

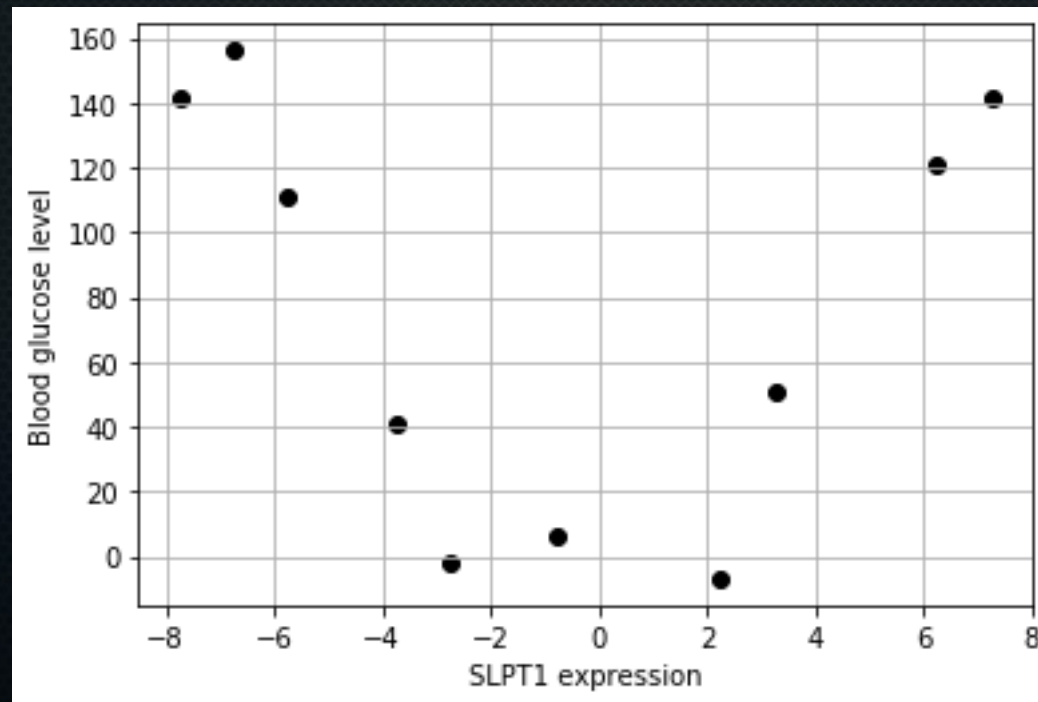
$$h_{\theta}(x) = \theta_0 + \theta_1 x_1 + \theta_2 x_2$$

Sample #	SGLT1_linear (x1)	SGLT1_square (x2)	Blood glucose level (mg/dL)
1	3	9	155
2	8	64	55
3	12	144	101
4	2	4	200

Note on feature scaling

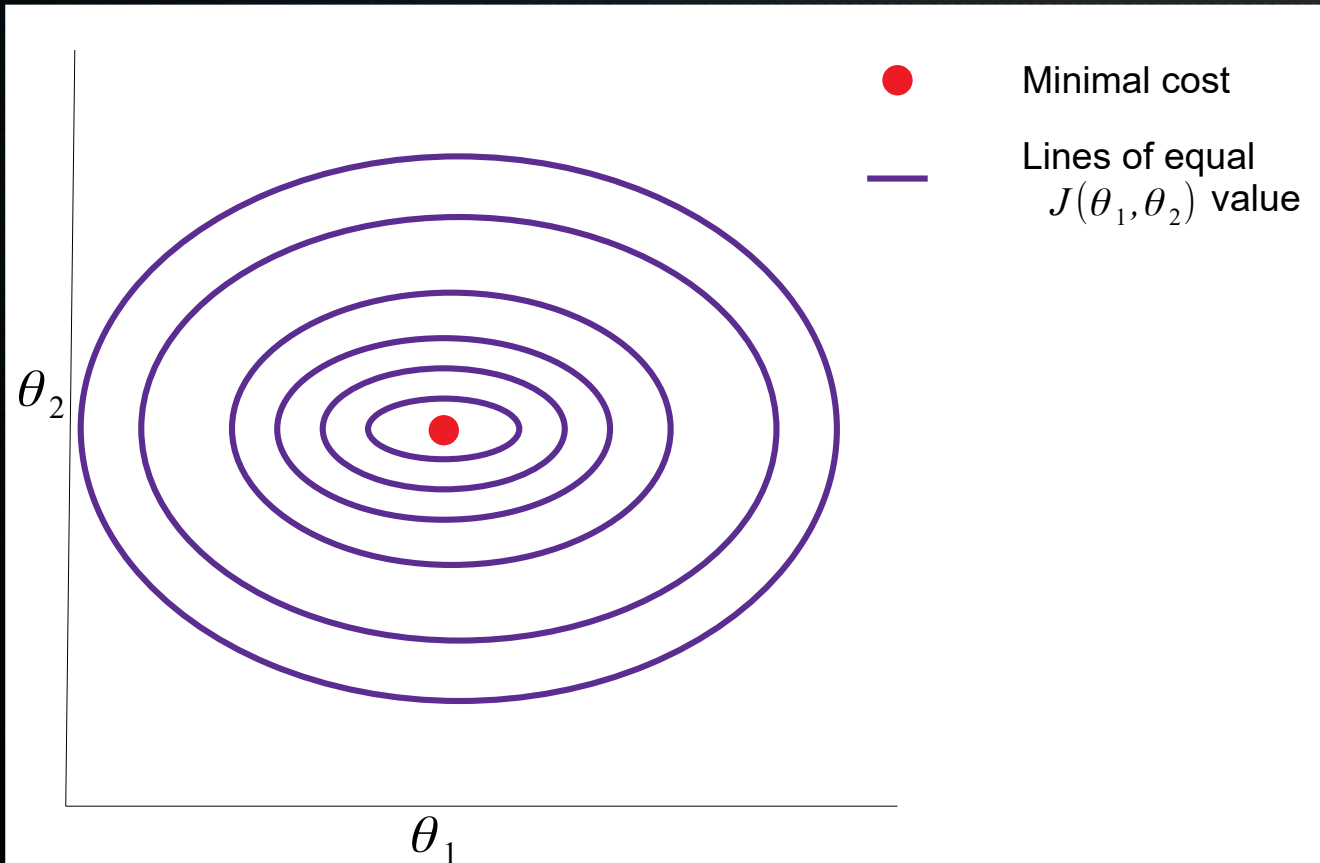
$$h_{\theta}(x) = \cancel{\theta_0} + \theta_1 x_1 + \theta_2 x_2$$

Sample #	SGLT1_linear (x1)	SGLT1_square (x2)	Blood glucose level (mg/dL)
1	3	9	155
2	8	64	55
3	12	144	101
4	2	4	200

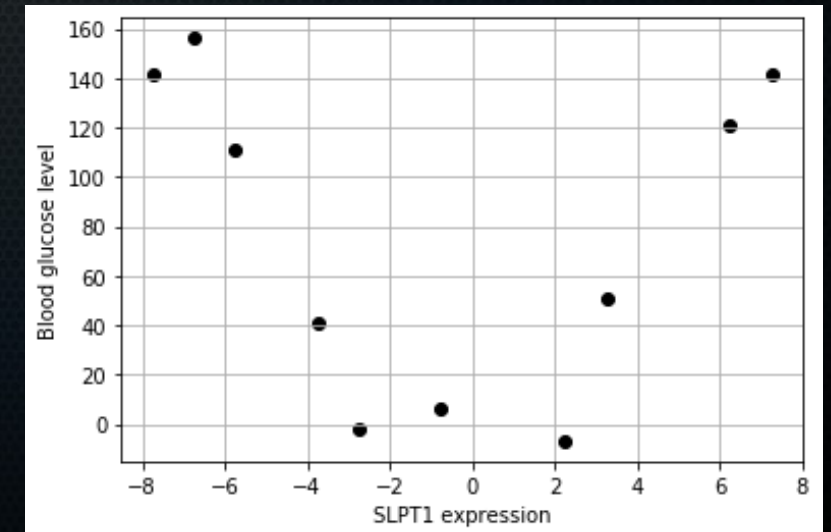


Note on feature scaling

$$h_{\theta}(x) = \cancel{\theta_0} + \theta_1 x_1 + \theta_2 x_2$$

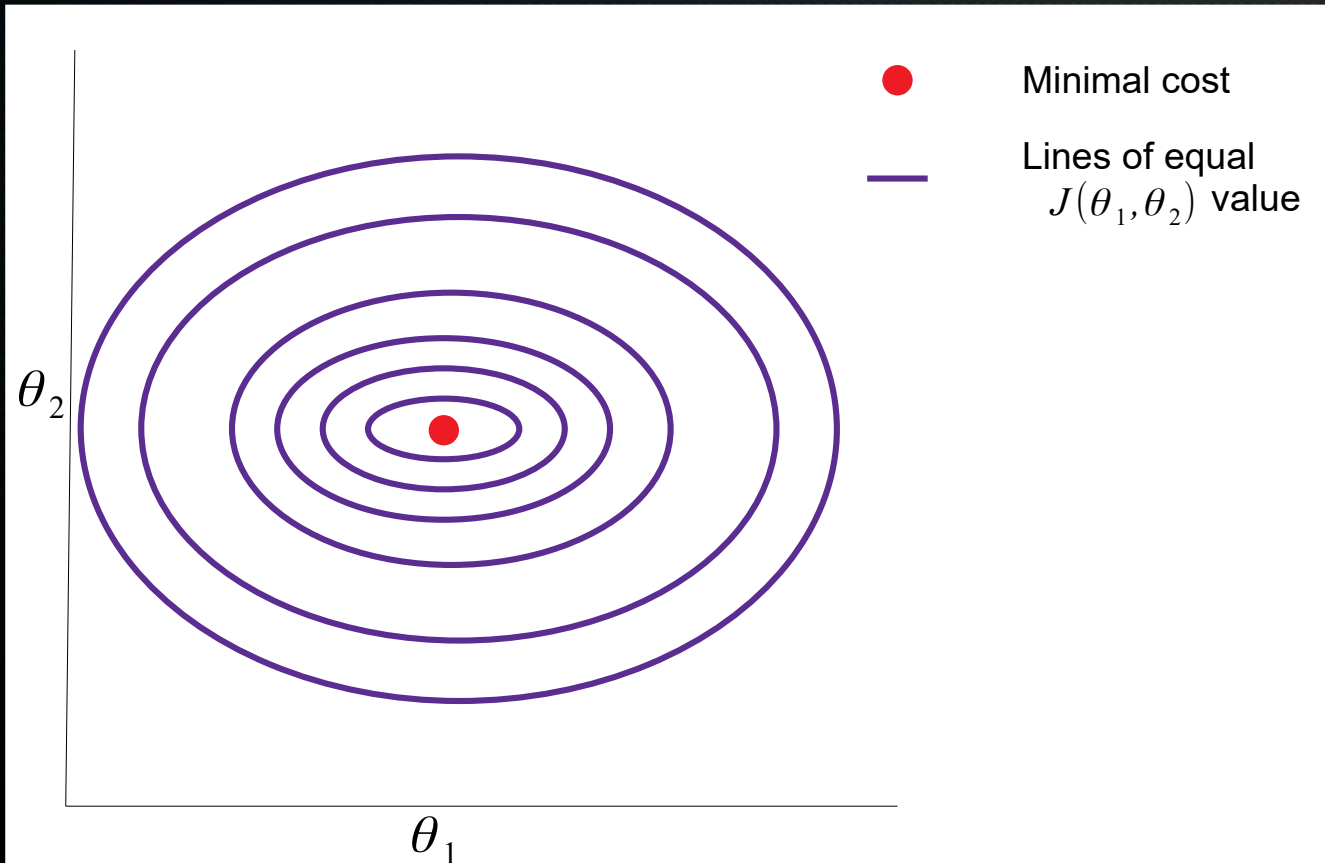


Sample #	SGLT1_linear (x1)	SGLT1_square (x2)	Blood glucose level (mg/dL)
1	3	9	155
2	8	64	55
3	12	144	101
4	2	4	200



Note on feature scaling

$$h_{\theta}(x) = \cancel{\theta_0} + \theta_1 x_1 + \theta_2 x_2$$



Sample #	SGLT1_linear (x1)	SGLT1_square (x2)	Blood glucose level (mg/dL)
1	3	9	155
2	8	64	55
3	12	144	101
4	2	4	200

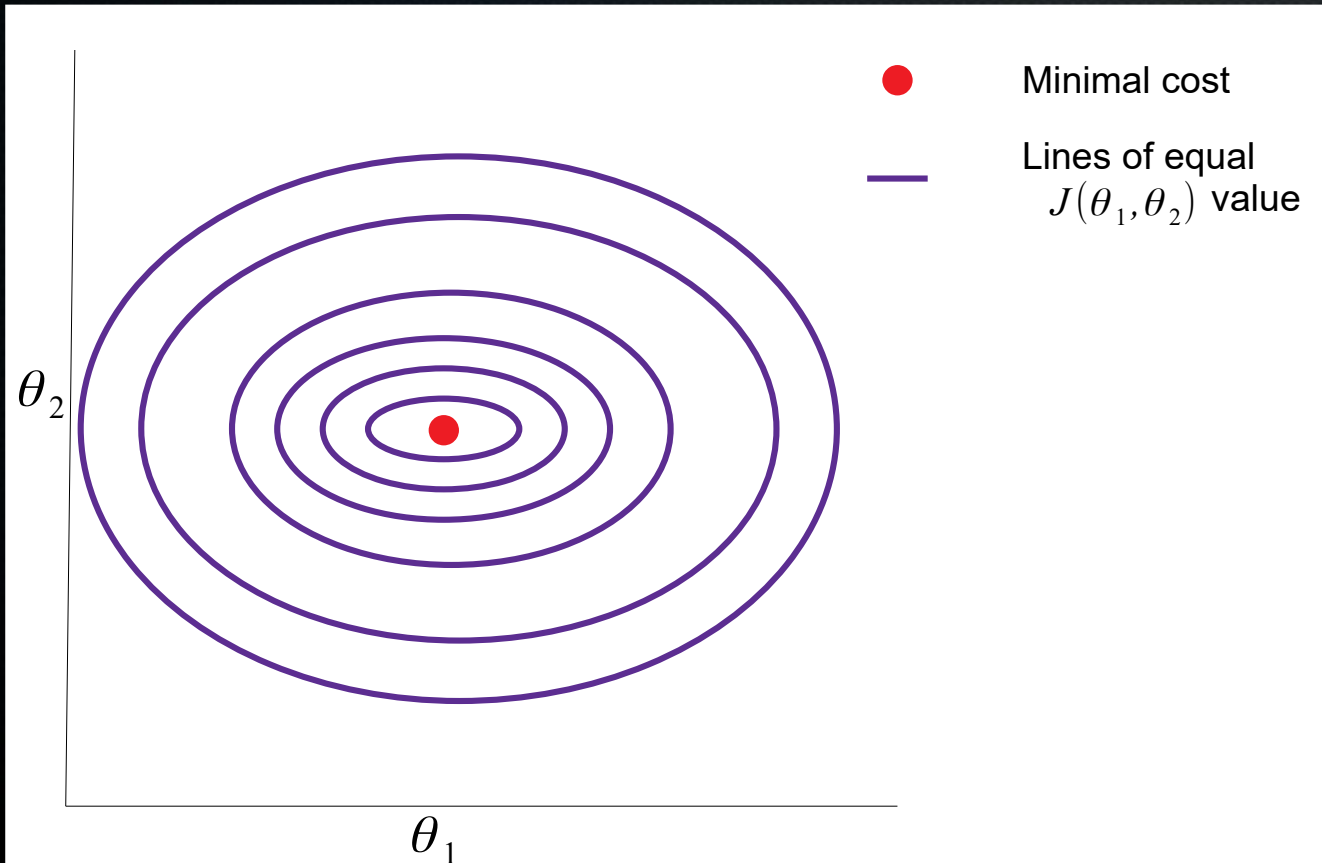
Range: 2-12

Range: 4-144

$$\theta_n = \theta_n - \frac{\alpha}{m} \sum_{i=1}^m ((h_{\theta}(x^{(i)}) - y^{(i)}) \cdot x_n^{(i)})$$

Note on feature scaling

$$h_{\theta}(x) = \cancel{\theta_0} + \theta_1 x_1 + \theta_2 x_2$$



Sample #	SGLT1_linear (x1)	SGLT1_square (x2)	Blood glucose level (mg/dL)
1	3	9	155
2	8	64	55
3	12	144	101
4	2	4	200

Range: 2-12

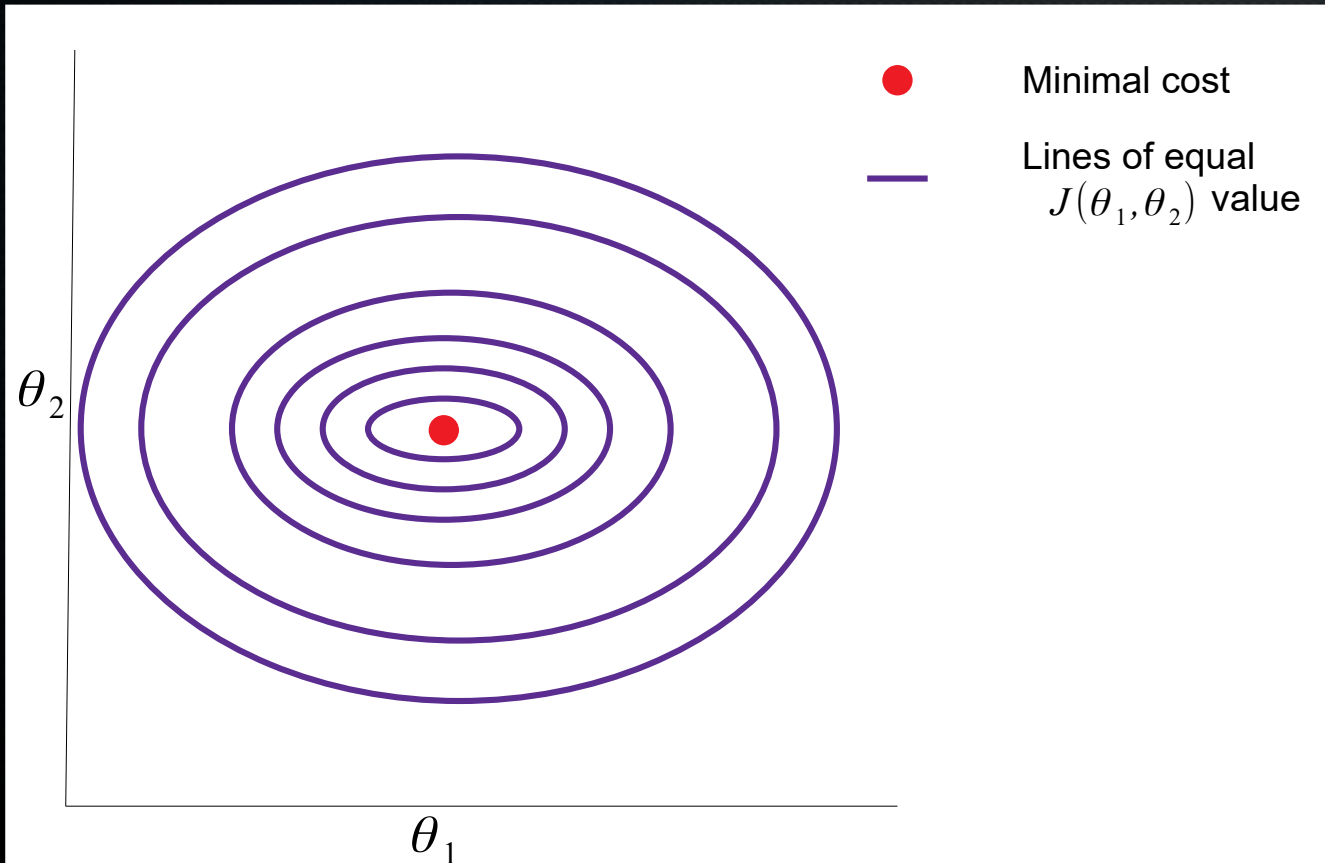
Range: 4-144

$$\theta_n = \theta_n - \frac{\alpha}{m} \sum_{i=1}^m ((h_{\theta}(x^{(i)}) - y^{(i)}) \cdot x_n^{(i)})$$

Small steps

Note on feature scaling

$$h_{\theta}(x) = \cancel{\theta_0} + \theta_1 x_1 + \theta_2 x_2$$



Sample #	SGLT1_linear (x1)	SGLT1_square (x2)	Blood glucose level (mg/dL)
1	3	9	155
2	8	64	55
3	12	144	101
4	2	4	200

Range: 2-12

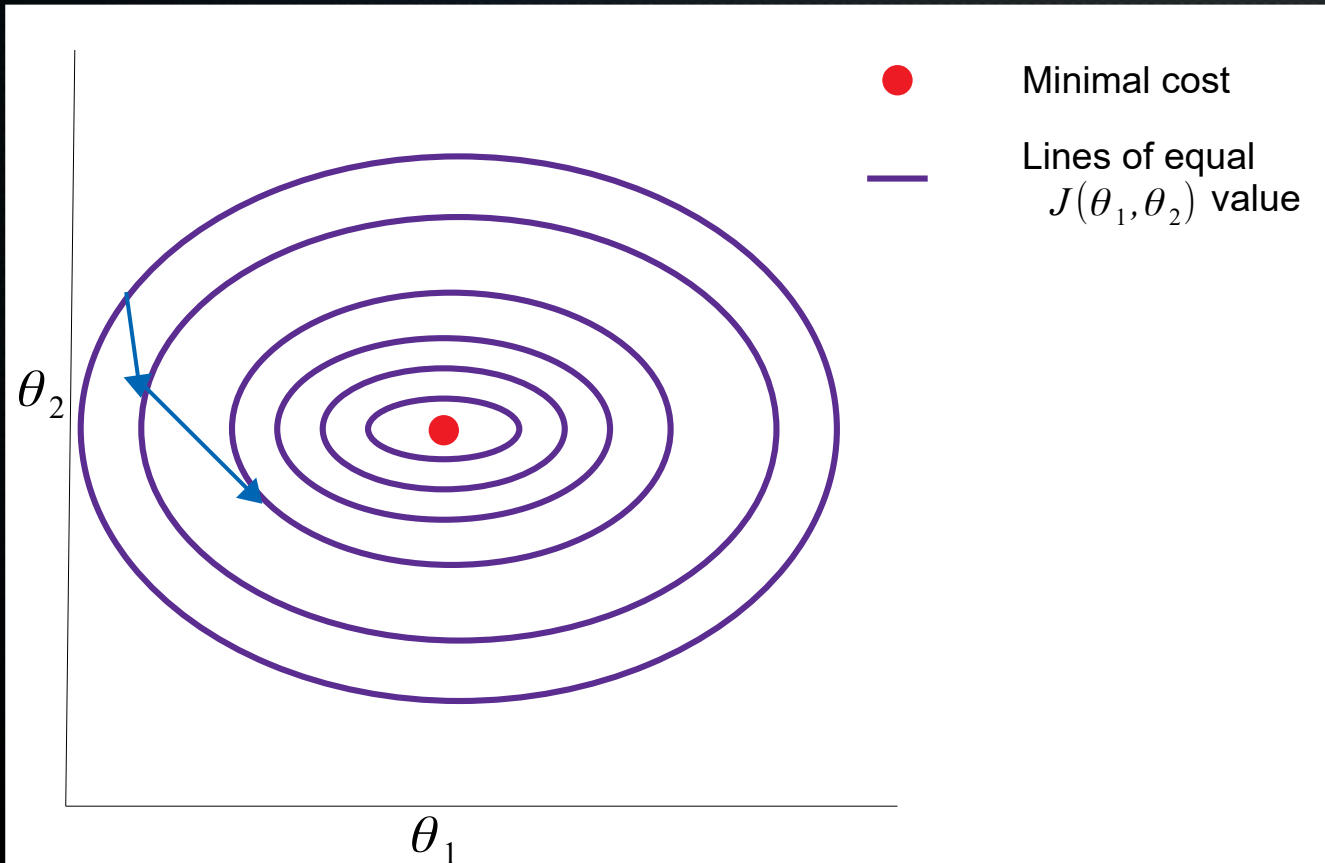
Range: 4-144

$$\theta_n = \theta_n - \frac{\alpha}{m} \sum_{i=1}^m ((h_{\theta}(x^{(i)}) - y^{(i)}) \cdot x_n^{(i)})$$

Large steps

Note on feature scaling

$$h_{\theta}(x) = \cancel{\theta_0} + \theta_1 x_1 + \theta_2 x_2$$



Sample #	SGLT1_linear (x1)	SGLT1_square (x2)	Blood glucose level (mg/dL)
1	3	9	155
2	8	64	55
3	12	144	101
4	2	4	200

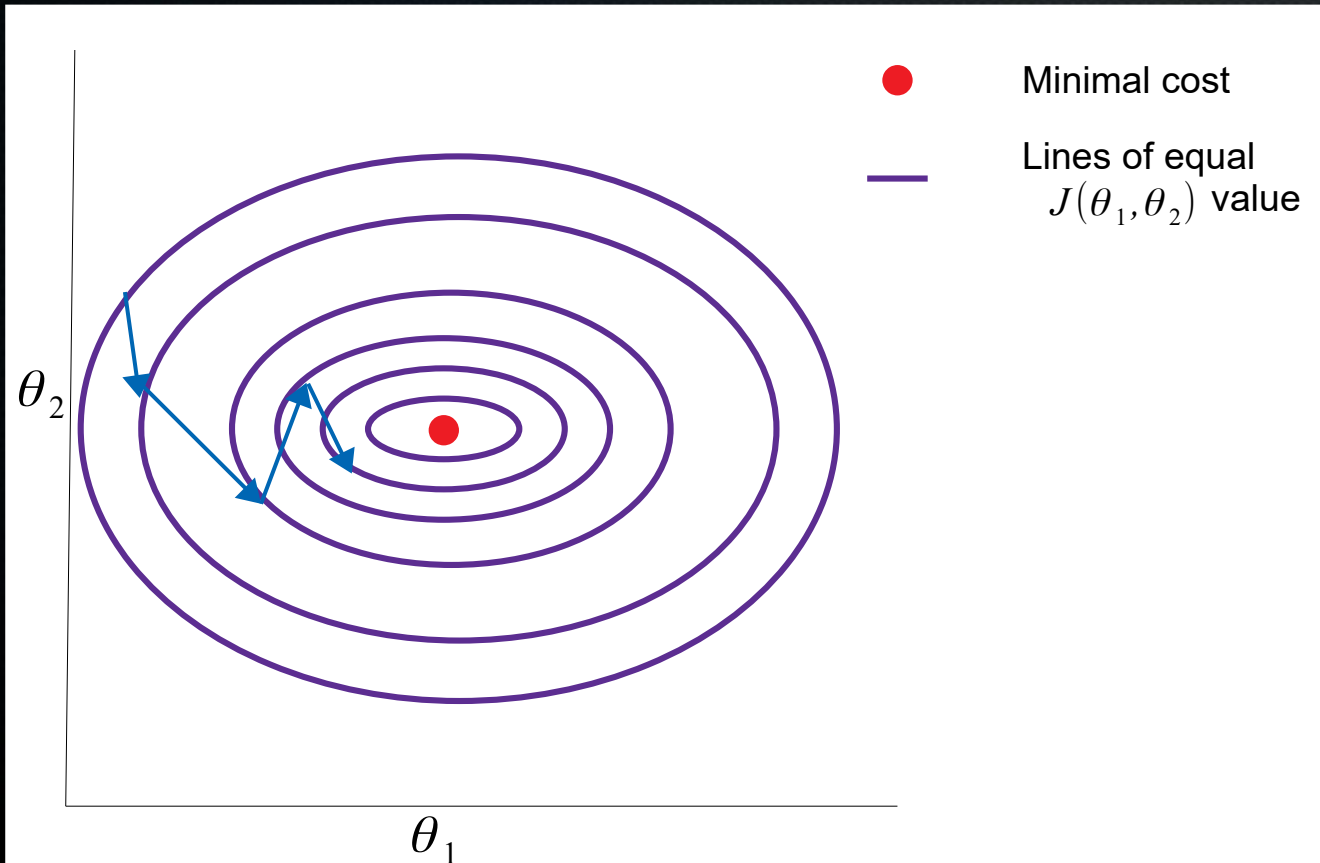
Range: 2-12

Range: 4-144

$$\theta_n = \theta_n - \frac{\alpha}{m} \sum_{i=1}^m ((h_{\theta}(x^{(i)}) - y^{(i)}) \cdot x_n^{(i)})$$

Note on feature scaling

$$h_{\theta}(x) = \cancel{\theta_0} + \theta_1 x_1 + \theta_2 x_2$$



Sample #	SGLT1_linear (x1)	SGLT1_square (x2)	Blood glucose level (mg/dL)
1	3	9	155
2	8	64	55
3	12	144	101
4	2	4	200

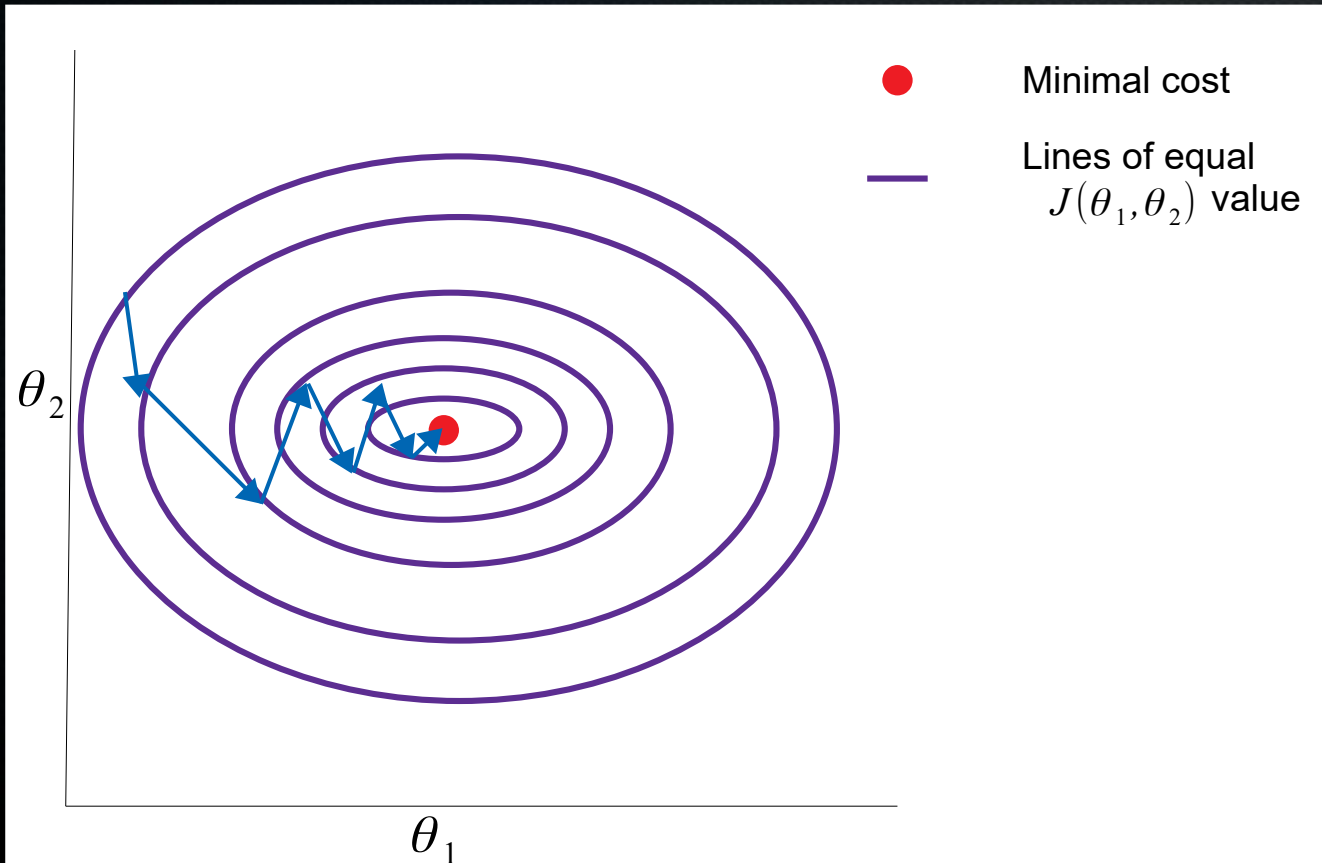
Range: 2-12

Range: 4-144

$$\theta_n = \theta_n - \frac{\alpha}{m} \sum_{i=1}^m ((h_{\theta}(x^{(i)}) - y^{(i)}) \cdot x_n^{(i)})$$

Note on feature scaling

$$h_{\theta}(x) = \cancel{\theta_0} + \theta_1 x_1 + \theta_2 x_2$$



Sample #	SGLT1_linear (x1)	SGLT1_square (x2)	Blood glucose level (mg/dL)
1	3	9	155
2	8	64	55
3	12	144	101
4	2	4	200

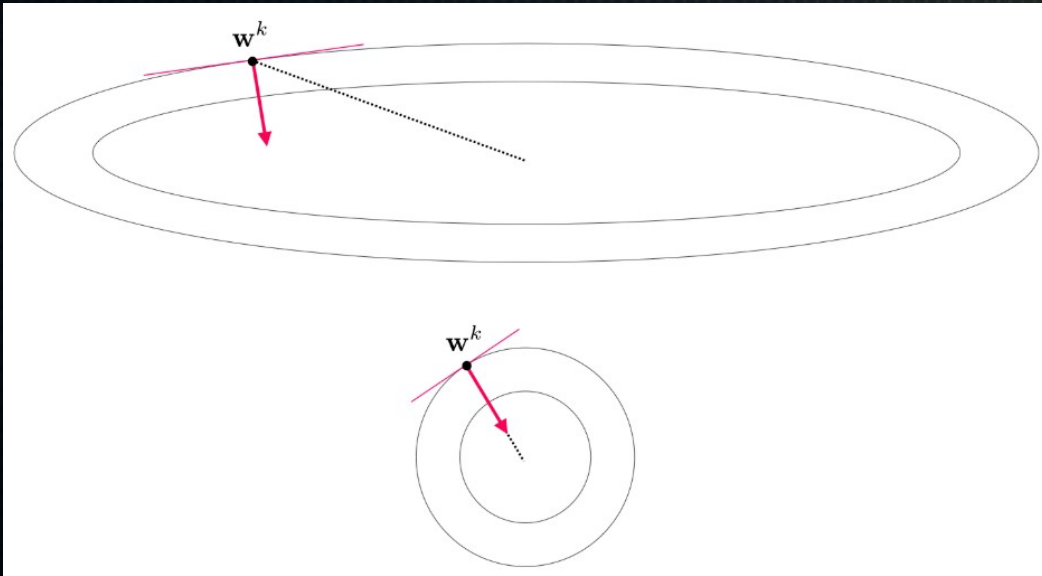
Range: 2-12

Range: 4-144

$$\theta_n = \theta_n - \frac{\alpha}{m} \sum_{i=1}^m ((h_{\theta}(x^{(i)}) - y^{(i)}) \cdot x_n^{(i)})$$

Note on feature scaling

$$h_{\theta}(x) = \cancel{\theta_0} + \theta_1 x_1 + \theta_2 x_2$$



Sample #	SGLT1_linear (x1)	SGLT1_square (x2)	Blood glucose level (mg/dL)
1	3	9	155
2	8	64	55
3	12	144	101
4	2	4	200

Range: 2-12

Range: 4-144

Note on feature scaling

$$h_{\theta}(x) = \cancel{\theta_0} + \theta_1 x_1 + \theta_2 x_2$$

$$x_j = \frac{x_j - \text{mean}(x_j)}{\text{std.dev}(x_j)}$$

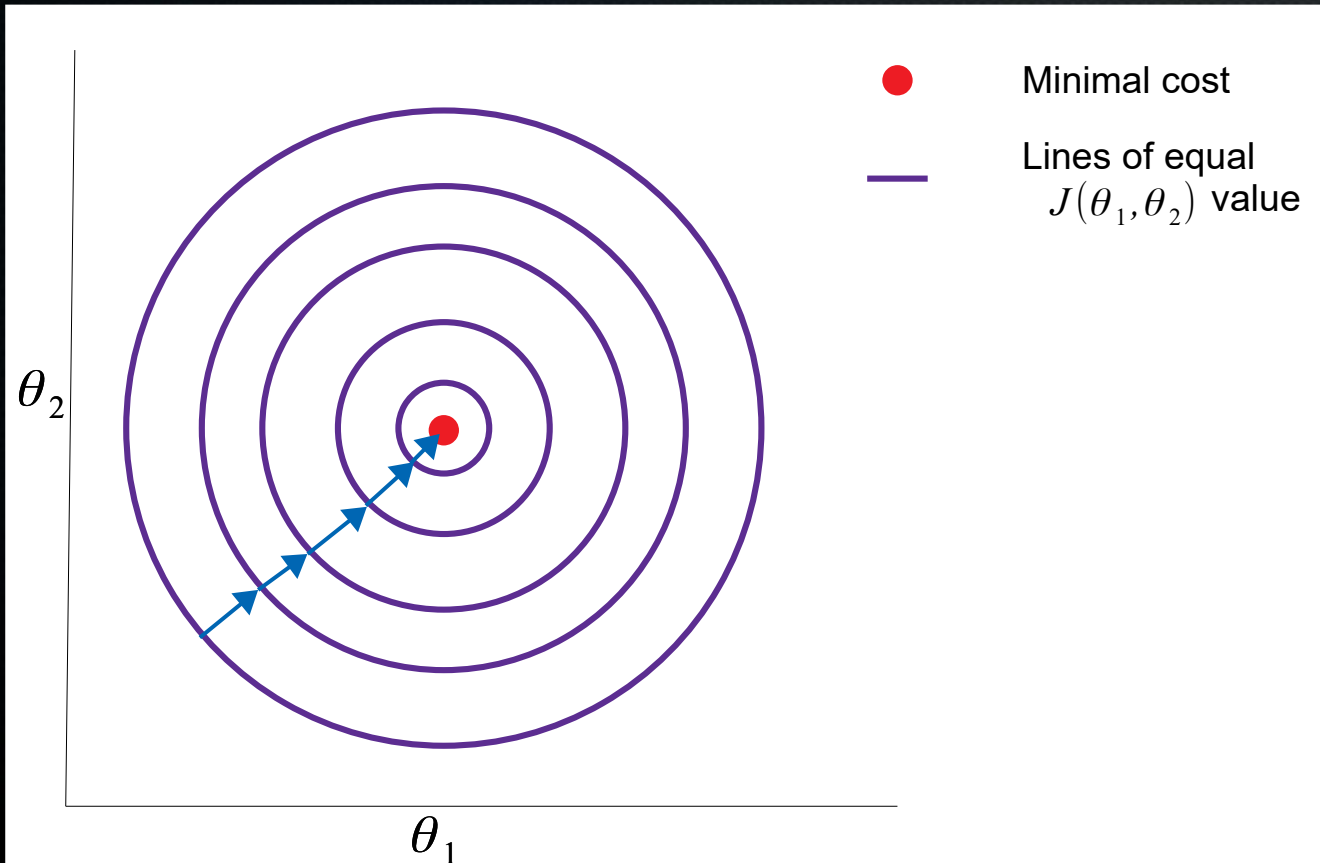
Sample #	SGLT1_linear (x1)	SGLT1_square (x2)	Blood glucose level (mg/dL)
1	-0,70	-0,71	155
2	0,38	0,13	55
3	1,24	1,36	101
4	-0,91	-0,79	200

Note on feature scaling

$$h_{\theta}(x) = \cancel{\theta_0} + \theta_1 x_1 + \theta_2 x_2$$

$$x_j = \frac{x_j - \text{mean}(x_j)}{\text{std.dev}(x_j)}$$

Sample #	SGLT1_linear (x1)	SGLT1_square (x2)	Blood glucose level (mg/dL)
1	-0,70	-0,71	155
2	0,38	0,13	55
3	1,24	1,36	101
4	-0,91	-0,79	200

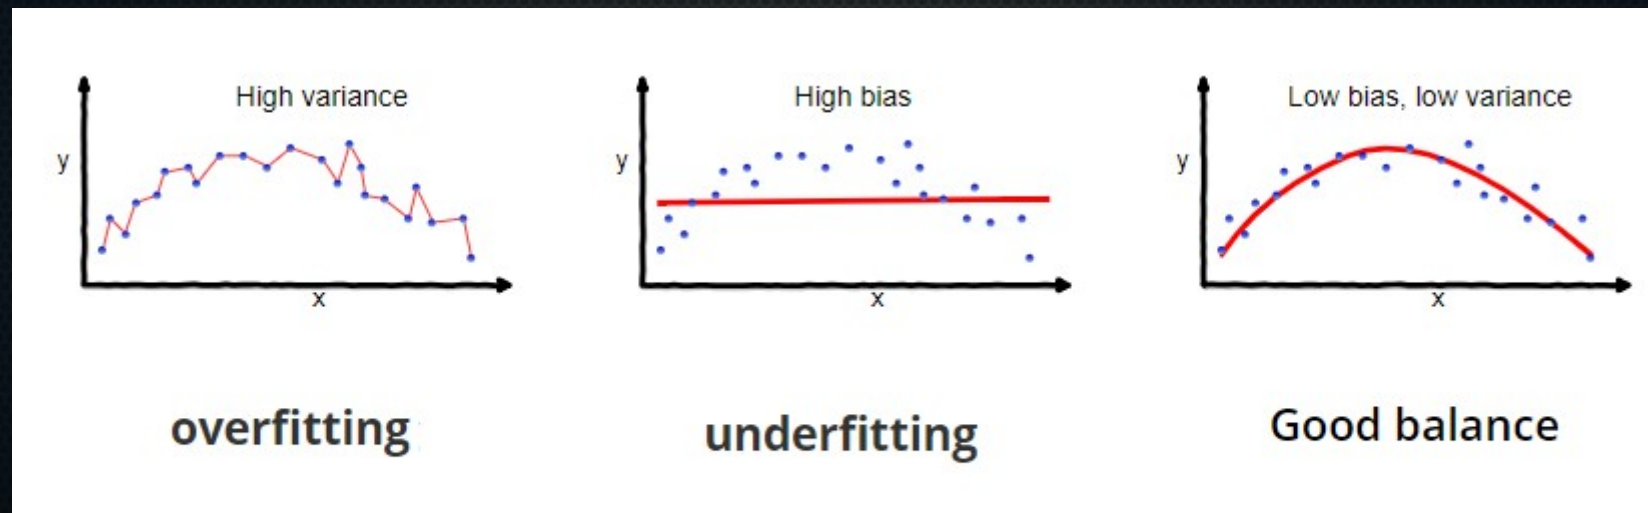


Summary

- Easily extendable to multiple features and polynomials
- Normalisation to help gradient descent converge faster

How to generalise well

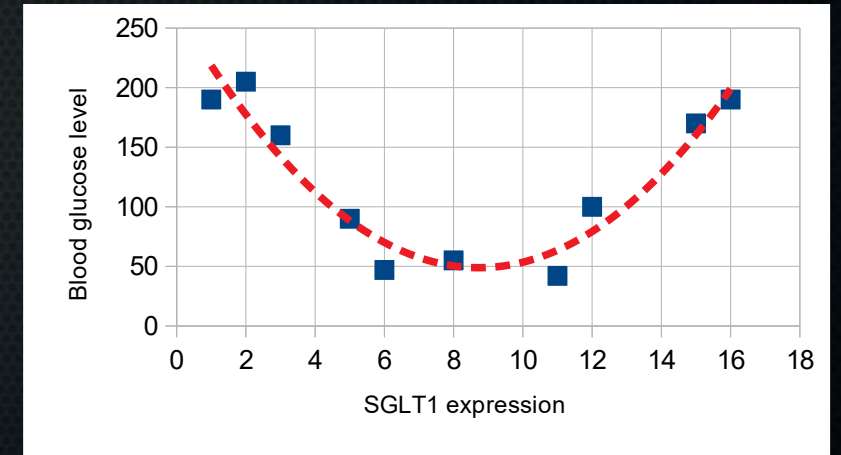
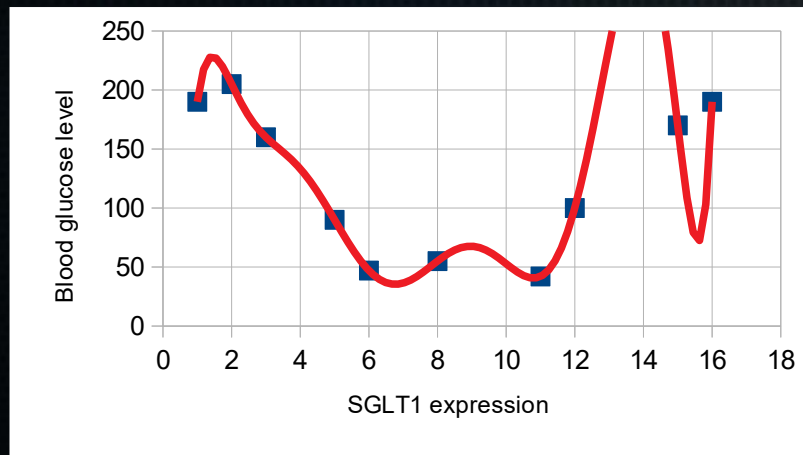
- Goal is not to fit the training data perfectly, but to generalise well



Source:
<https://towardsdatascience.com/understanding-the-bias-variance-tradeoff-165e6942b229>

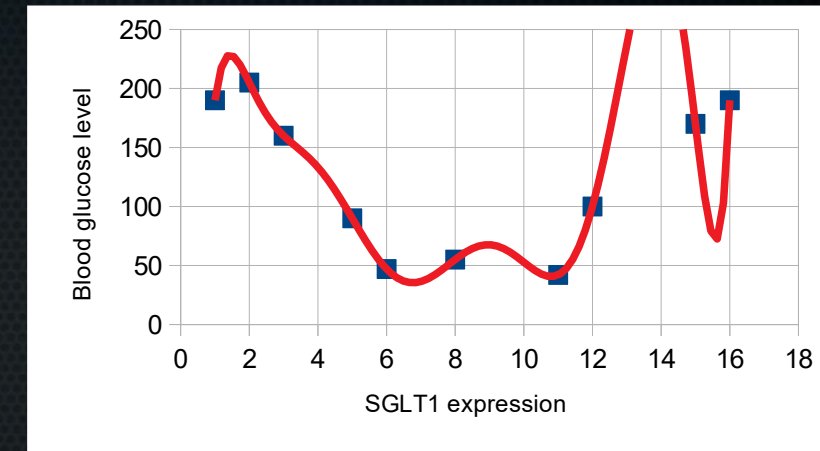
How to generalise well

- Goal is not to fit the training data perfectly, but to generalise well



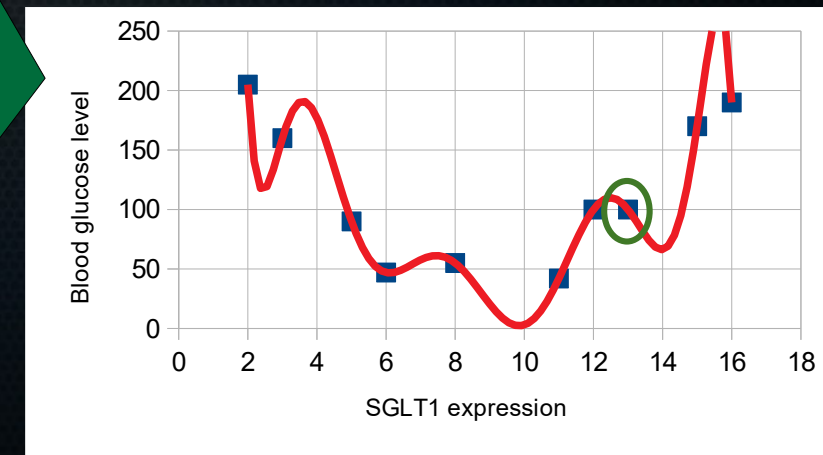
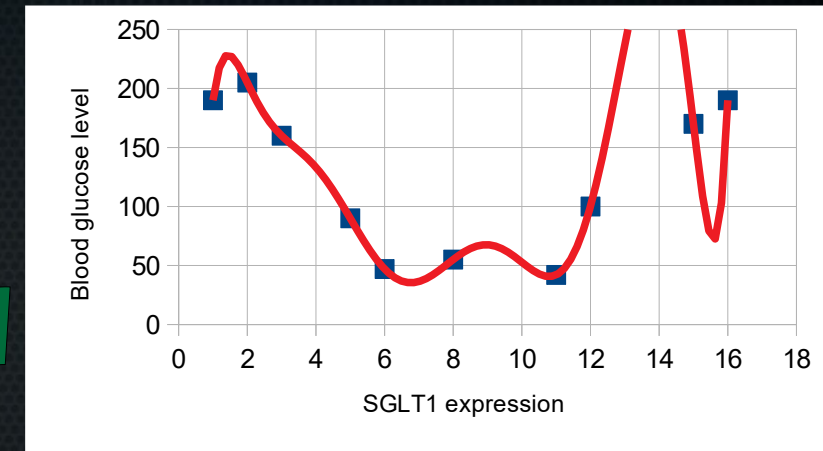
How to generalise well

- **High variance:**
 - how much our hypothesis function would change if we changed our training set
 - used x^1 - x^9 as features



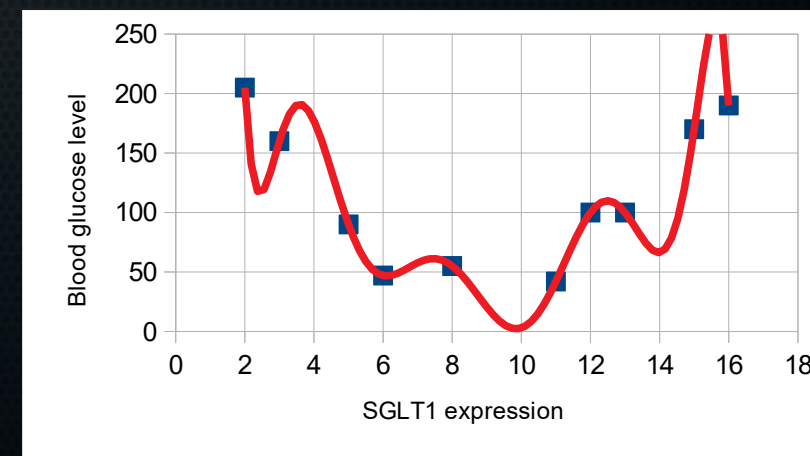
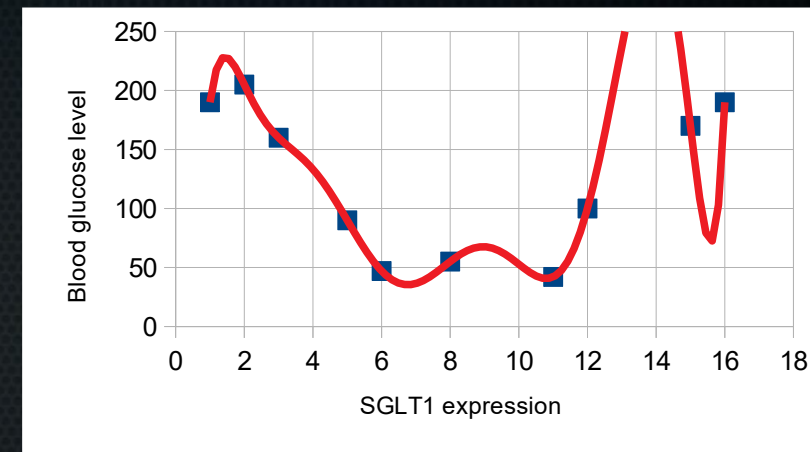
How to generalise well

- **High variance:**
 - how much our hypothesis function would change if we changed our training set
 - used x^1 - x^9 as features
 - If we add one point:



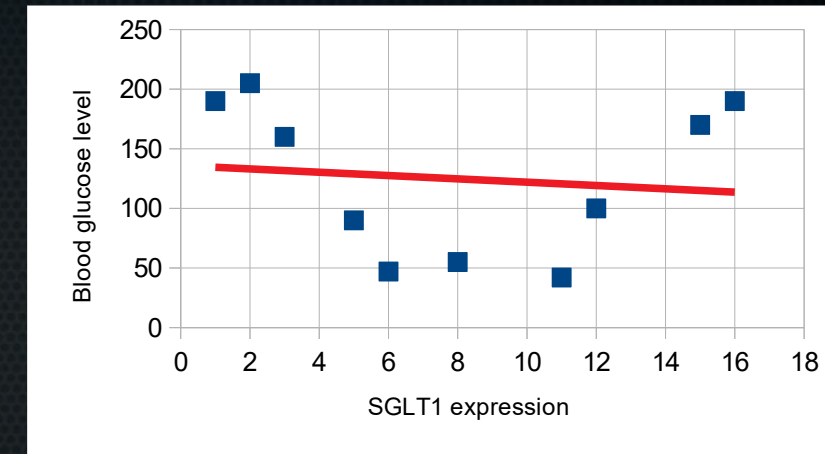
How to generalise well

- **High variance:**
 - Huge amount of possible functions that pass through all points: large hypothesis space, can basically fit all the training data we give perfectly!
 - Do great on this data, but will fare poorly when predicting unseen data



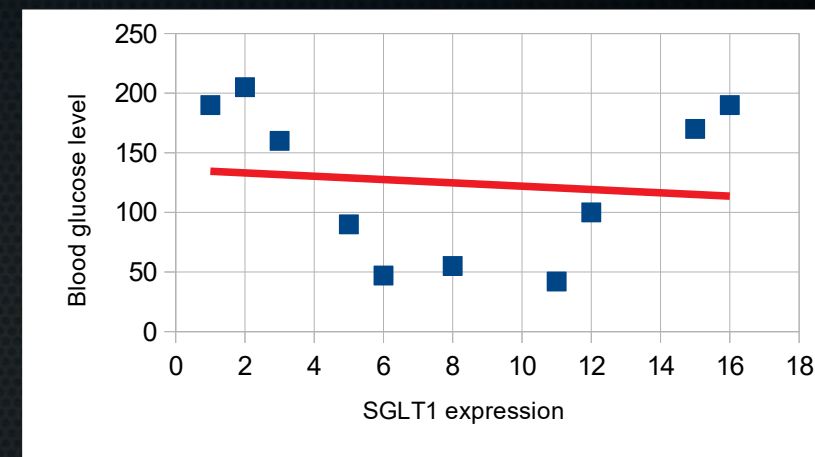
How to generalise well

- **High bias:**
 - error introduced by approximating a complex process with a simple model.
 - Only 1 feature (intercept + θ_1)



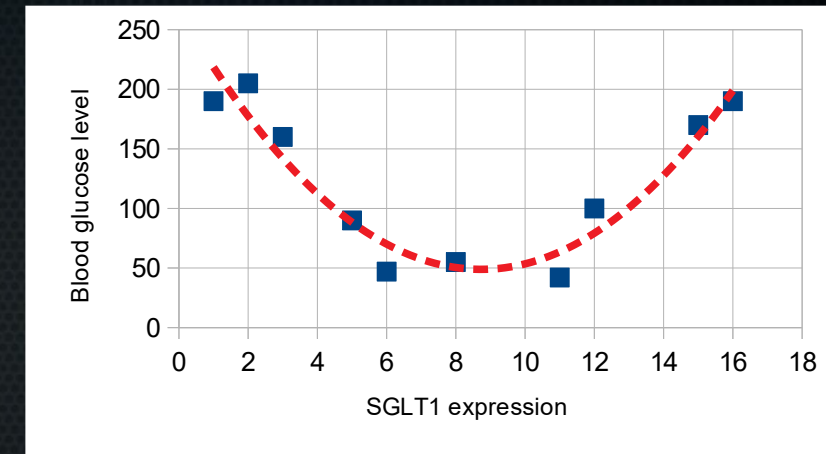
How to generalise well

- High **bias**:
 - error introduced by approximating a complex process with a simple model.
 - Only 1 feature (intercept + θ_1)
 - Data clearly shows that a linear curve is not the best fit, yet we keep our preconception, or *bias*, that it should adhere to a univariate linear regression
 - Does poorly on this data and will also fare poorly when predicting unseen data



How to generalise well

- Just right:
 - Not fit too closely to known examples, probably generalises well.



What can we do to find a good model?

- Find a way to approximate generalisation error: how well do you do on unseen data?
- See how error on seen and unseen data changes with amount of training data (plot learning curves)
- Reduce dimensionality by using only certain features
- Automatically constrain the fitting by penalising the cost function for too many/too large parameters

What can we do to find a good model?

- *Find a way to approximate generalisation error: how well do you do on unseen data?*
- See how error on seen and unseen data changes with amount of training data (plot learning curves)
- Automatically constrain the fitting by penalising the cost function for too many/too large parameters

Approximate generalisation error: split data

- Split data into training data, validation data, and a test set
- Test set: completely untouched until you are done training
- Train set: train your classifier on this
- Validation set: test your trained classifier on this



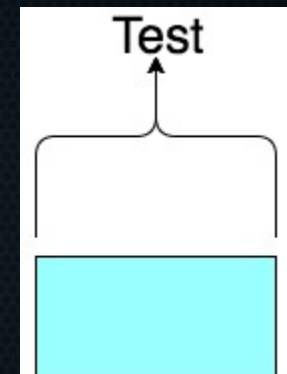
Source: <https://stackoverflow.com/questions/56099495/what-should-be-passed-as-input-parameter-when-using-train-test-split-function-tw>

Approximate generalisation error: split data

- K-fold cross-validation (often 10-fold):



Source: <https://www.statology.org/k-fold-cross-validation/>

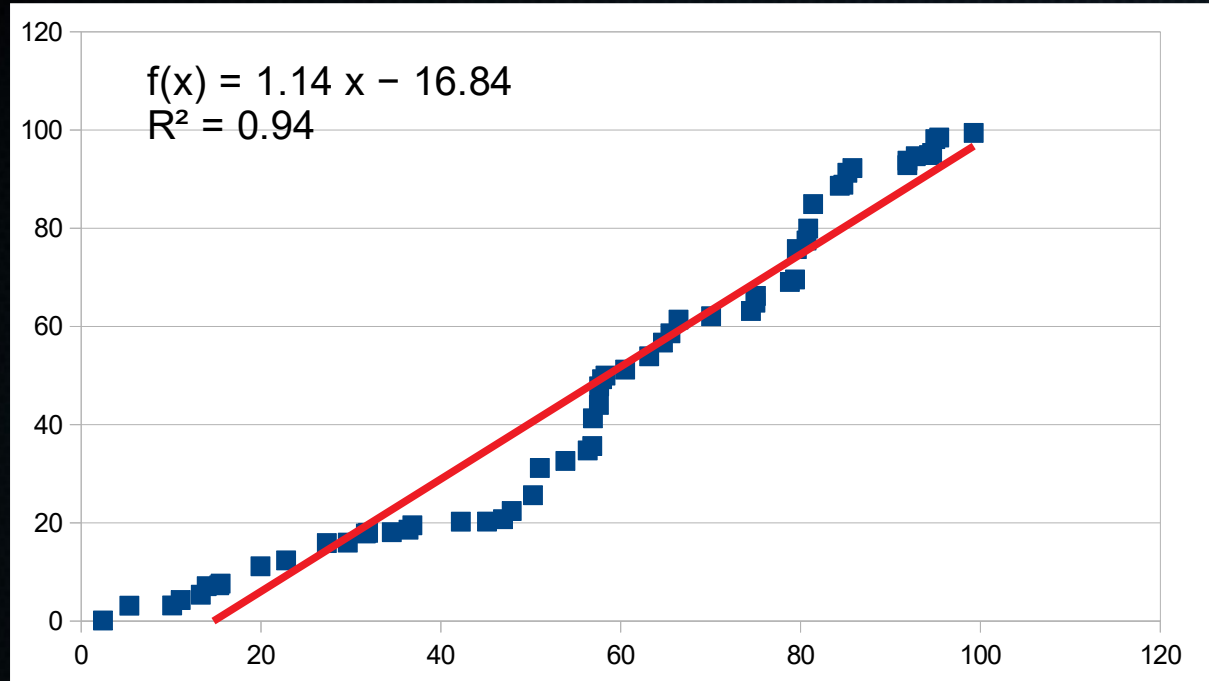


Approximate generalisation error: split data

- Procedure:
 - Shuffle the data
 - Divide into k folds (e.g. 10 folds of 100 training examples each)
 - For each fold:
 - find parameters by minimising cost function on training set with gradient descent
 - predict on validation data
 - calculate cost on validation data (you know the true values)
 - Average validation cost over folds \sim generalisation error

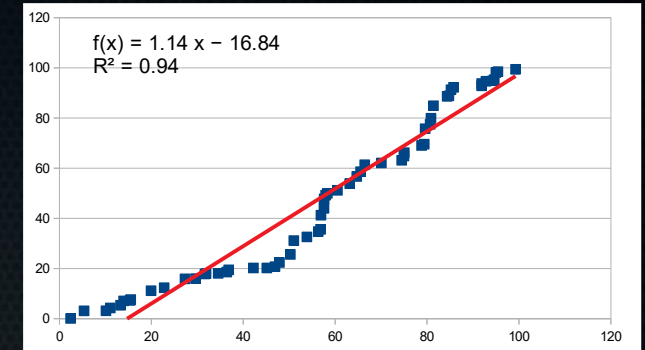
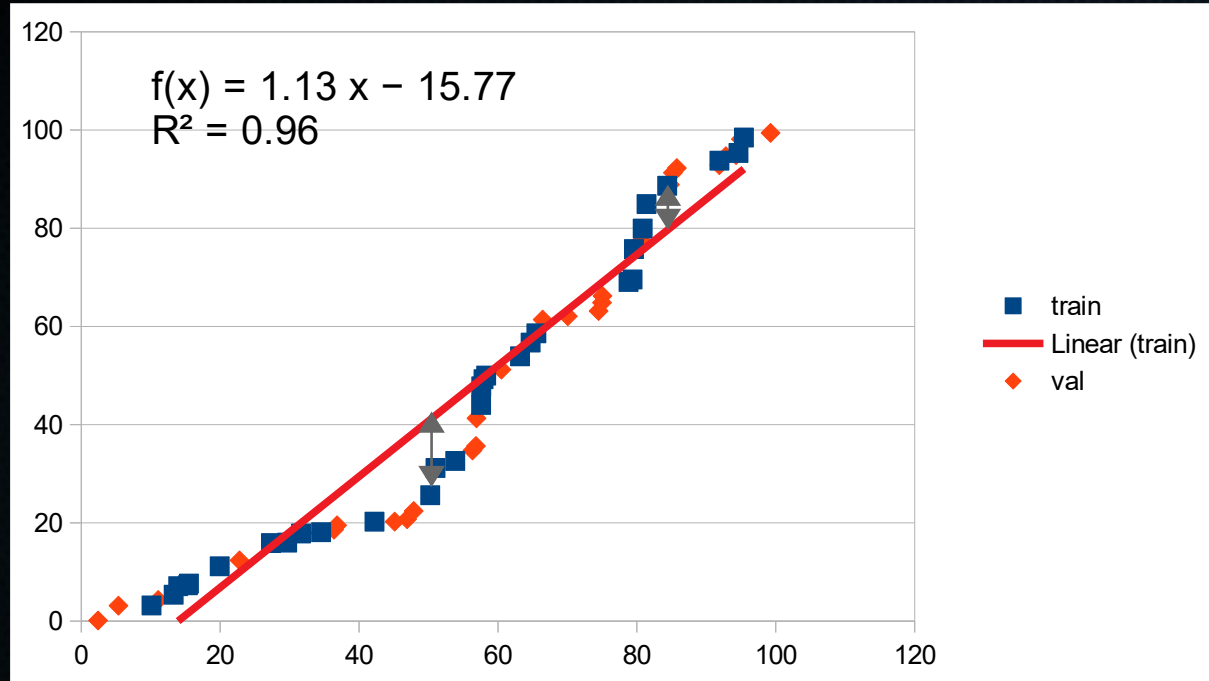


Example 2-fold cv on linear regression



All data

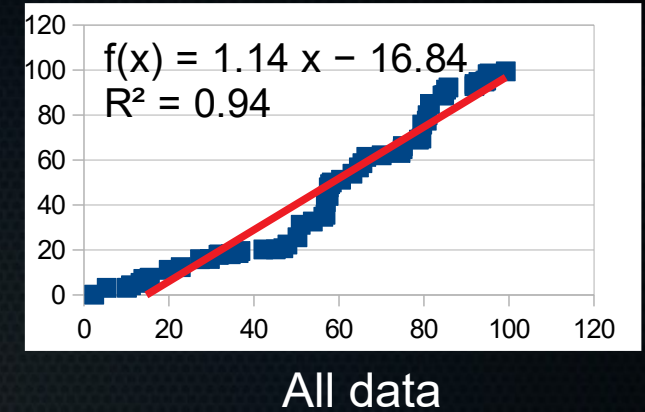
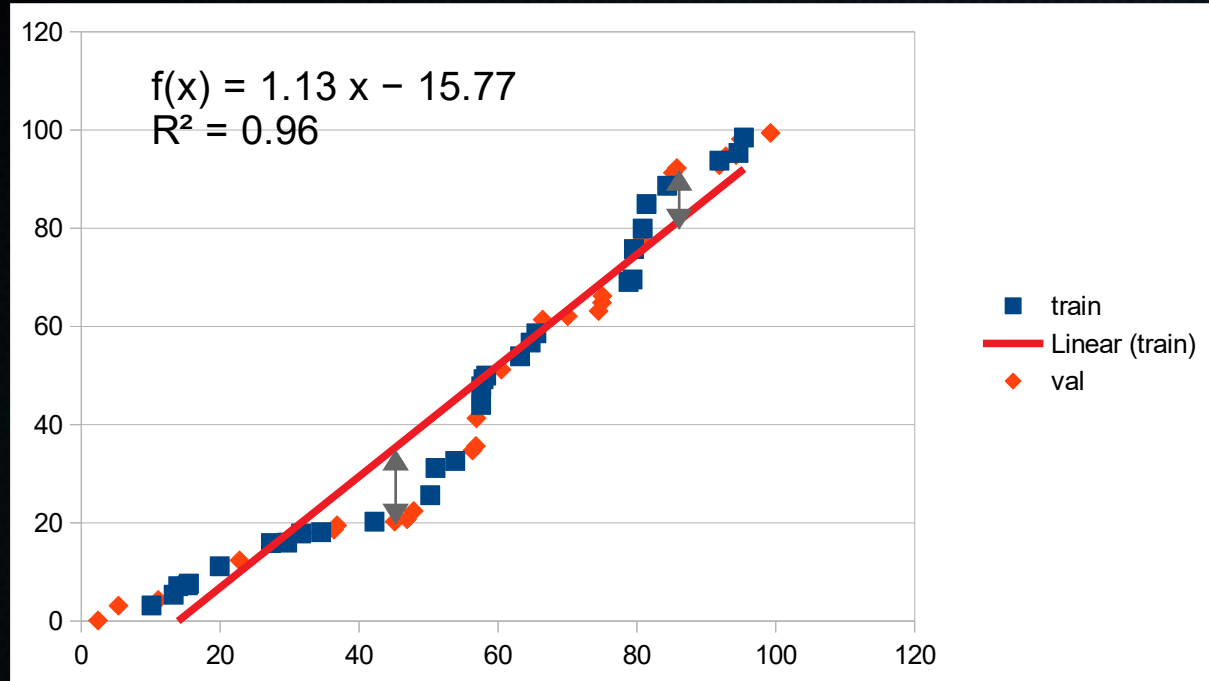
Example 2-fold cv on linear regression: fold 1



All data

$$J(\theta_0, \theta_1)_{train} = \frac{1}{2m_{train}} \sum_{i=1}^{m_{train}} (h_{\theta}(x^{(i)}) - y^{(i)})^2 = 41,66$$

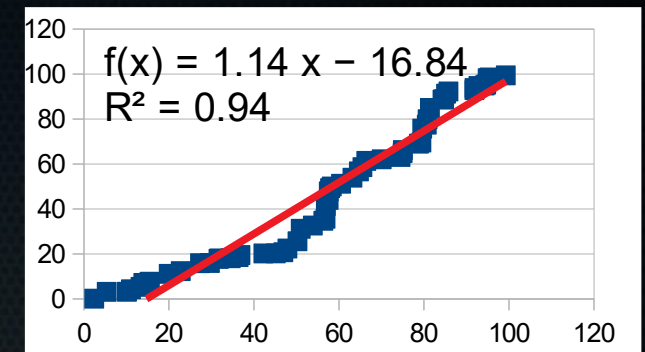
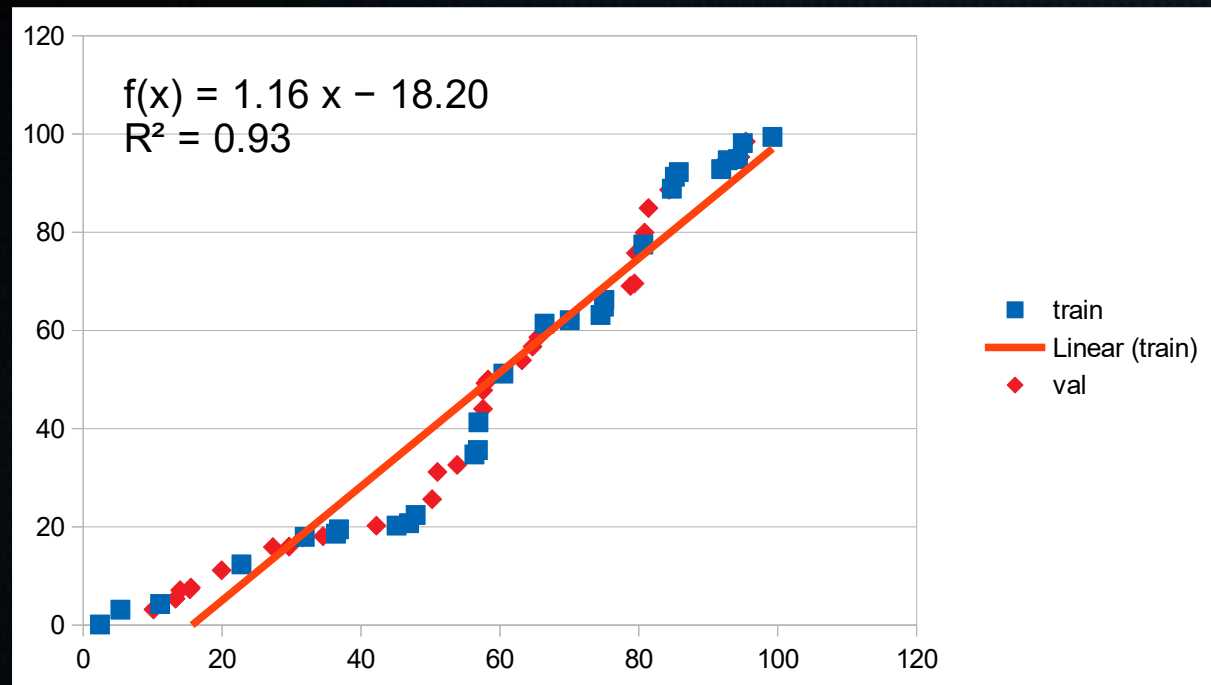
Example 2-fold cv on linear regression: fold 1



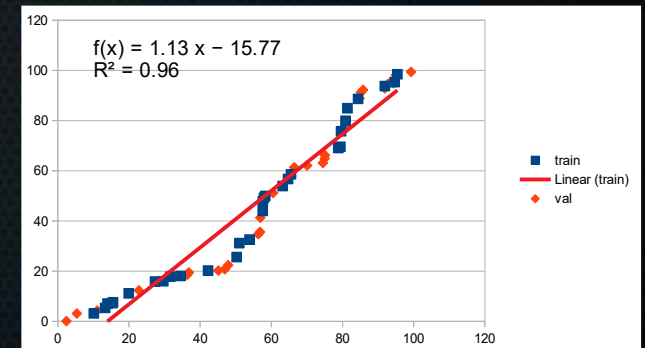
$$J(\theta_0, \theta_1)_{train} = \frac{1}{2m_{train}} \sum_{i=1}^{m_{train}} (h_{\theta}(x^{(i)}) - y^{(i)})^2 = 41,66$$

$$J(\theta_0, \theta_1)_{val} = \frac{1}{2m_{val}} \sum_{i=1}^{m_{val}} (h_{\theta}(x^{(i)}) - y^{(i)})^2 = 52,34$$

Example 2-fold cv on linear regression: fold 2



All data



$$J(\theta_0, \theta_1)_{train} = 41,66$$

$$J(\theta_0, \theta_1)_{val} = 52,34$$

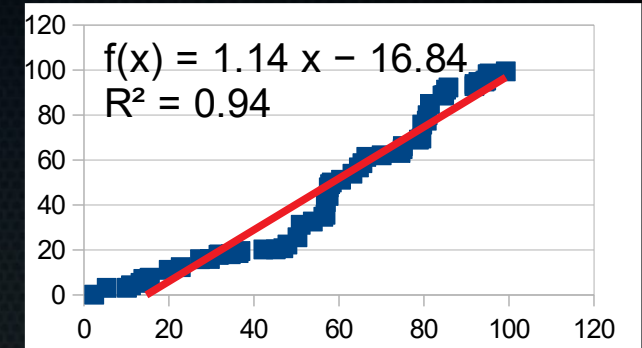
Fold 1

$$J(\theta_0, \theta_1)_{train} = \frac{1}{2m_{train}} \sum_{i=1}^{m_{train}} (h_{\theta}(x^{(i)}) - y^{(i)})^2 = 75,25$$

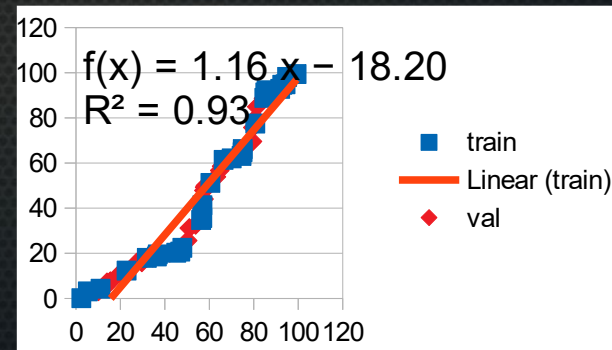
$$J(\theta_0, \theta_1)_{val} = \frac{1}{2m_{val}} \sum_{i=1}^{m_{val}} (h_{\theta}(x^{(i)}) - y^{(i)})^2 = 43$$

Example 2-fold cv on linear regression

- Avg. train error: 58,5
- Avg. Validation error: 47,7
- Perform better on unseen than seen data: generalises well.
- Finally: would train on all data and test that on test set.

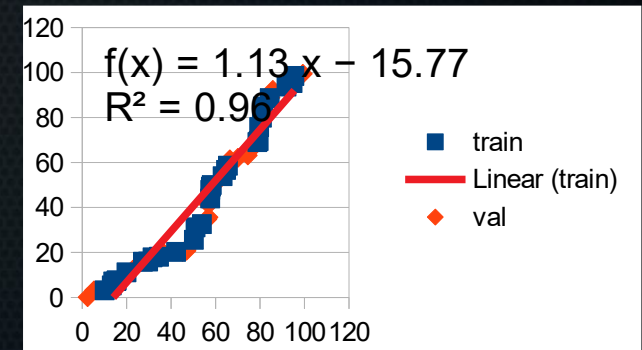


All data



$$J(\theta_0, \theta_1)_{train} = 75,25$$
$$J(\theta_0, \theta_1)_{val} = 43$$

Fold 2



$$J(\theta_0, \theta_1)_{train} = 41,66$$
$$J(\theta_0, \theta_1)_{val} = 52,34$$

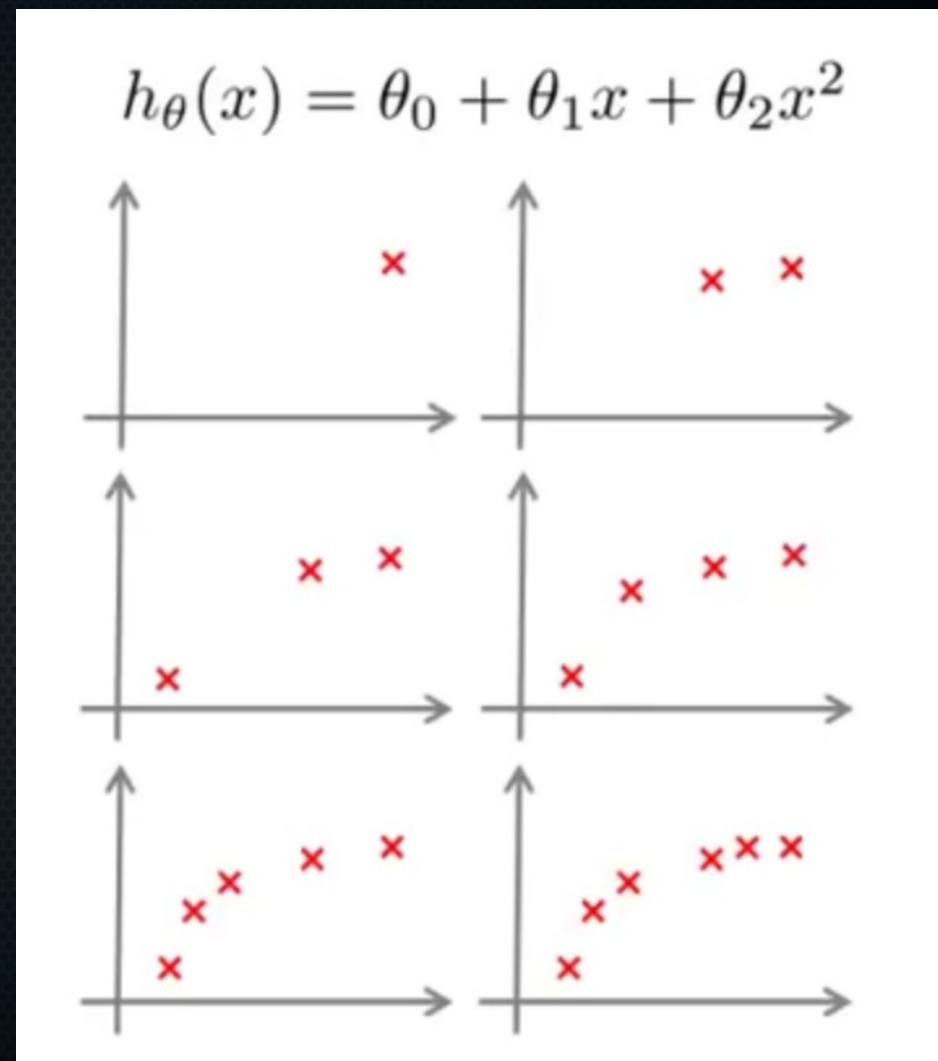
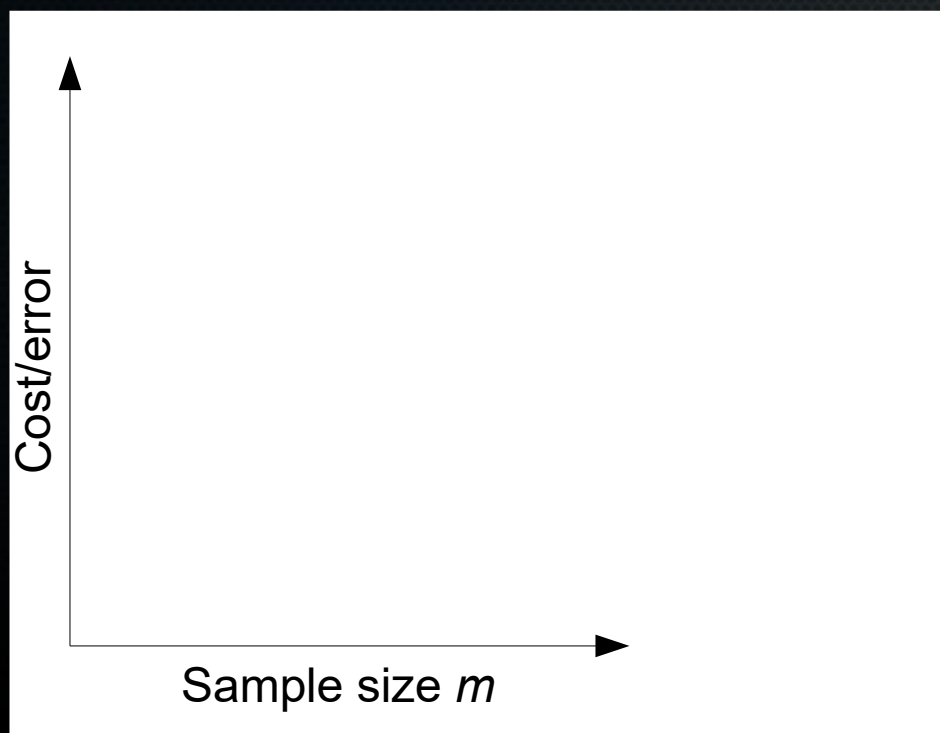
Fold 1

What can we do to find a good model?

- ~~Find a way to approximate generalisation error: how well do you do on unseen data?~~
- **See how error on seen and unseen data changes with amount of training data (plot learning curves)**
- Automatically constrain the fitting by penalising the cost function for too many/too large parameters

Learning curves

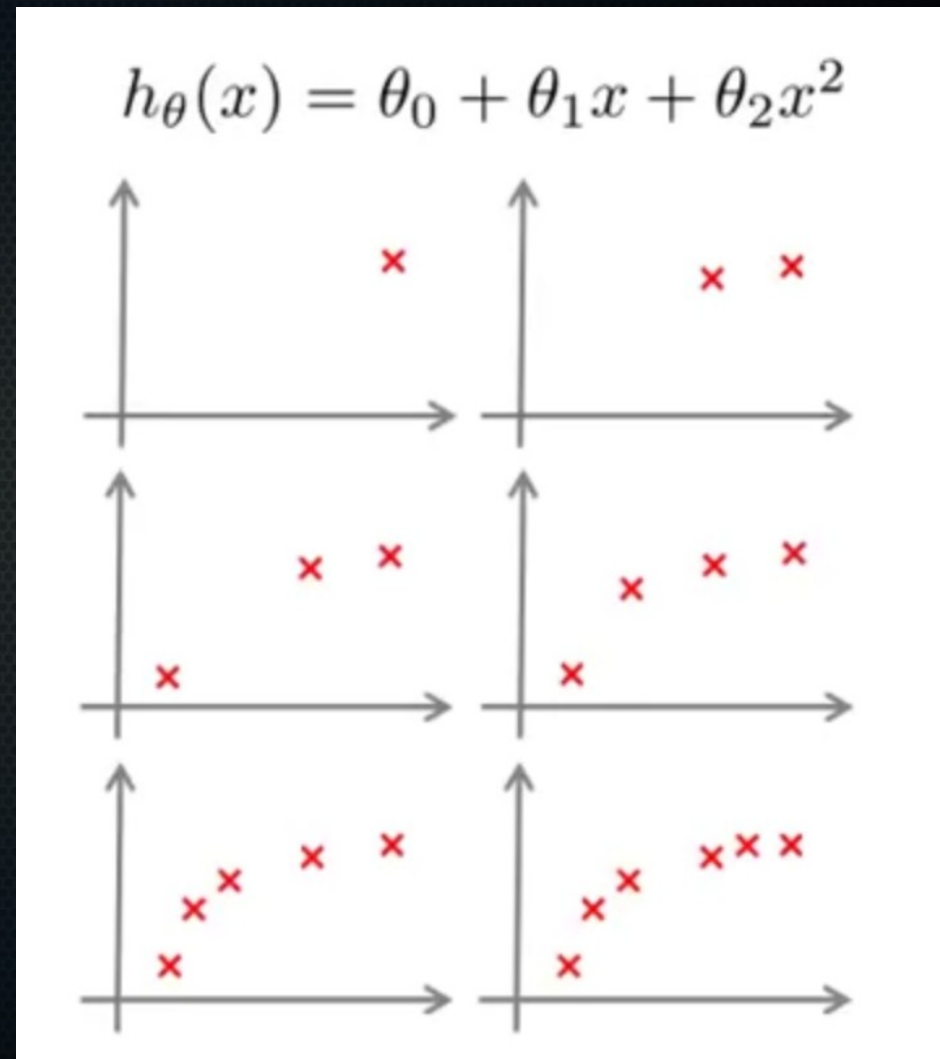
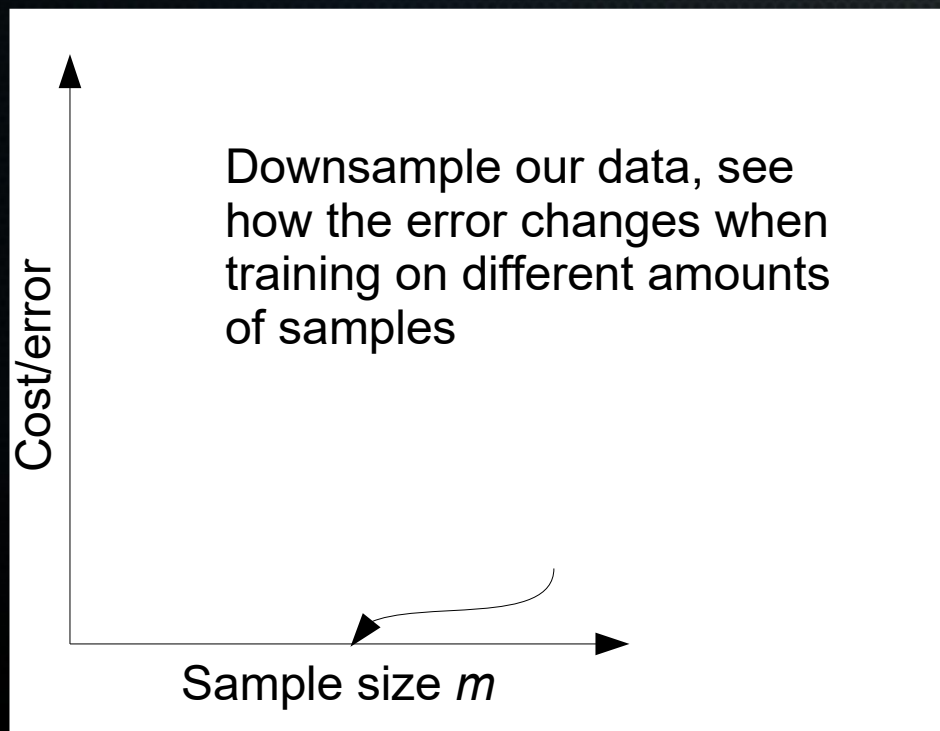
$$J_{train} = \frac{1}{2m_{train}} \sum_{i=1}^{m_{train}} (h_{\theta}(x^{(i)}) - y^{(i)})^2$$
$$J_{val} = \frac{1}{2m_{val}} \sum_{i=1}^{m_{val}} (h_{\theta}(x^{(i)}) - y^{(i)})^2$$



Source: Andrew Ng, Coursera

Learning curves

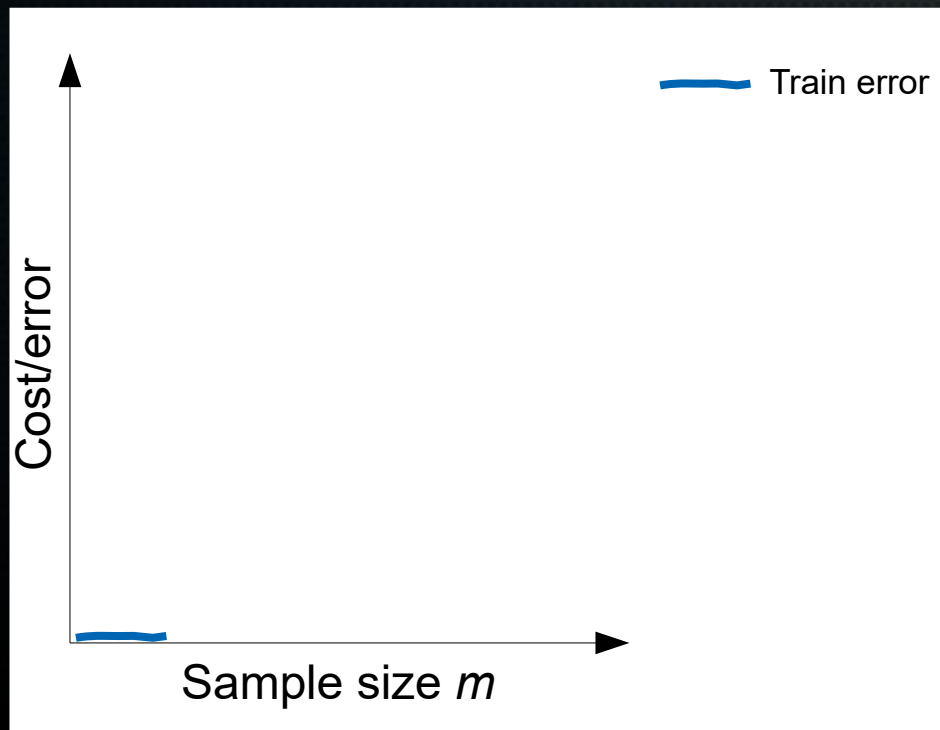
$$J_{train} = \frac{1}{2m_{train}} \sum_{i=1}^{m_{train}} (h_{\theta}(x^{(i)}) - y^{(i)})^2$$
$$J_{val} = \frac{1}{2m_{val}} \sum_{i=1}^{m_{val}} (h_{\theta}(x^{(i)}) - y^{(i)})^2$$



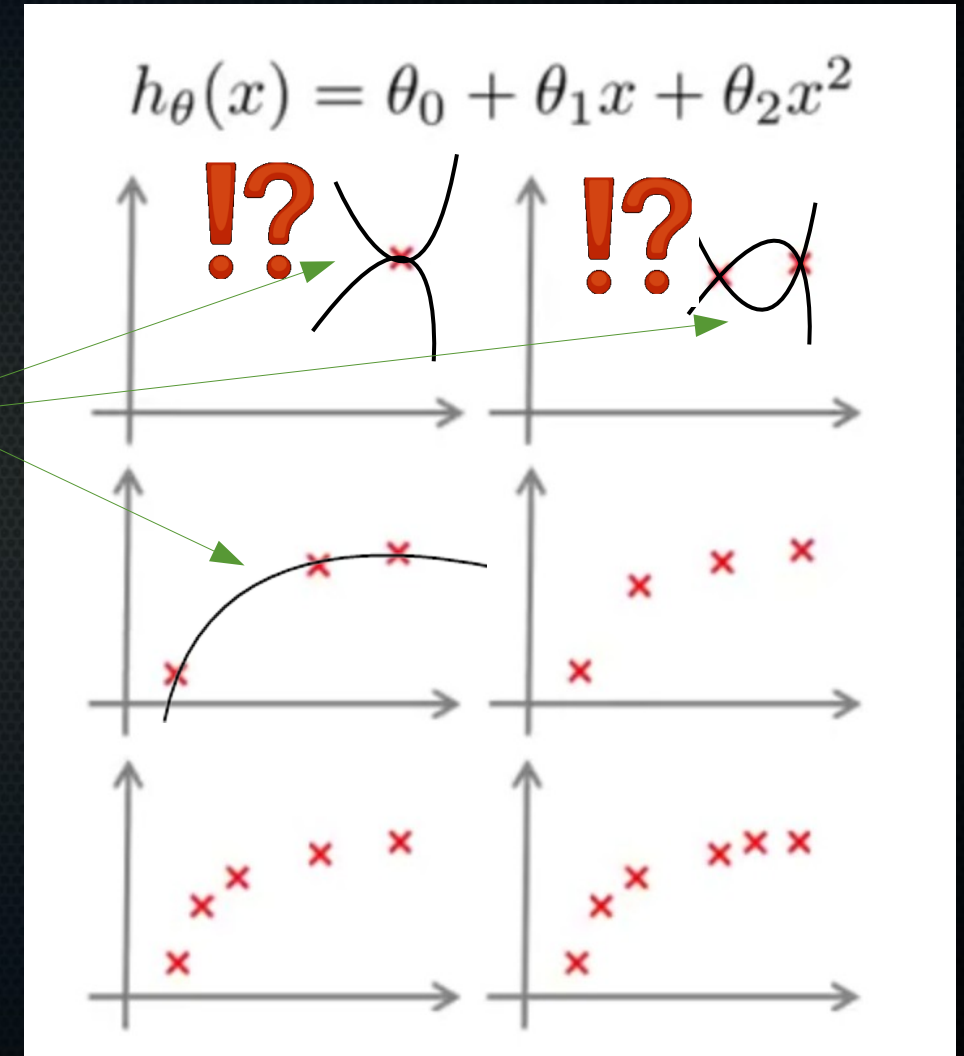
Source: Andrew Ng, Coursera

Learning curves

$$J_{train} = \frac{1}{2m_{train}} \sum_{i=1}^{m_{train}} (h_{\theta}(x^{(i)}) - y^{(i)})^2$$
$$J_{val} = \frac{1}{2m_{val}} \sum_{i=1}^{m_{val}} (h_{\theta}(x^{(i)}) - y^{(i)})^2$$



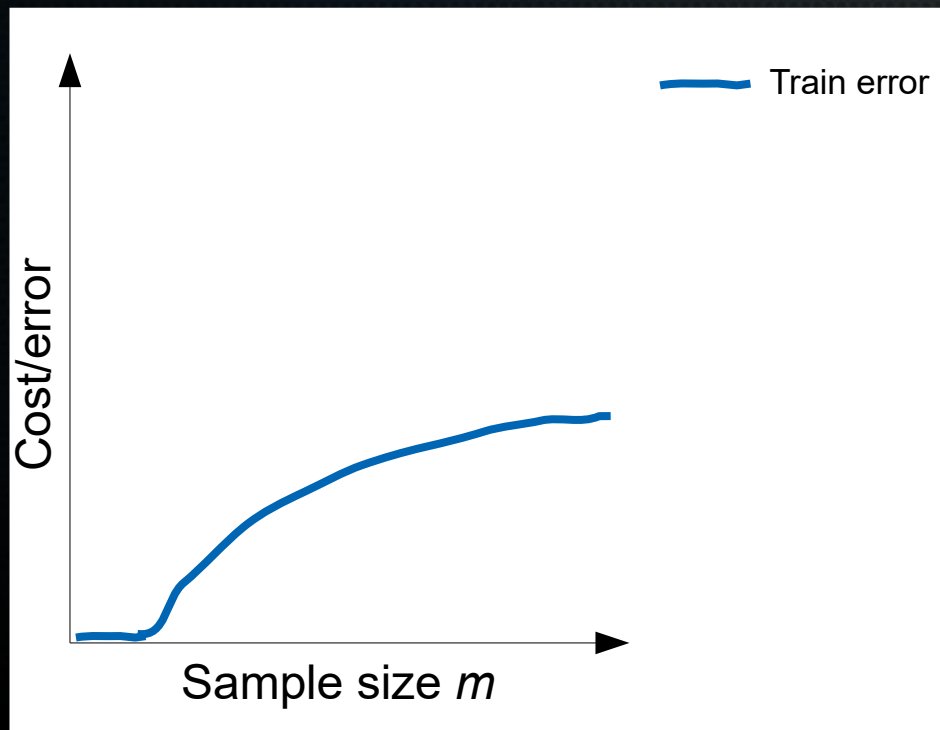
Easy to fit few datapoints perfectly



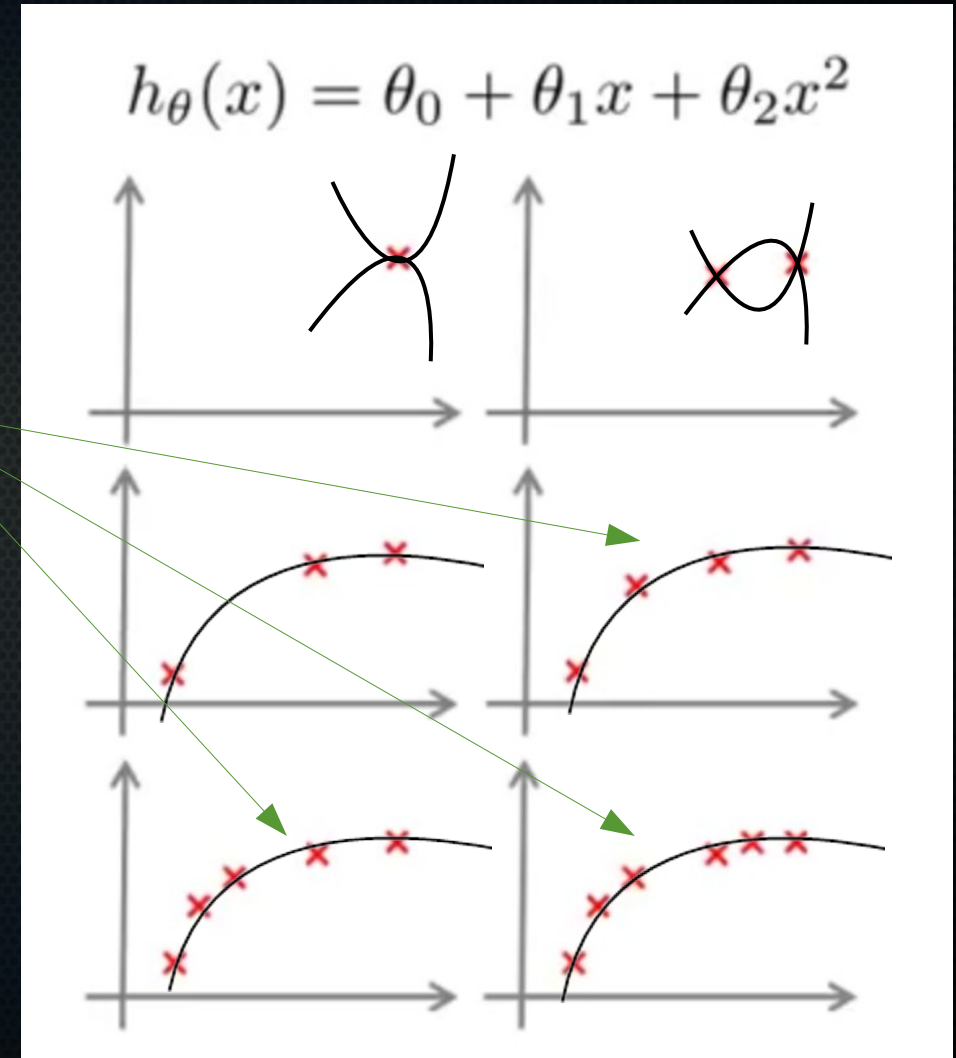
Source: Andrew Ng, Coursera

Learning curves

$$J_{train} = \frac{1}{2m_{train}} \sum_{i=1}^{m_{train}} (h_{\theta}(x^{(i)}) - y^{(i)})^2$$
$$J_{val} = \frac{1}{2m_{val}} \sum_{i=1}^{m_{val}} (h_{\theta}(x^{(i)}) - y^{(i)})^2$$



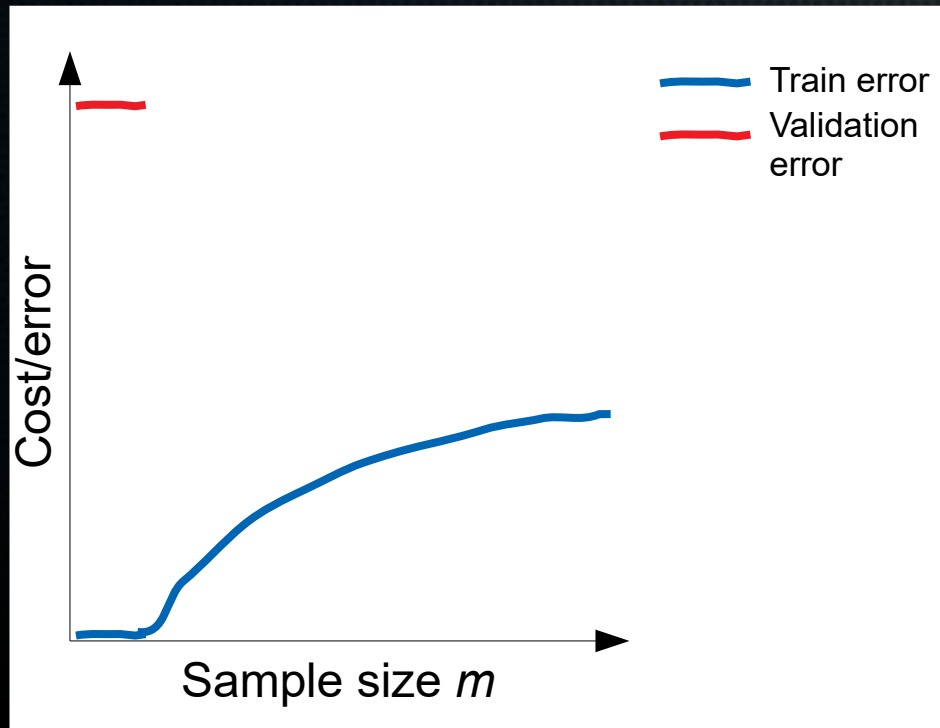
Harder and harder to fit everything perfectly



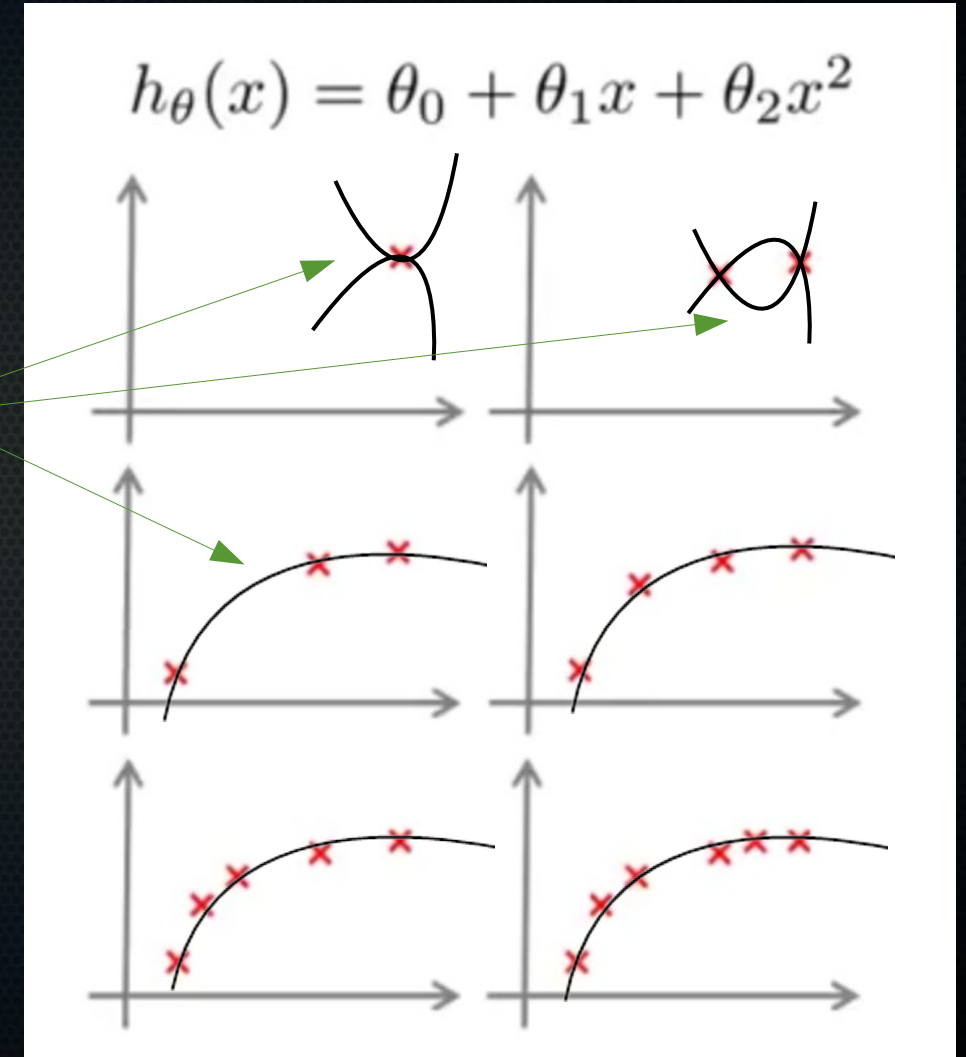
Source: Andrew Ng, Coursera

Learning curves

$$J_{train} = \frac{1}{2m_{train}} \sum_{i=1}^{m_{train}} (h_{\theta}(x^{(i)}) - y^{(i)})^2$$
$$J_{val} = \frac{1}{2m_{val}} \sum_{i=1}^{m_{val}} (h_{\theta}(x^{(i)}) - y^{(i)})^2$$



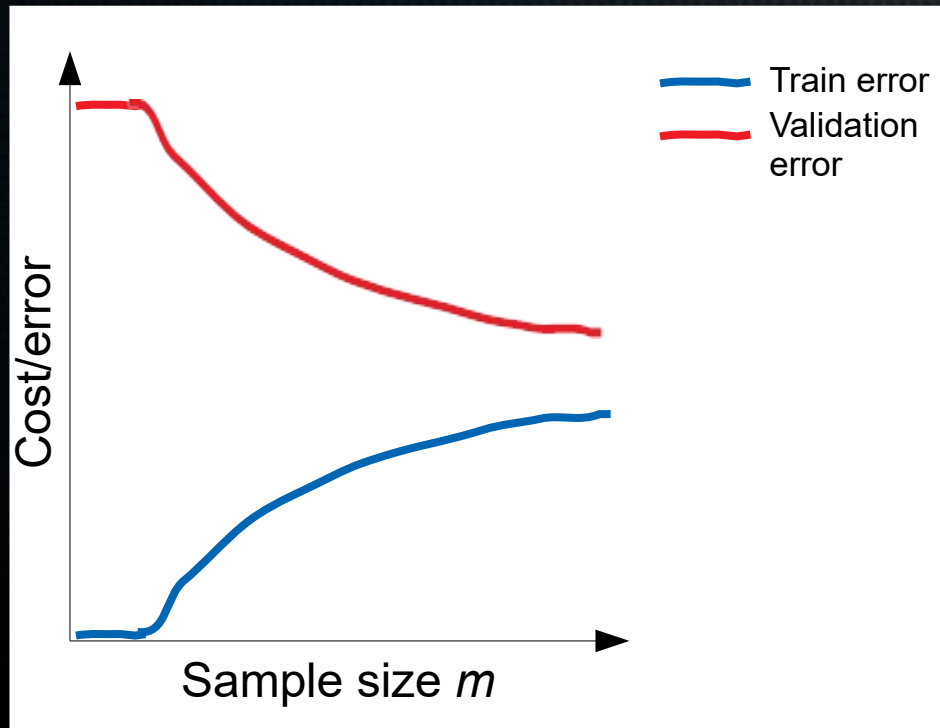
Generalises poorly to new data



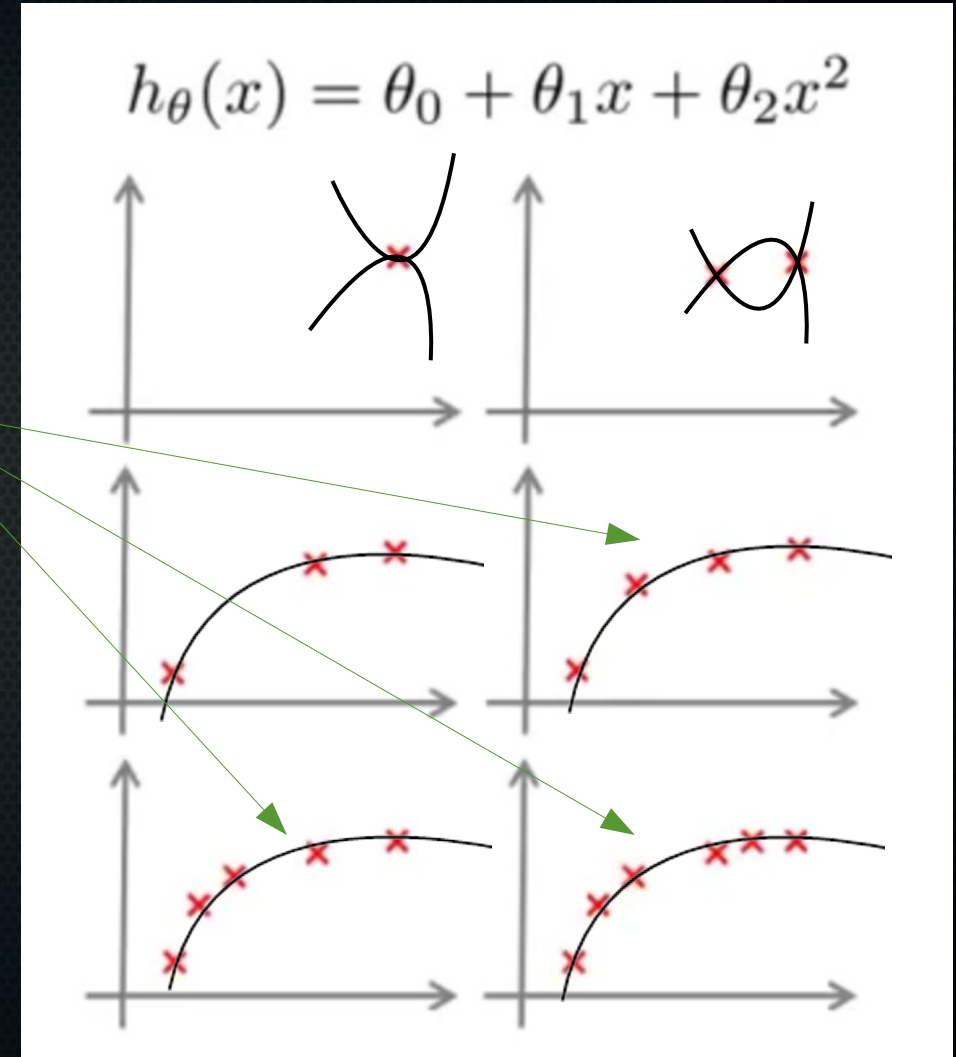
Source: Andrew Ng, Coursera

Learning curves

$$J_{train} = \frac{1}{2m_{train}} \sum_{i=1}^{m_{train}} (h_{\theta}(x^{(i)}) - y^{(i)})^2$$
$$J_{val} = \frac{1}{2m_{val}} \sum_{i=1}^{m_{val}} (h_{\theta}(x^{(i)}) - y^{(i)})^2$$



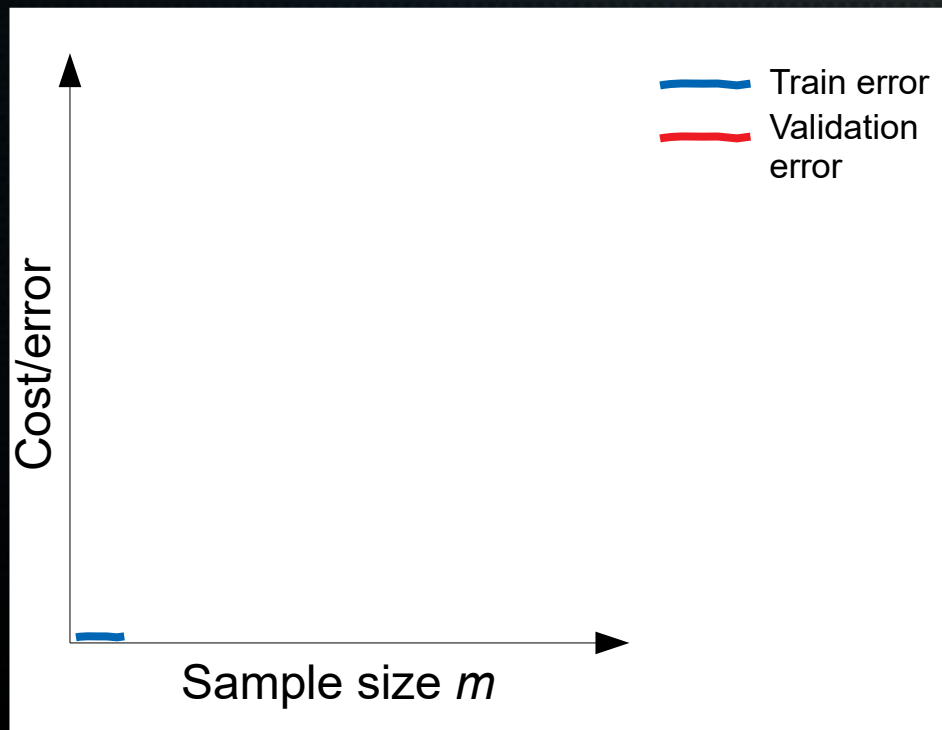
Generalises better
to new data



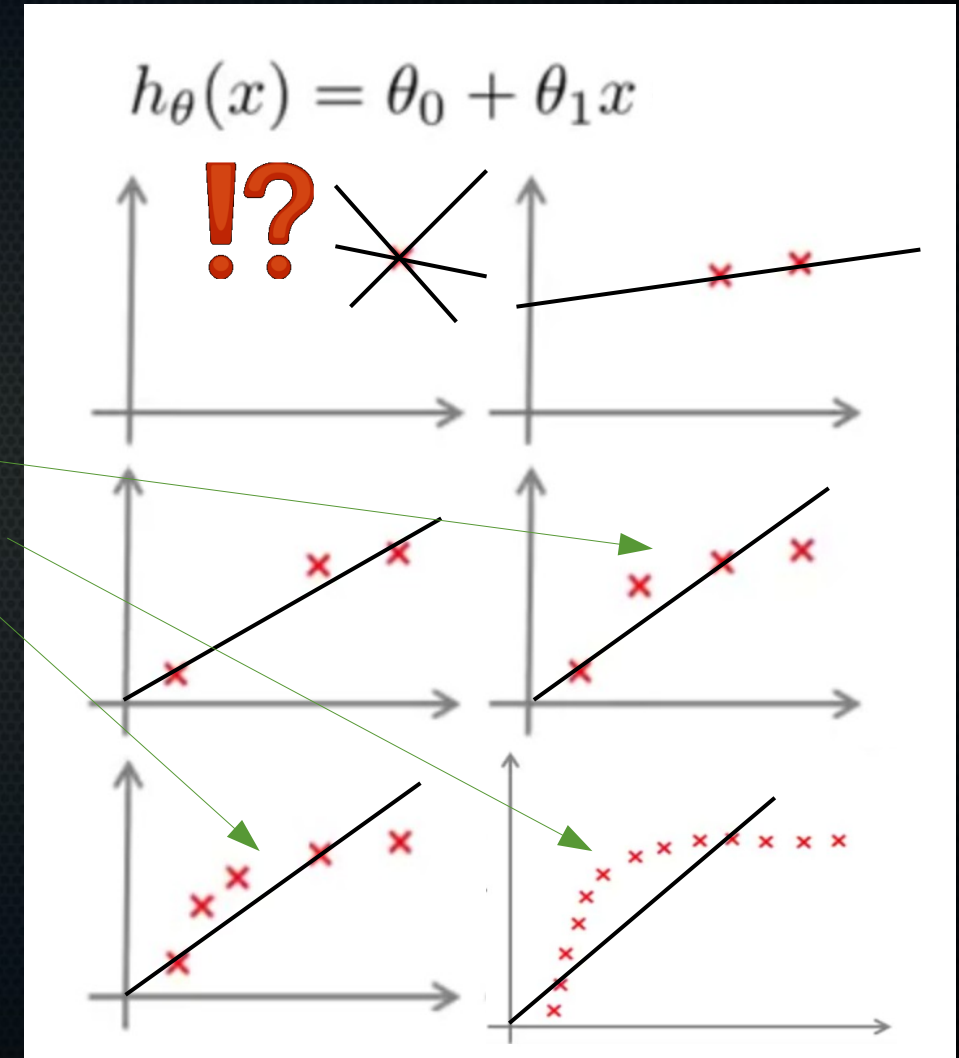
Source: Andrew Ng, Coursera

Learning curves: high bias

$$J_{train} = \frac{1}{2m_{train}} \sum_{i=1}^{m_{train}} (h_{\theta}(x^{(i)}) - y^{(i)})^2$$
$$J_{val} = \frac{1}{2m_{val}} \sum_{i=1}^{m_{val}} (h_{\theta}(x^{(i)}) - y^{(i)})^2$$



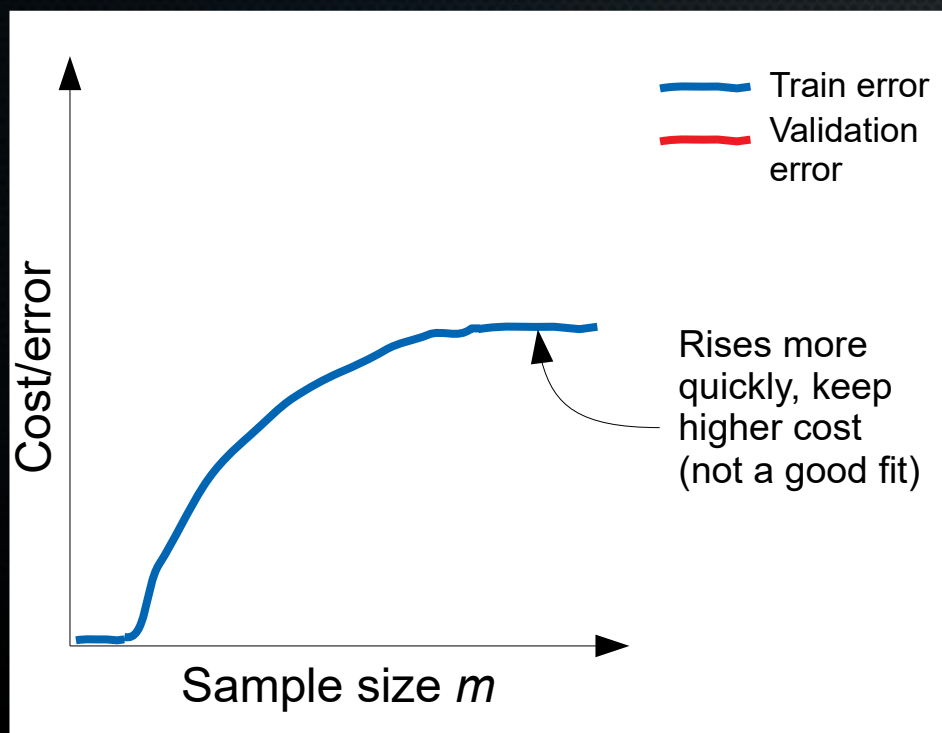
Keep pretty much the same line even as data increases



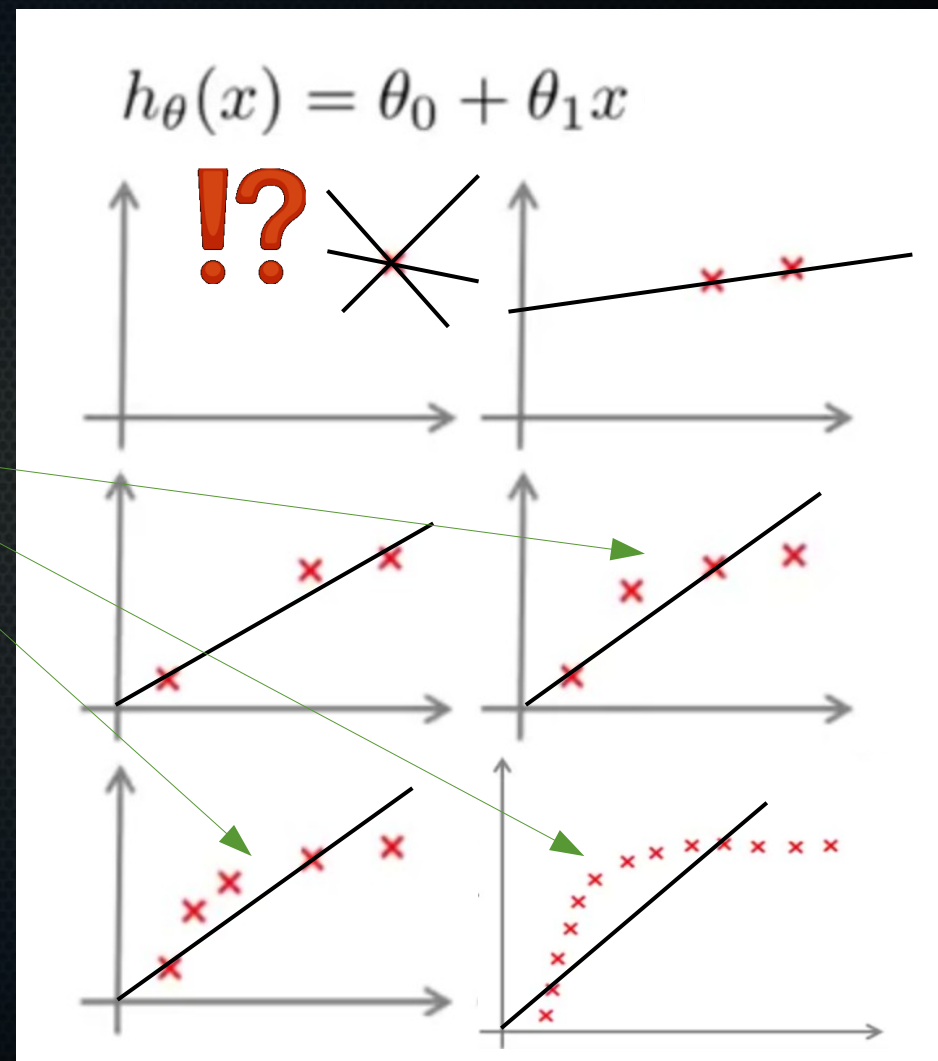
Source: Andrew Ng, Coursera

Learning curves: high bias

$$J_{train} = \frac{1}{2m_{train}} \sum_{i=1}^{m_{train}} (h_{\theta}(x^{(i)}) - y^{(i)})^2$$
$$J_{val} = \frac{1}{2m_{val}} \sum_{i=1}^{m_{val}} (h_{\theta}(x^{(i)}) - y^{(i)})^2$$



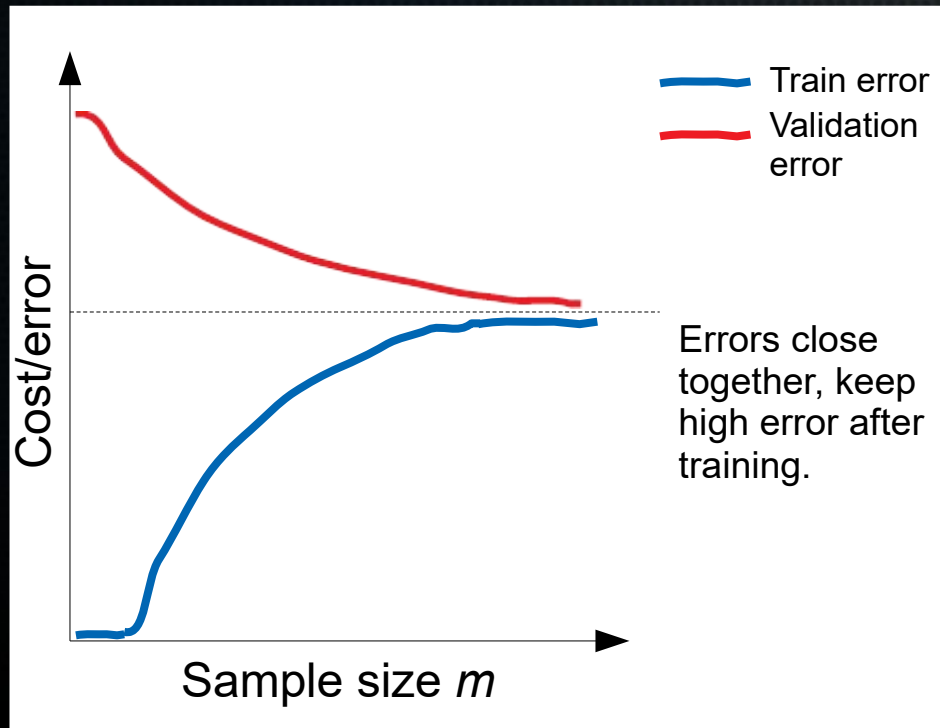
Keep pretty much the same line even as data increases



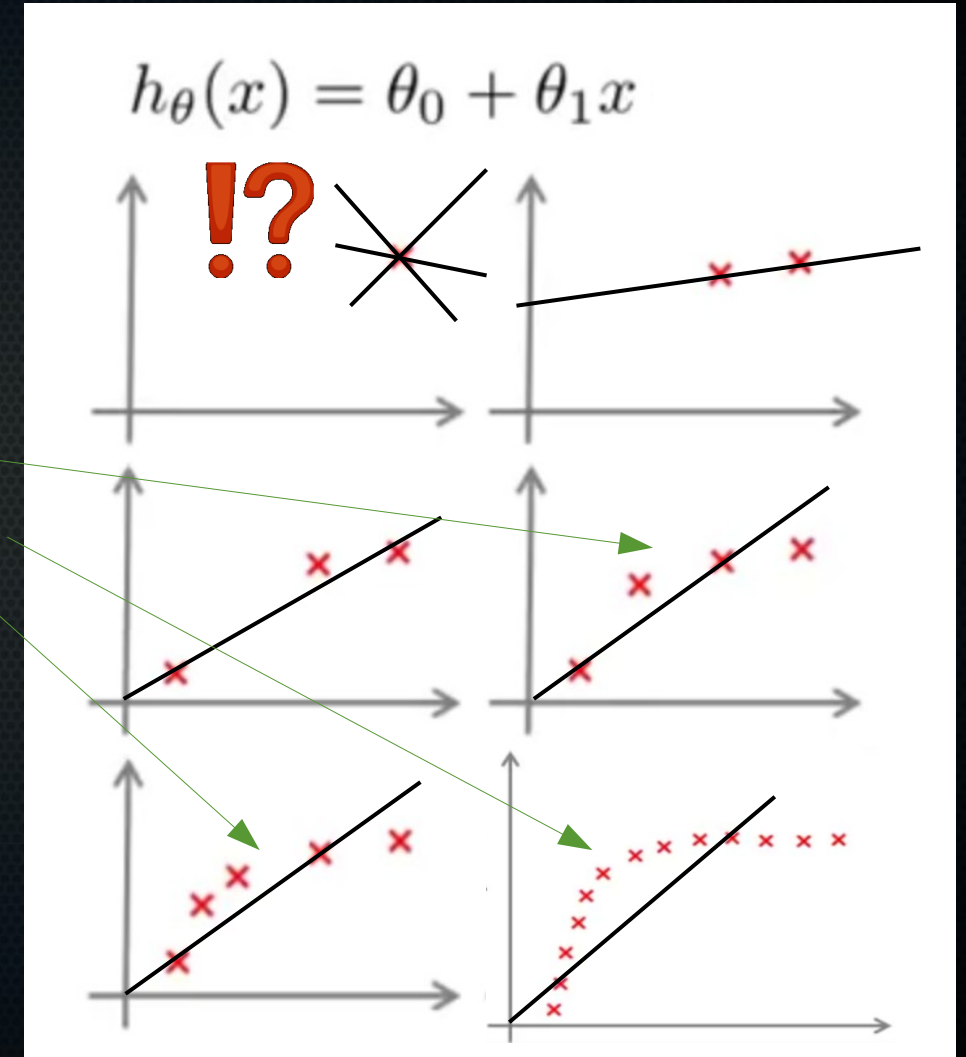
Source: Andrew Ng, Coursera

Learning curves: high bias

$$J_{train} = \frac{1}{2m_{train}} \sum_{i=1}^{m_{train}} (h_{\theta}(x^{(i)}) - y^{(i)})^2$$
$$J_{val} = \frac{1}{2m_{val}} \sum_{i=1}^{m_{val}} (h_{\theta}(x^{(i)}) - y^{(i)})^2$$



Keep pretty much the same line even as data increases



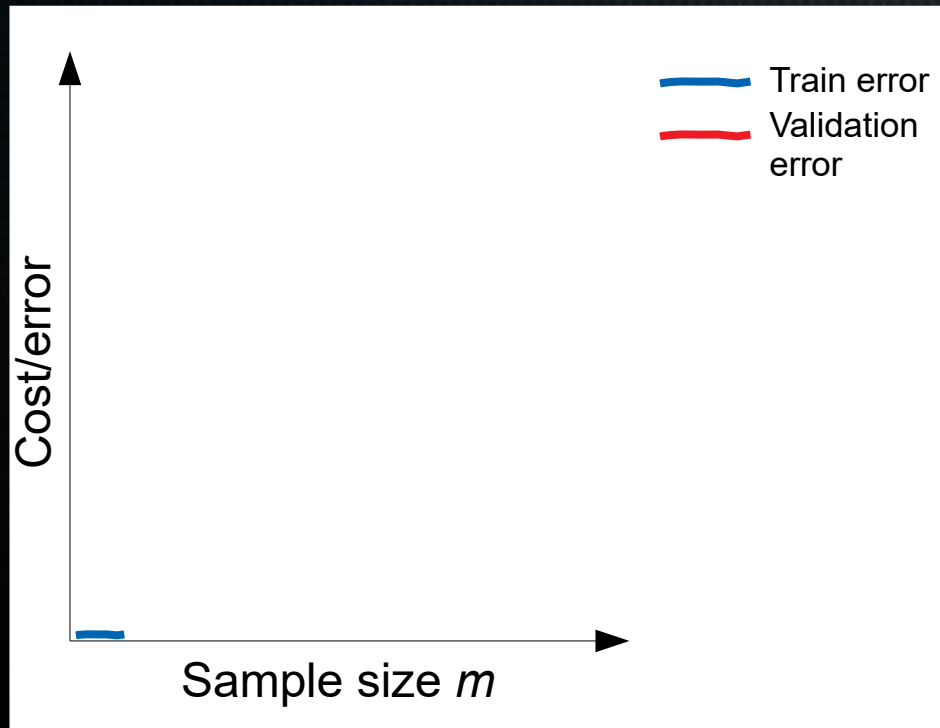
Source: Andrew Ng, Coursera

Learning curves: high bias

- If your learning algorithm is biased, getting more training data will not help!

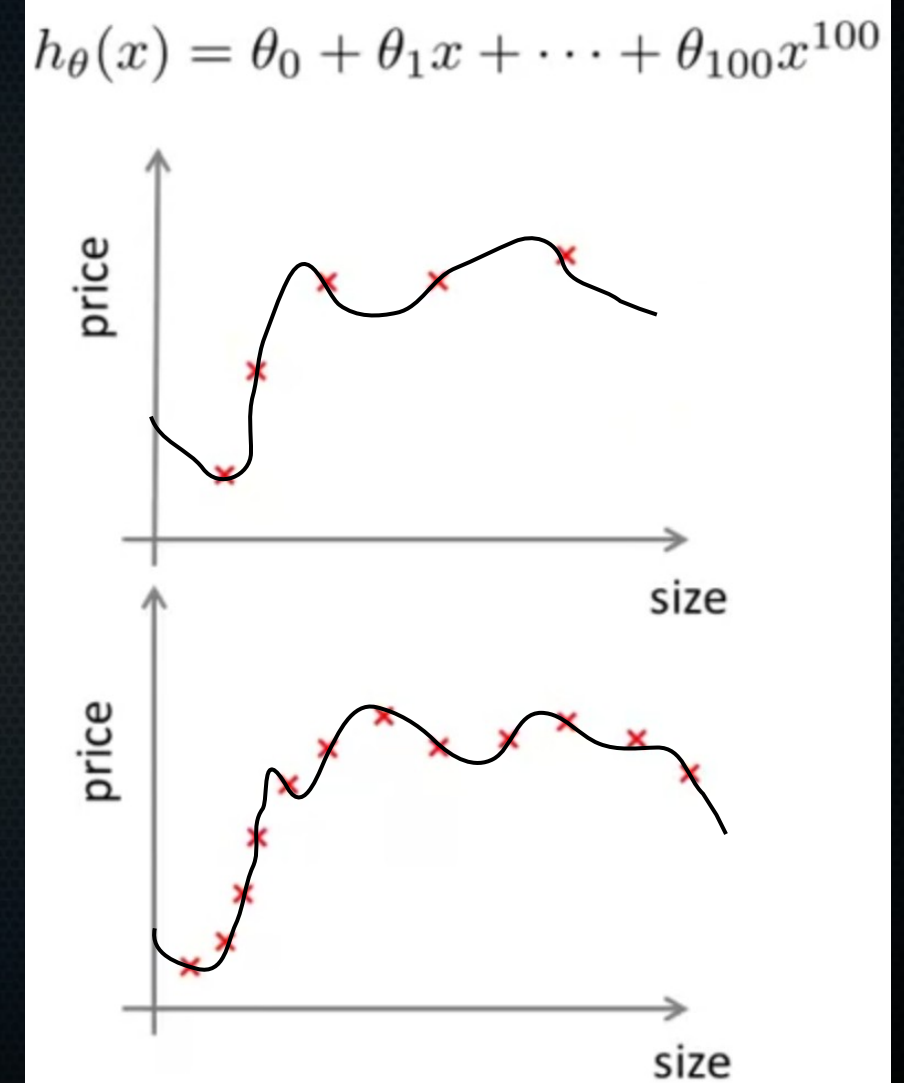
Learning curves: high variance

$$J_{train} = \frac{1}{2m_{train}} \sum_{i=1}^{m_{train}} (h_{\theta}(x^{(i)}) - y^{(i)})^2$$
$$J_{val} = \frac{1}{2m_{val}} \sum_{i=1}^{m_{val}} (h_{\theta}(x^{(i)}) - y^{(i)})^2$$



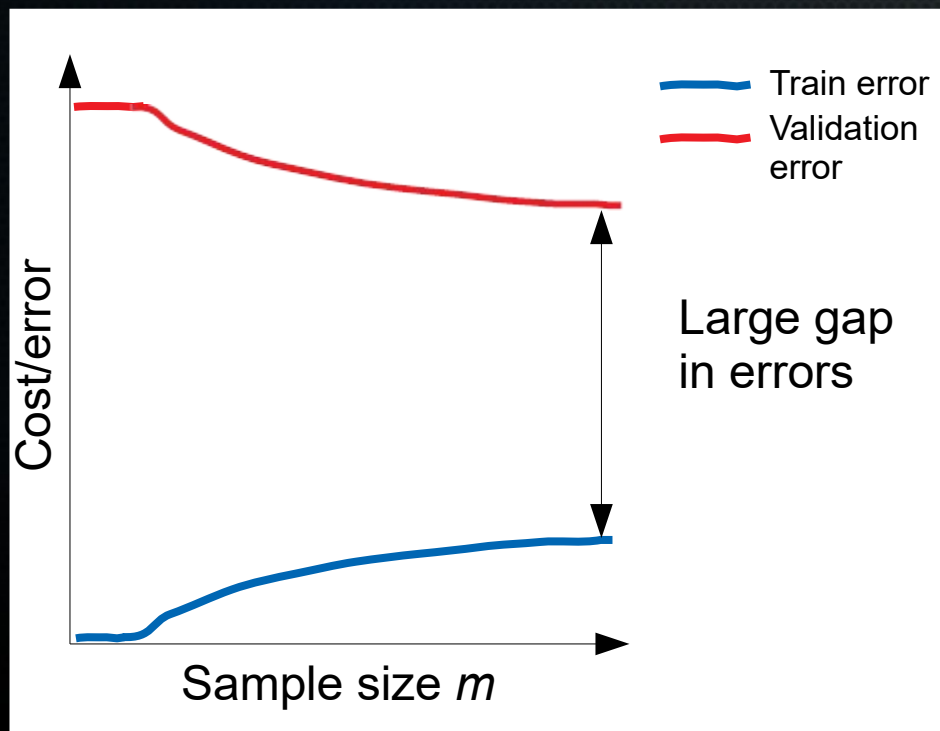
Fit perfectly

Fit very well,
but not perfect



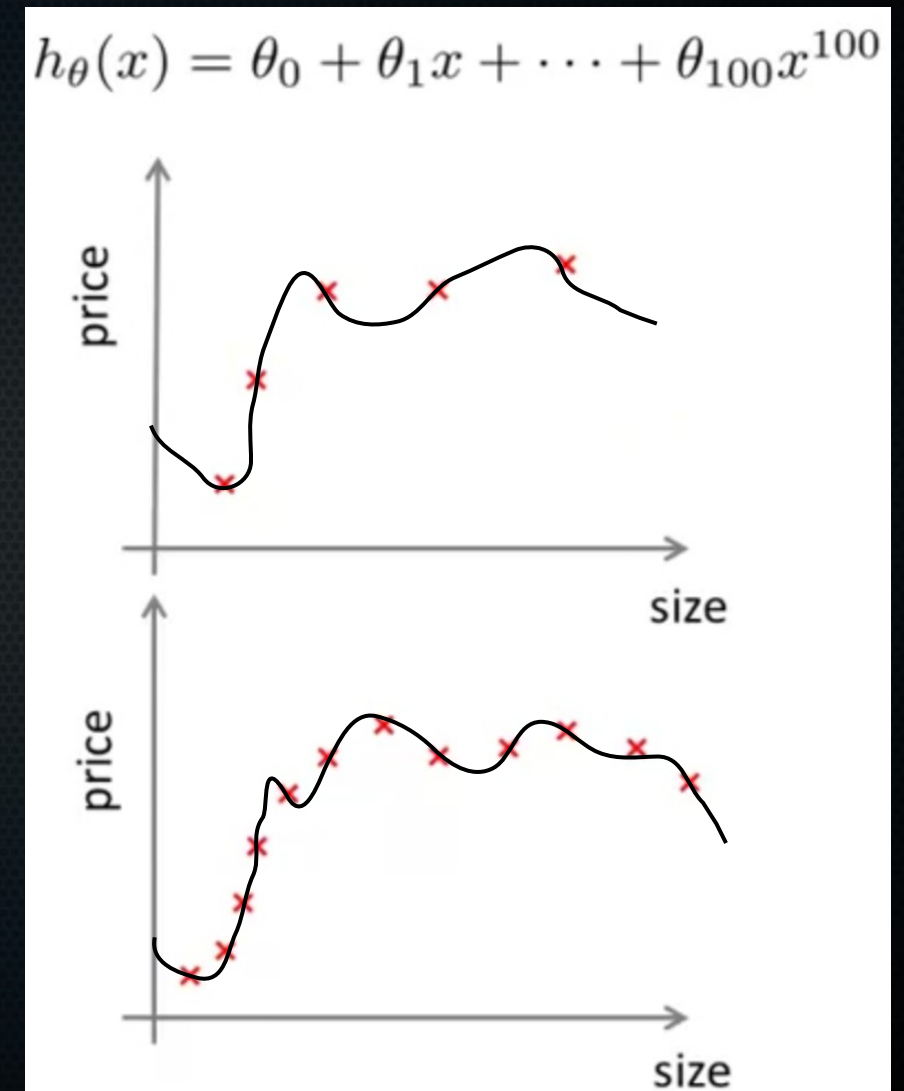
Learning curves: high variance

$$J_{train} = \frac{1}{2m_{train}} \sum_{i=1}^{m_{train}} (h_{\theta}(x^{(i)}) - y^{(i)})^2$$
$$J_{val} = \frac{1}{2m_{val}} \sum_{i=1}^{m_{val}} (h_{\theta}(x^{(i)}) - y^{(i)})^2$$



Fit perfectly

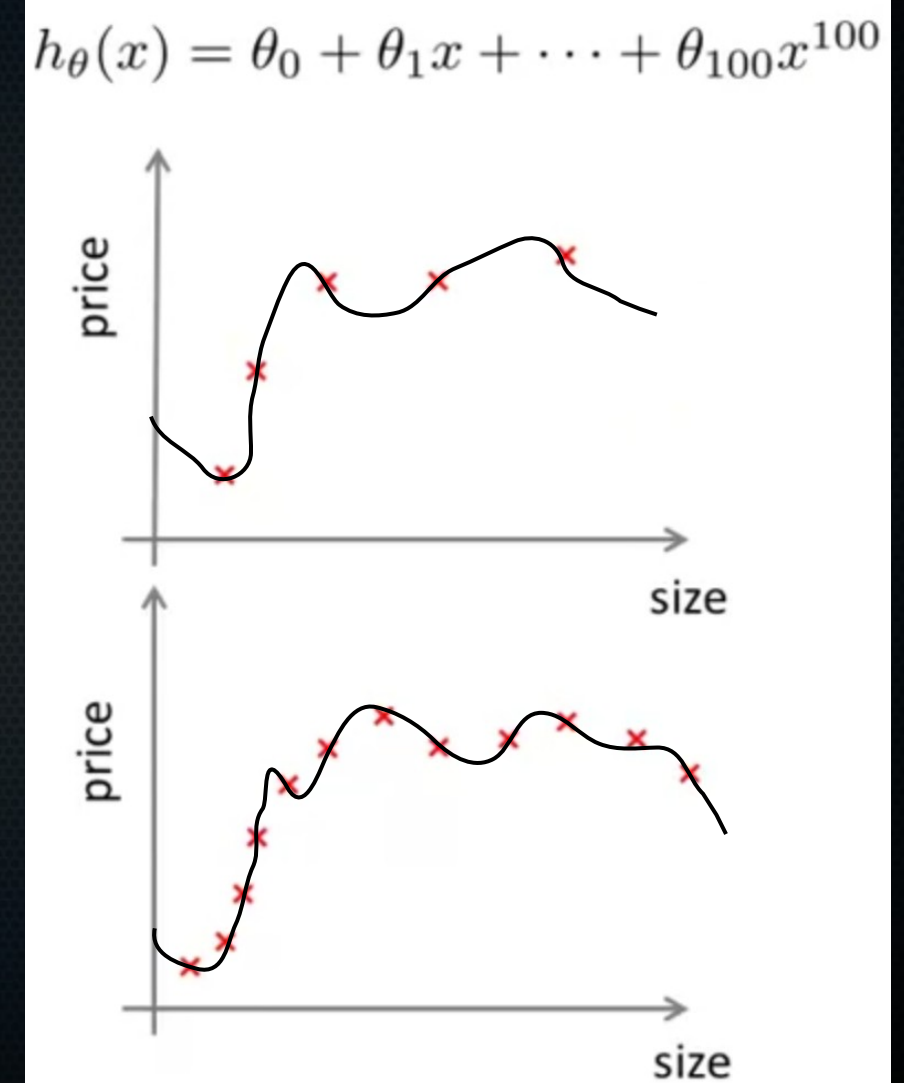
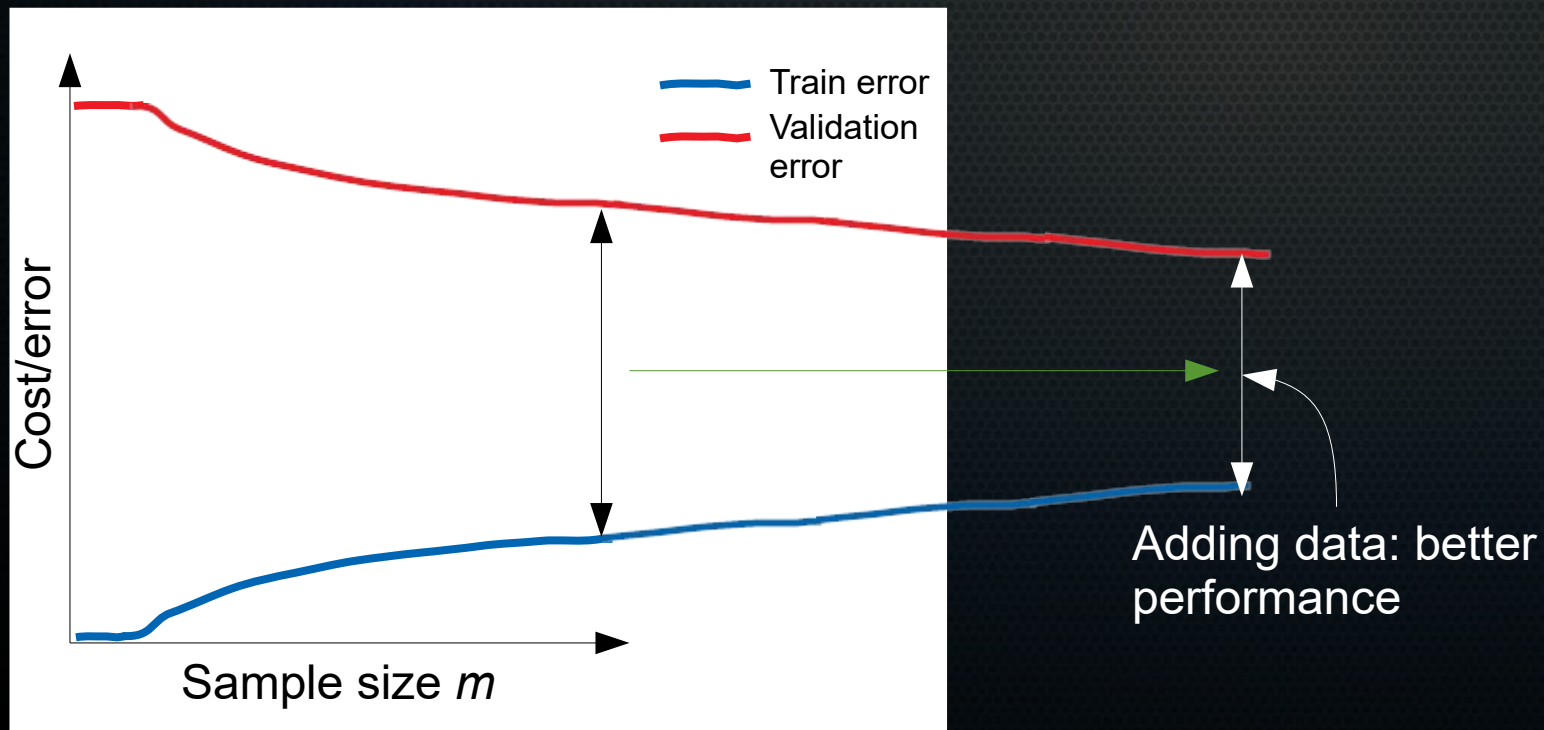
Fit very well,
but not perfect



Learning curves: high variance

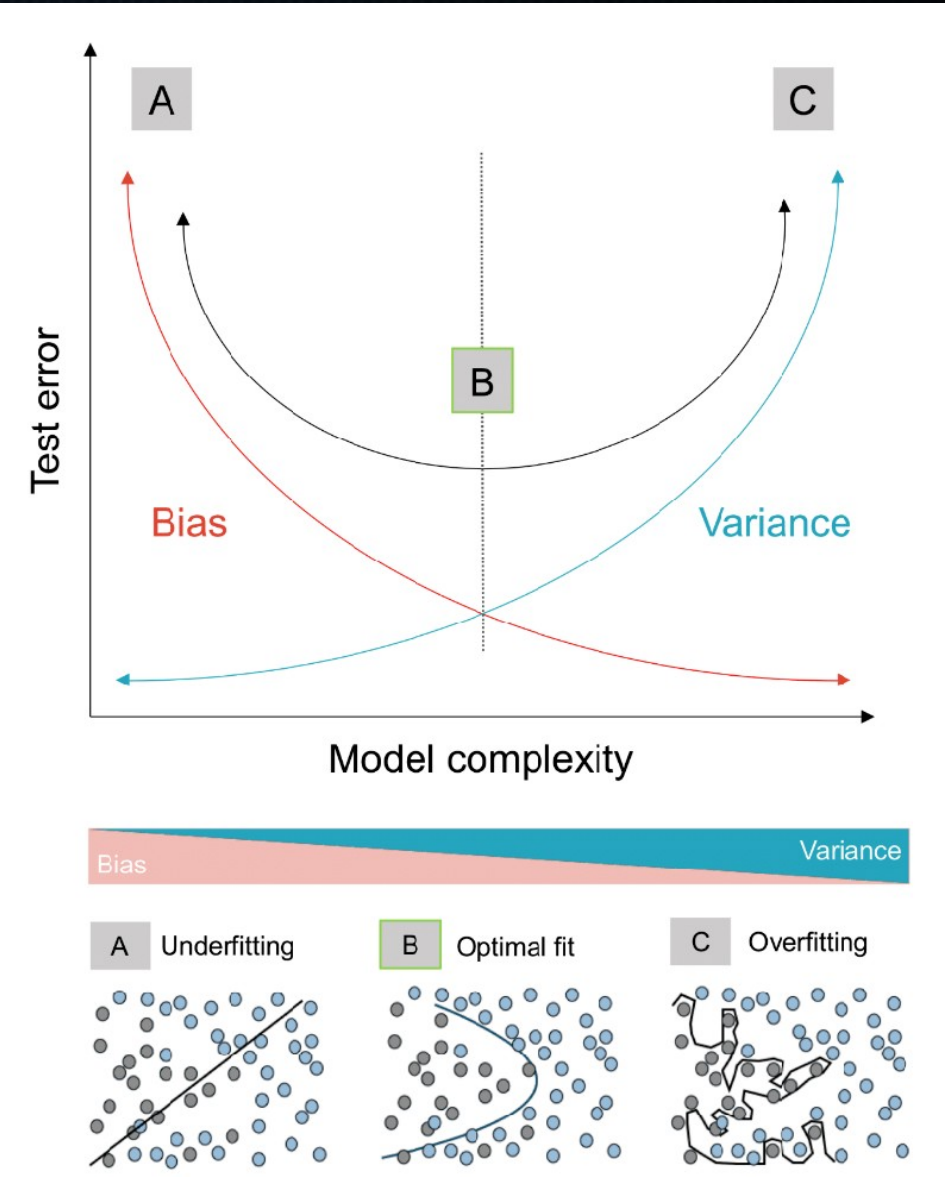
$$J_{train} = \frac{1}{2m_{train}} \sum_{i=1}^{m_{train}} (h_{\theta}(x^{(i)}) - y^{(i)})^2$$

$$J_{val} = \frac{1}{2m_{val}} \sum_{i=1}^{m_{val}} (h_{\theta}(x^{(i)}) - y^{(i)})^2$$



Learning curves

- There's a trade-off between bias and variance
- Learning curves allow you to diagnose what your algorithm might be suffering from



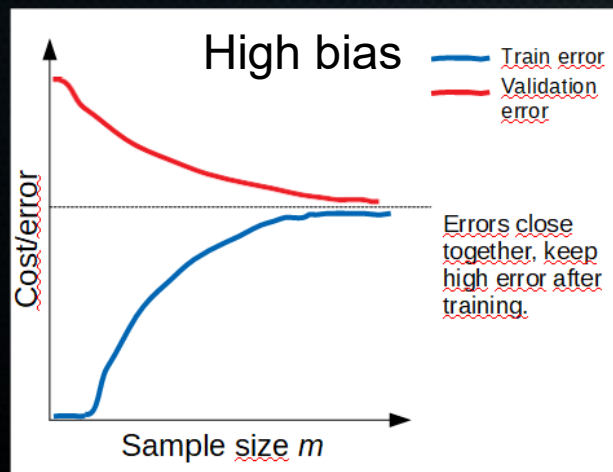
Source: Kernbach, J. M., & Staartjes, V. E. (2020). Machine learning-based clinical prediction modeling--A practical guide for clinicians. arXiv preprint arXiv:2006.15069.

Summary: cross-validation and learning curves

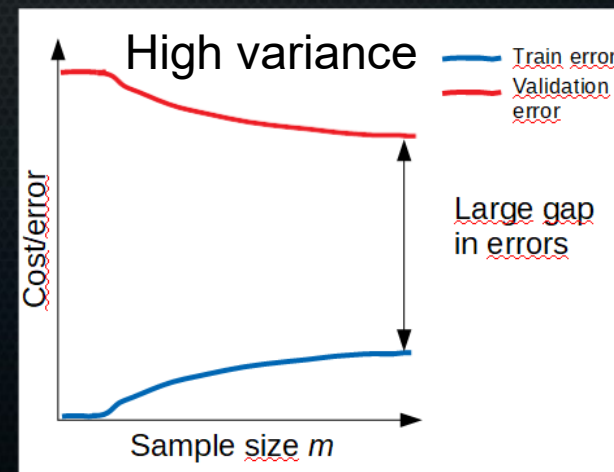
- Data is split into data to train on and a test set that you don't touch at all
- Training data is split into k folds, with (average) cross-validation error as proxy for how your algorithm performs on unseen data

Summary: cross-validation and learning curves

- Data is split into data to train on and a test set that you don't touch at all
- Training data is split into k folds, with (average) cross-validation error as proxy for how your algorithm performs on unseen data
- Learning curves allow you to diagnose whether your algorithm suffers from high bias or high variance



Underfitting:
use more
complex model



Overfitting: use
less complex
model or supply
more training
data

What can we do to find a good model?

- ~~Find a way to approximate generalisation error: how well do you do on unseen data?~~
- ~~See how error on seen and unseen data changes with amount of training data (plot learning curves)~~
- ***Automatically constrain the fitting by penalising the cost function for too many/too large parameters***

Regularisation

- Change the cost function to apply a cost for complexity

Regularisation

- Change the cost function to apply a cost for complexity

$$J(\theta_0, \dots, \theta_n) = \frac{1}{2m} \sum_{i=1}^m (h_{\theta}(x) - y^{(i)})^2$$

Regularisation: ridge regression

- Change the cost function to apply a cost for complexity

$$J(\theta_0, \dots, \theta_n) = \frac{1}{2m} \sum_{i=1}^m (h_{\theta}(x) - y^{(i)})^2$$

$$J(\theta_0, \dots, \theta_n) = \frac{1}{2m} \sum_{i=1}^m (h_{\theta}(x) - y^{(i)})^2 + \lambda \sum_{j=1}^n (\theta_j)^2$$

- Make J a function of both the error of predictions given some parameters *and* the magnitude of those parameters themselves

Regularisation: ridge regression

- Change the cost function to apply a cost for complexity

$$J(\theta_0, \dots, \theta_n) = \frac{1}{2m} \sum_{i=1}^m (h_{\theta}(x) - y^{(i)})^2$$

$$J(\theta_0, \dots, \theta_n) = \frac{1}{2m} \sum_{i=1}^m (h_{\theta}(x) - y^{(i)})^2 + \lambda \sum_{j=1}^n (\theta_j)^2 \quad \left. \vphantom{\sum_{j=1}^n} \right\} \begin{array}{l} \text{By convention: don't} \\ \text{shrink bias/intercept} \\ \text{term} \end{array}$$

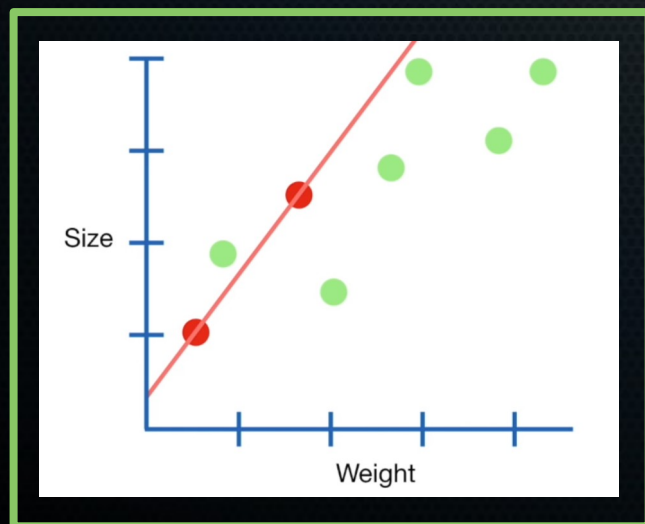
- Make J a function of both the error of predictions given some parameters *and* the magnitude of those parameters themselves

Regularisation: ridge regression

- Change the cost function to apply a cost for complexity

$$J(\theta_0, \dots, \theta_n) = \frac{1}{2m} \sum_{i=1}^m (h_{\theta}(x) - y^{(i)})^2 + \lambda \sum_{j=1}^n (\theta_j)^2 \quad h_{\theta}(x) = \theta_0 + \theta_1 x$$

- Add some **bias** (constrain hypothesis to a set with small parameter values) but reduces **variance**:

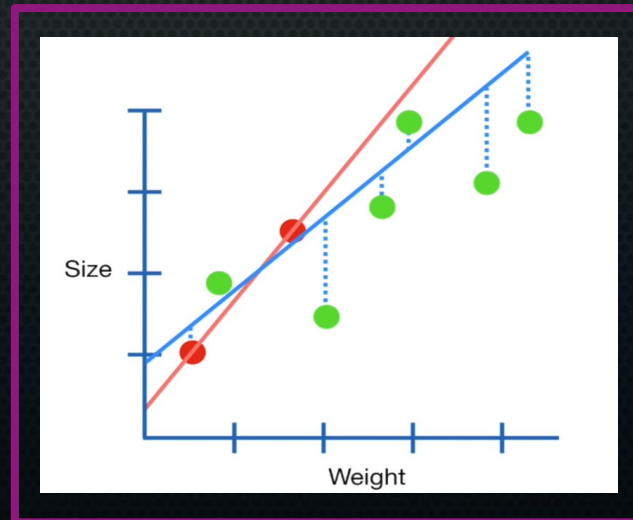
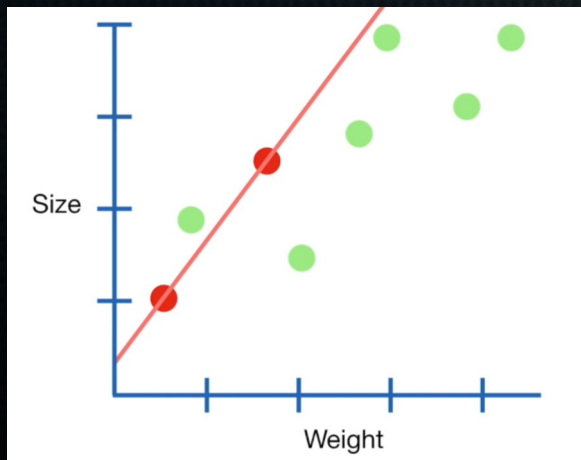


Regularisation: ridge regression

- Change the cost function to apply a cost for complexity

$$J(\theta_0, \dots, \theta_n) = \frac{1}{2m} \sum_{i=1}^m (h_{\theta}(x) - y^{(i)})^2 + \lambda \sum_{j=1}^n (\theta_j)^2 \quad h_{\theta}(x) = \theta_0 + \theta_1 x$$

- Add some **bias** (constrain hypothesis to a set with small parameter values) but reduces **variance**:



Constrained how much the line may increase with Weight (biased) → generalises better to test set

Regularisation: LASSO regression

- Same idea, slightly different execution:

$$J(\theta_0, \dots, \theta_n) = \frac{1}{2m} \sum_{i=1}^m (h_{\theta}(x) - y^{(i)})^2 + \lambda \sum_{j=1}^n (\theta_j)^2$$

Regularisation: LASSO regression

- Same idea, slightly different execution:

$$J(\theta_0, \dots, \theta_n) = \frac{1}{2m} \sum_{i=1}^m (h_{\theta}(x) - y^{(i)})^2 + \lambda \sum_{j=1}^n (\theta_j)^2$$

Regularisation: LASSO regression

- Same idea, slightly different execution:

$$J(\theta_0, \dots, \theta_n) = \frac{1}{2m} \sum_{i=1}^m (h_{\theta}(x) - y^{(i)})^2 + \lambda \sum_{j=1}^n |\theta_j|$$

Regularisation: LASSO regression

- Same idea, slightly different execution:

$$J(\theta_0, \dots, \theta_n) = \frac{1}{2m} \sum_{i=1}^m (h_{\theta}(x) - y^{(i)})^2 + \lambda \sum_{j=1}^n |\theta_j|$$

- Ridge regression *shrinks* parameters/weights to 0, LASSO can make them 0 outright, i.e. simply removes uninformative features.

Regularisation: LASSO regression

- Same idea, slightly different execution:

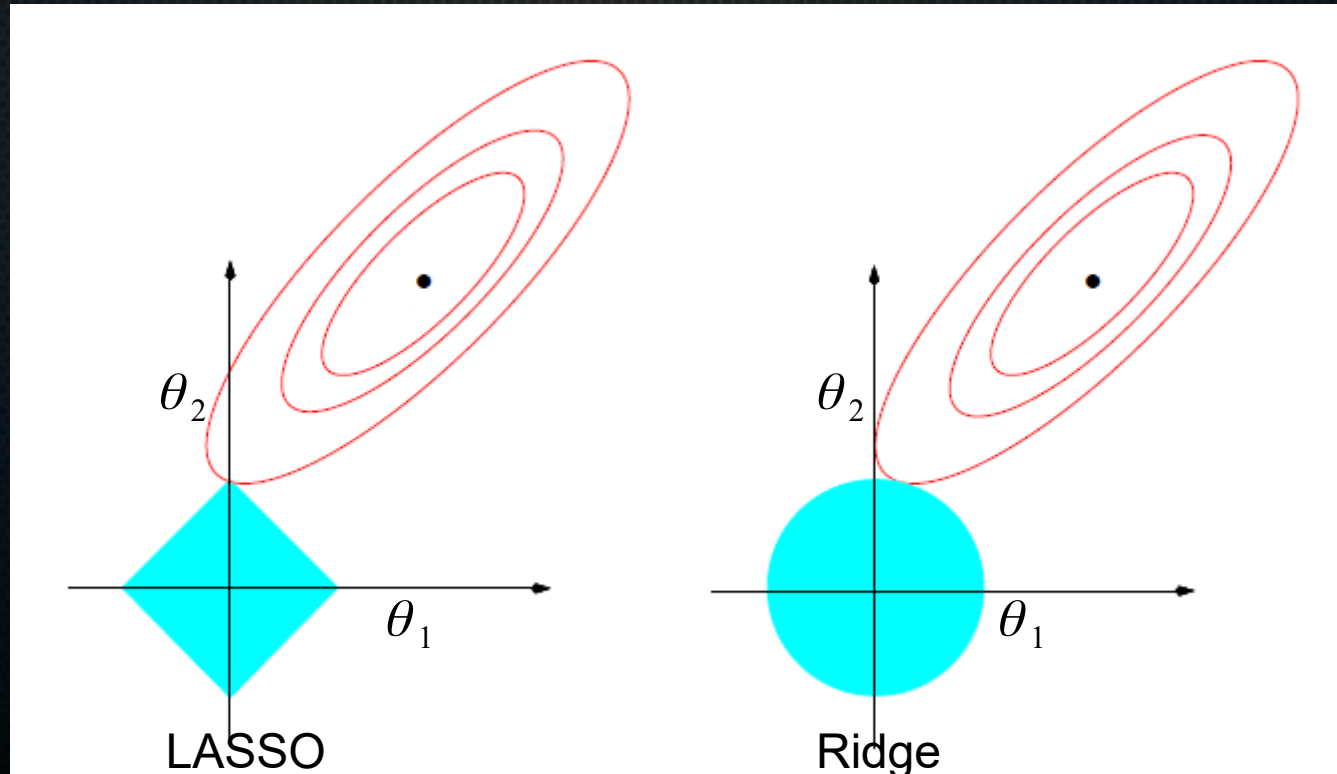
$$J(\theta_0, \dots, \theta_n) = \frac{1}{2m} \sum_{i=1}^m (h_{\theta}(x) - y^{(i)})^2 + \lambda \sum_{j=1}^n |\theta_j|$$

- Ridge regression *shrinks* parameters/weights to 0, LASSO can make them 0 outright, i.e. simply removes uninformative features. → Why?

Regularisation: Ridge vs. LASSO regression

- Ridge regression *shrinks* parameters/weights to 0, LASSO can make them 0 outright, i.e. simply removes uninformative features. → Why?

$$\frac{1}{2m} \sum_{i=1}^m (h_{\theta}(x) - y^{(i)})^2$$

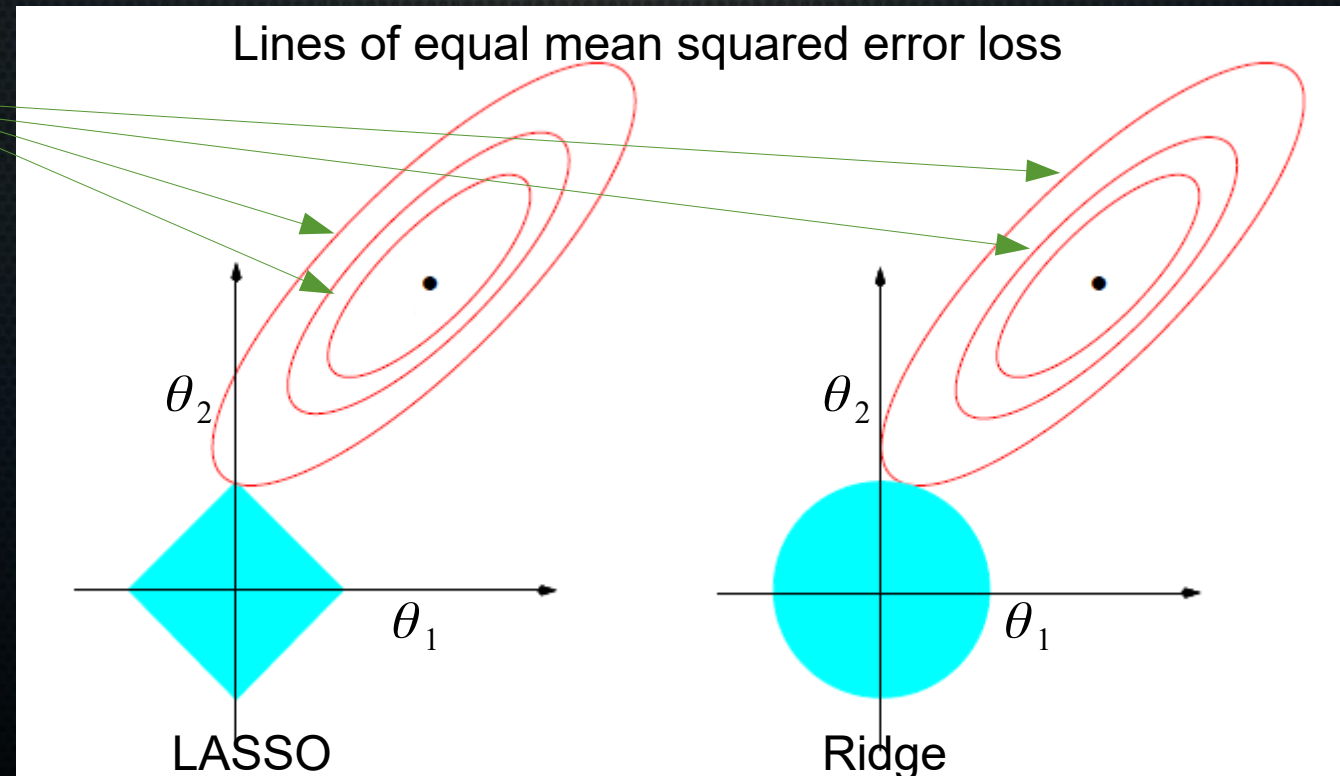


Source: Hastie, T., Tibshirani, R., & Friedman, J. (2009). The elements of statistical learning: data mining, inference, and prediction. Springer Science & Business Media.

Regularisation: Ridge vs. LASSO regression

- Ridge regression *shrinks* parameters/weights to 0, LASSO can make them 0 outright, i.e. simply removes uninformative features. → Why?

$$\frac{1}{2m} \sum_{i=1}^m (h_{\theta}(x) - y^{(i)})^2$$

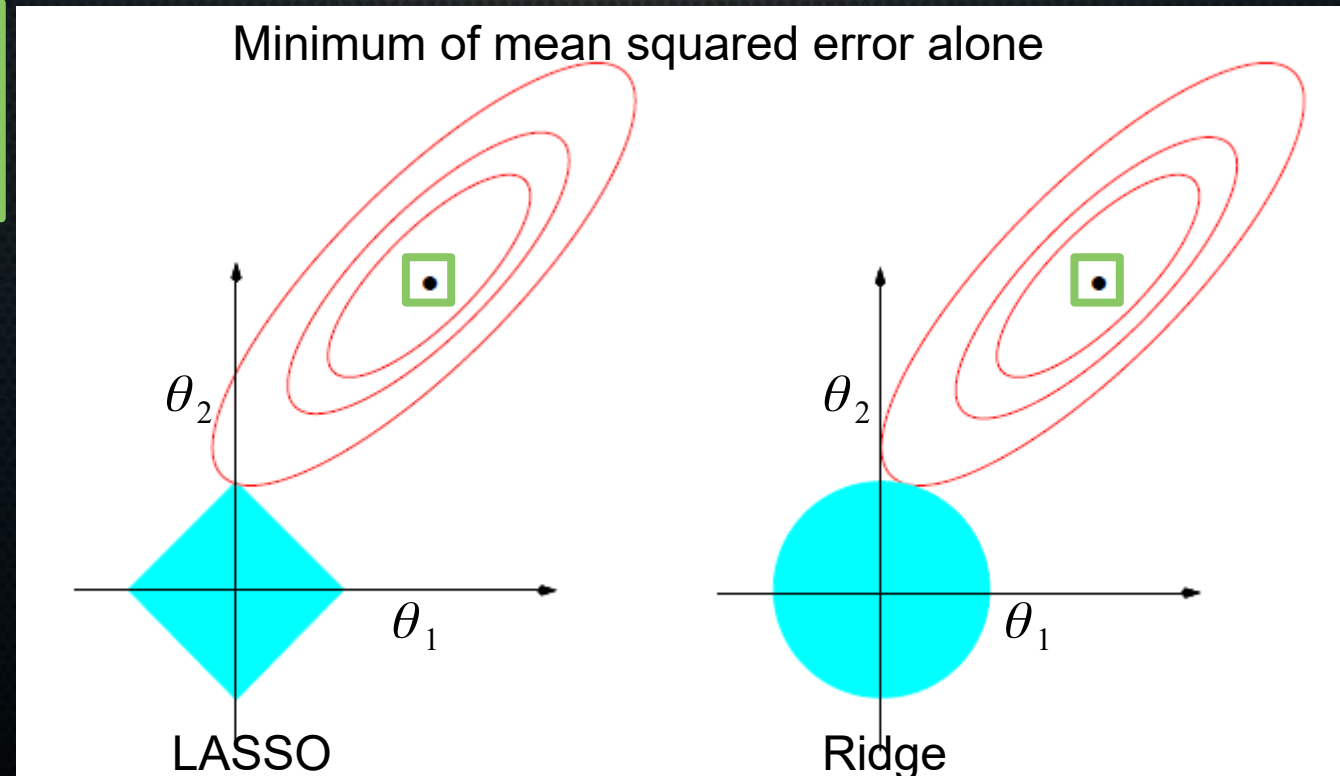


Source: Hastie, T., Tibshirani, R., & Friedman, J. (2009). The elements of statistical learning: data mining, inference, and prediction. Springer Science & Business Media.

Regularisation: Ridge vs. LASSO regression

- Ridge regression *shrinks* parameters/weights to 0, LASSO can make them 0 outright, i.e. simply removes uninformative features. → Why?

$$\frac{1}{2m} \sum_{i=1}^m (h_{\theta}(x) - y^{(i)})^2$$



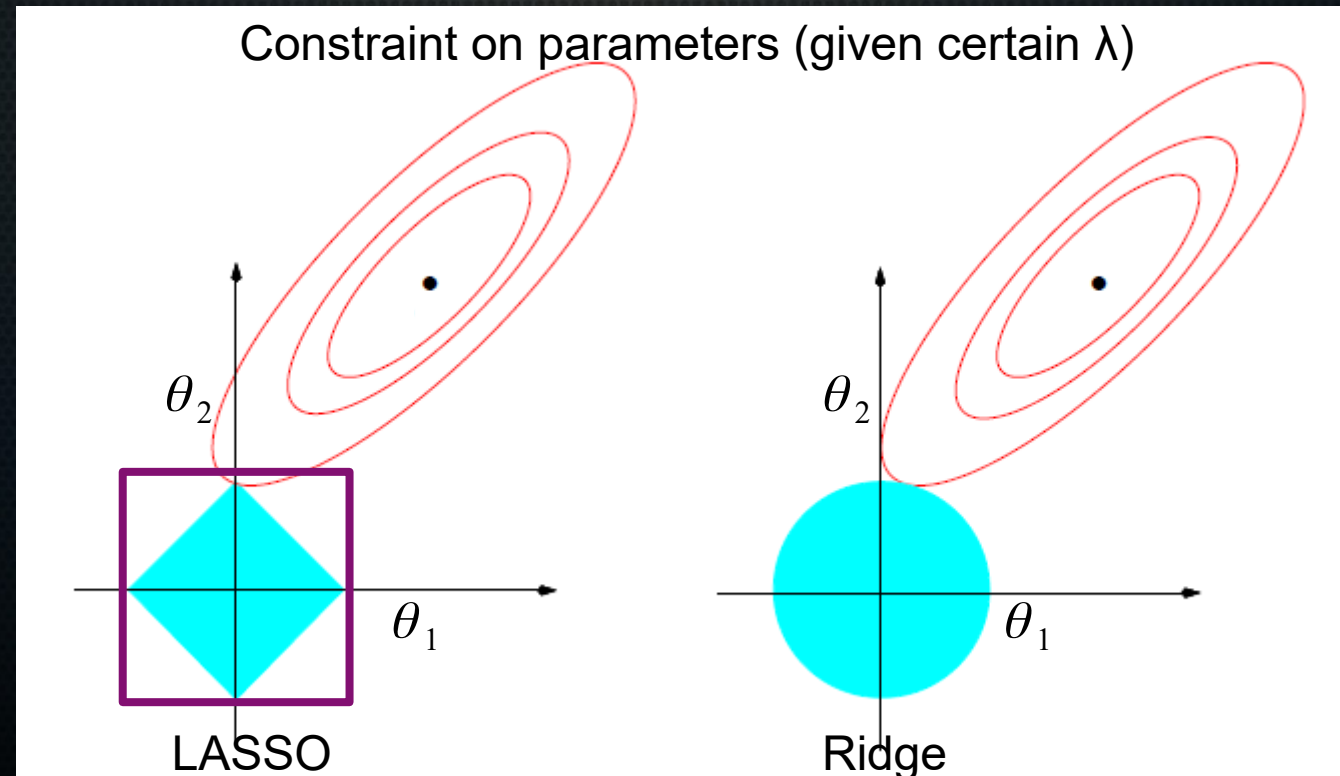
Source: Hastie, T., Tibshirani, R., & Friedman, J. (2009). The elements of statistical learning: data mining, inference, and prediction. Springer Science & Business Media.

Regularisation: Ridge vs. LASSO regression

- Ridge regression *shrinks* parameters/weights to 0, LASSO can make them 0 outright, i.e. simply removes uninformative features. → Why?

$$\frac{1}{2m} \sum_{i=1}^m (h_{\theta}(x) - y^{(i)})^2$$

$$\lambda \sum_{j=1}^n |\theta_j|$$



Source: Hastie, T., Tibshirani, R., & Friedman, J. (2009). The elements of statistical learning: data mining, inference, and prediction. Springer Science & Business Media.

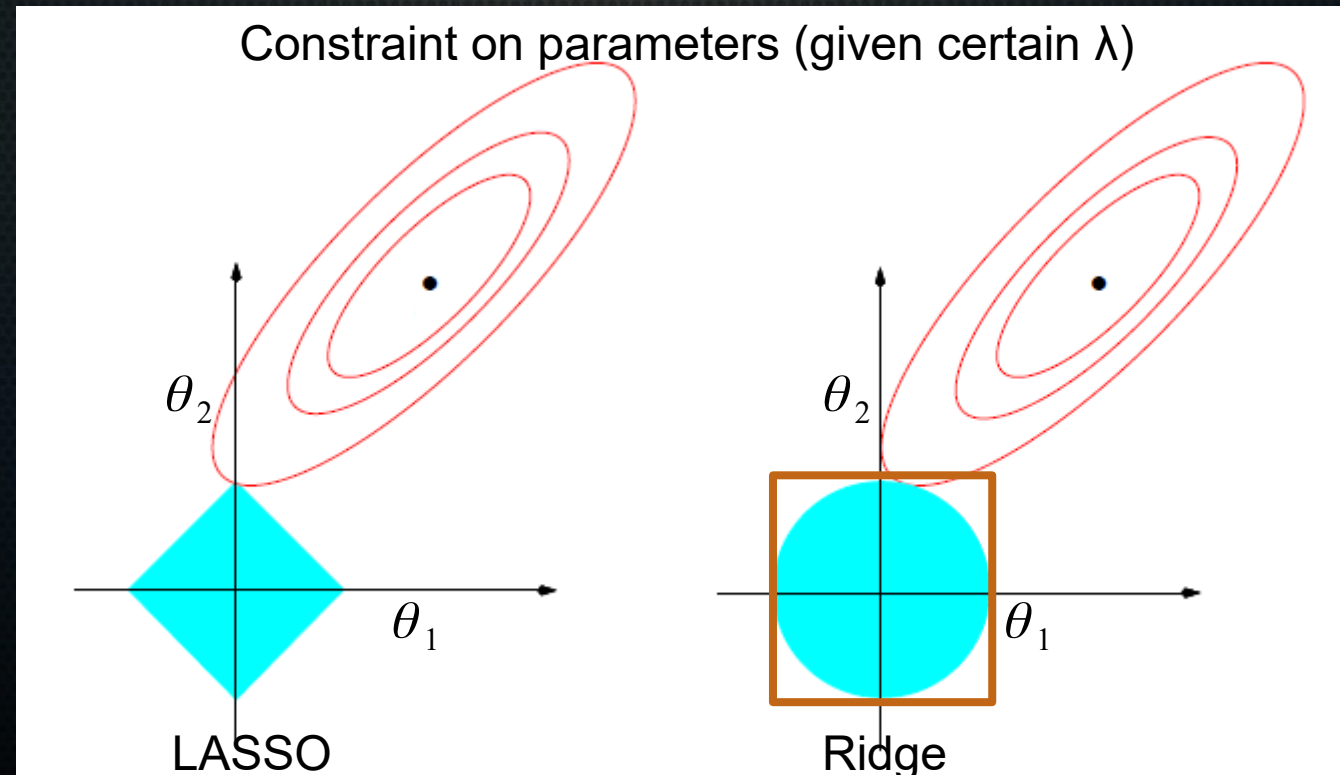
Regularisation: Ridge vs. LASSO regression

- Ridge regression *shrinks* parameters/weights to 0, LASSO can make them 0 outright, i.e. simply removes uninformative features. → Why?

$$\frac{1}{2m} \sum_{i=1}^m (h_{\theta}(x) - y^{(i)})^2$$

$$\lambda \sum_{j=1}^n |\theta_j|$$

$$\lambda \sum_{j=1}^n (\theta_j)^2$$



Source: Hastie, T., Tibshirani, R., & Friedman, J. (2009). The elements of statistical learning: data mining, inference, and prediction. Springer Science & Business Media.

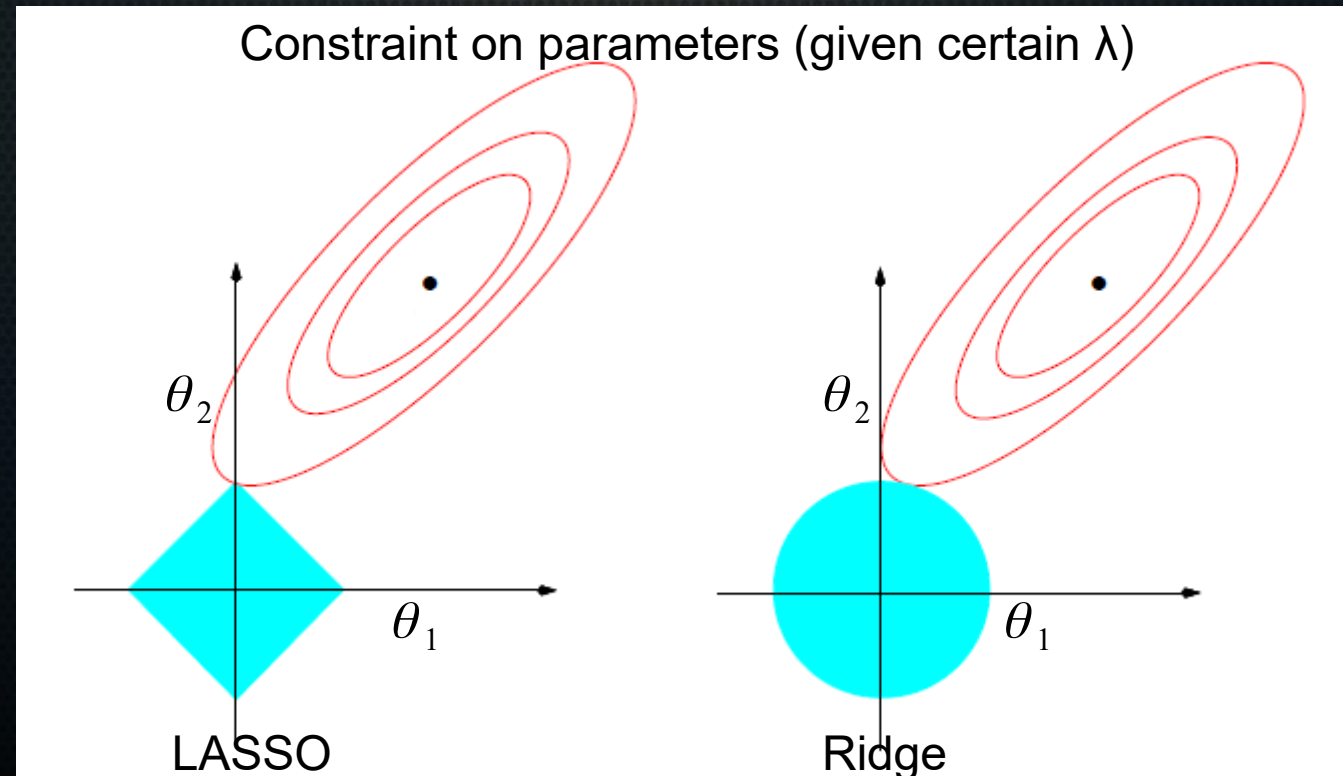
Regularisation: Ridge vs. LASSO regression

- Optimum = best least squares fit given constraint = intersect red lines with blue area.

$$\frac{1}{2m} \sum_{i=1}^m (h_{\theta}(x) - y^{(i)})^2$$

$$\lambda \sum_{j=1}^n |\theta_j|$$

$$\lambda \sum_{j=1}^n (\theta_j)^2$$



Source: Hastie, T., Tibshirani, R., & Friedman, J. (2009). The elements of statistical learning: data mining, inference, and prediction. Springer Science & Business Media.

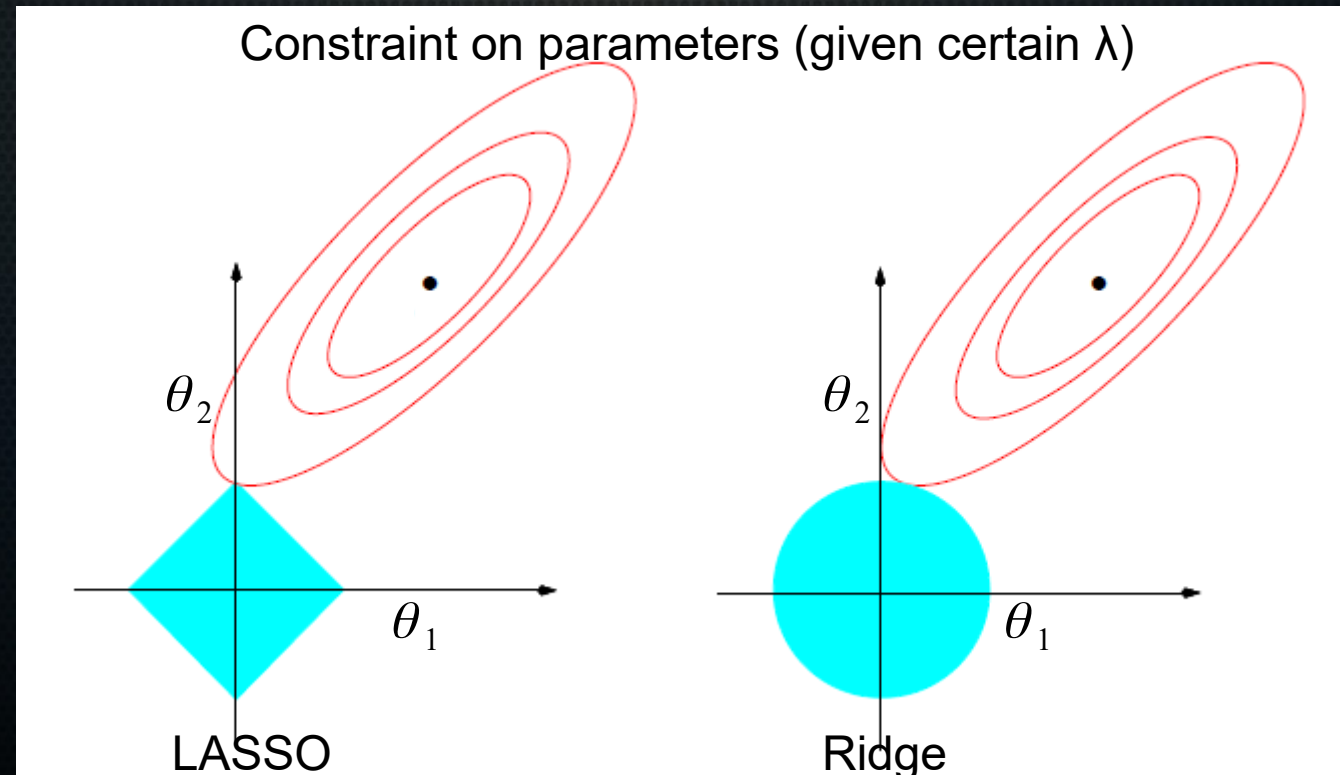
Regularisation: Ridge vs. LASSO regression

- Optimum = best least squares fit given constraint = intersect red lines with blue area.
- LASSO: can be at tip of rhombus, where $\theta_1 = 0$.
- Ridge: intersection ellips with circle never where either = 0

$$\frac{1}{2m} \sum_{i=1}^m (h_{\theta}(x) - y^{(i)})^2$$

$$\lambda \sum_{j=1}^n |\theta_j|$$

$$\lambda \sum_{j=1}^n (\theta_j)^2$$



Source: Hastie, T., Tibshirani, R., & Friedman, J. (2009). The elements of statistical learning: data mining, inference, and prediction. Springer Science & Business Media.

What about λ ?

$$J(\theta_0, \dots, \theta_n) = \frac{1}{2m} \sum_{i=1}^m (h_{\theta}(x) - y^{(i)})^2 + \lambda \sum_{j=1}^n |\theta_j|$$

- Lambda is another hyperparameter
- Intuition: higher lambda \rightarrow constrain parameters more, i.e. increase **bias**.
lower lambda \rightarrow constrain parameters less, i.e. increase **variance**

What about λ ?

$$J(\theta_0, \dots, \theta_n) = \frac{1}{2m} \sum_{i=1}^m (h_{\theta}(x) - y^{(i)})^2 + \lambda \sum_{j=1}^n |\theta_j|$$

- How to pick a good value?

What about λ ?

$$J(\theta_0, \dots, \theta_n) = \frac{1}{2m} \sum_{i=1}^m (h_{\theta}(x) - y^{(i)})^2 + \lambda \sum_{j=1}^n |\theta_j|$$

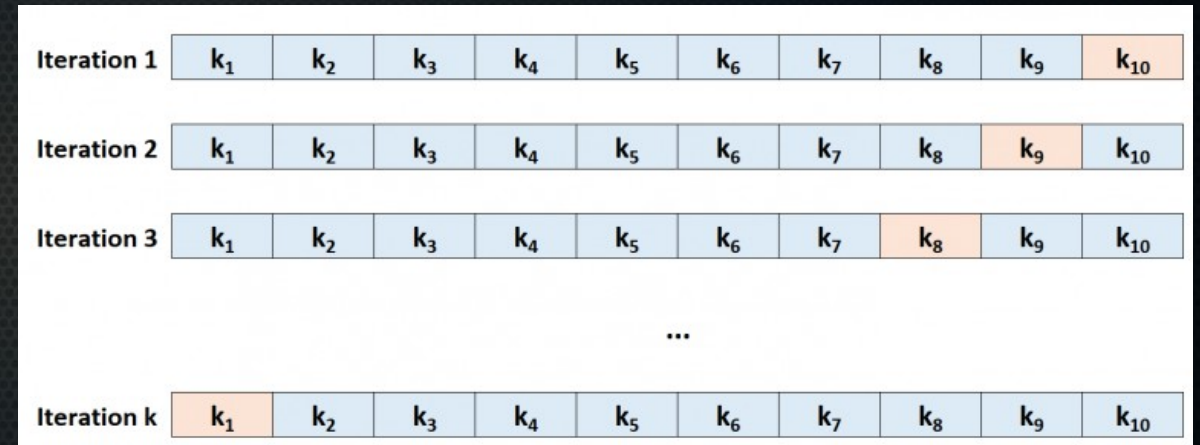
- How to pick a good value?
- Cross-validation!

Iteration 1	k ₁	k ₂	k ₃	k ₄	k ₅	k ₆	k ₇	k ₈	k ₉	k ₁₀
Iteration 2	k ₁	k ₂	k ₃	k ₄	k ₅	k ₆	k ₇	k ₈	k ₉	k ₁₀
Iteration 3	k ₁	k ₂	k ₃	k ₄	k ₅	k ₆	k ₇	k ₈	k ₉	k ₁₀
...										
Iteration k	k ₁	k ₂	k ₃	k ₄	k ₅	k ₆	k ₇	k ₈	k ₉	k ₁₀

What about λ ?

$$J(\theta_0, \dots, \theta_n) = \frac{1}{2m} \sum_{i=1}^m (h_{\theta}(x) - y^{(i)})^2 + \lambda \sum_{j=1}^n |\theta_j|$$

- How to pick a good value?
- Cross-validation!
 - Vary lambda over orders of magnitude (0.3, 1, 3, 7, 10, 30, 100)
 - For each lambda:
 - Train on training set using regularisation term



What about λ ?

$$J(\theta_0, \dots, \theta_n) = \frac{1}{2m} \sum_{i=1}^m (h_{\theta}(x) - y^{(i)})^2 + \lambda \sum_{j=1}^n |\theta_j|$$

- How to pick a good value?
- Cross-validation!
 - Vary lambda over orders of magnitude (0.3, 1, 3, 7, 10, 30, 100)
 - For each lambda:
 - Train on training set using regularisation term
 - Calculate validation error ***without regularisation term***



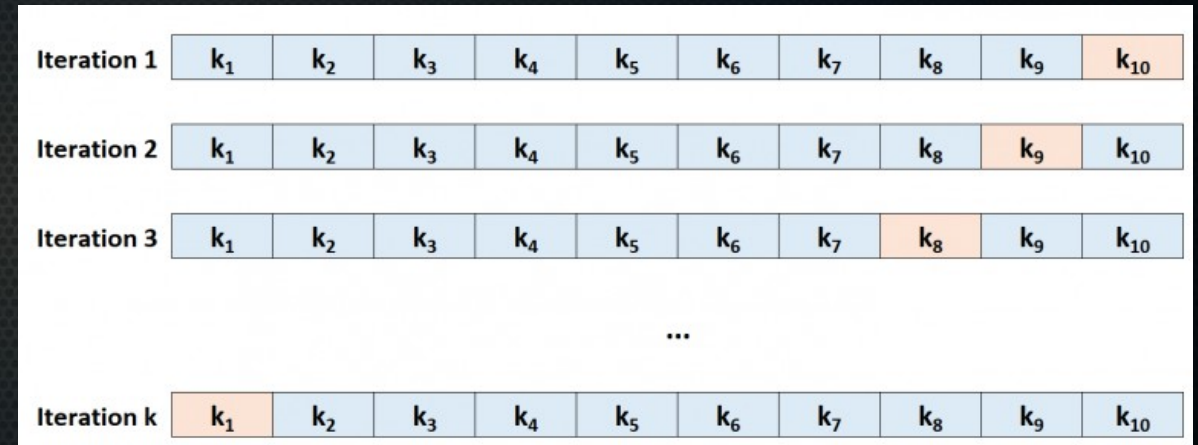
Don't care about parameter sizes,
just care about how good your
predictions are!



What about λ ?

$$J(\theta_0, \dots, \theta_n) = \frac{1}{2m} \sum_{i=1}^m (h_{\theta}(x) - y^{(i)})^2 + \lambda \sum_{j=1}^n |\theta_j|$$

- How to pick a good value?
- Cross-validation!
 - Vary lambda over orders of magnitude (0.3, 1, 3, 7, 10, 30, 100)
 - For each lambda:
 - Train on training set using regularisation term
 - Calculate validation error ***without regularisation term***
 - Pick lambda that gives best (average) validation error!



What can we do to find a good model?

- ~~Find a way to approximate generalisation error: how well do you do on unseen data?~~
- ~~See how error on seen and unseen data changes with amount of training data (plot learning curves)~~
- ~~Automatically constrain the fitting by penalising the cost function for too many/too large parameters~~

All done!

Summary

- We want our models to generalise well but have to contend with **bias** and **variance**
- We can measure (by proxy) how well we generalise using cross-validation
- We can plot learning curves for different subsamplings of the data to diagnose bias and variance
- To avoid overfitting we can use regularisation terms in the cost function, (slightly) increasing bias but decreasing variance
- Hyperparameters (such as λ) can be chosen by trying different ones and seeing what gives best results on the validation data